

ejabberd Docs

ejabberd Community Server

Copyright © 2008 - 2024 ProcessOne

Table of contents

Overview	5
Getting started 🖐	5
Features	8
Frequently Asked Questions	10
ejabberd Use Cases	12
GNU GENERAL PUBLIC LICENSE	15
Security Policy	19
Glossary	20
Readme	23
Install	25
Installation	25
ejabberd Container Images	26
Binary Installers	42
Operating System Packages	44
Install ejabberd from Source Code	45
Install ejabberd on macOS	53
Installing ejabberd development environment on OSX	54
Next Steps	57
Configure	59
Configuring ejabberd	59
File format	60
Basic Configuration	66
Authentication	77
Database Configuration	82
LDAP Configuration	87
Listen Modules	97
Listen Options	106
Top-Level Options	113
Modules Options	141
Advanced	0
Advanced ejabberd Administration	0
Architecture	0
Clustering	0
Understanding ejabberd and its dependencies	0
Erlang Distribution	0

Managing an ejabberd server	0
Get More Modules	0
Securing ejabberd	0
Troubleshooting ejabberd	0
Upgrade Procedure for ejabberd	0
ejabberd and XMPP tutorials	0
Getting started with MIX	0
MQTT Support	0
MUC Hats	0
Setting vCards / Avatars for MUC rooms	0
Using ejabberd with MySQL	0
Development	0
ejabberd for Developers	0
ejabberd Developer Guide	0
PubSub overview	0
Roster versioning	0
ejabberd Stanza Routing	0
ejabberd SQL Database Schema	0
External authentication	0
Main contribution repository	0
ejabberd API libraries	0
Old / obsolete contributions	0
Contributing to ejabberd	0
Contributor Covenant Code of Conduct	0
Contributors	0
ejabberd Docs Source Code	0
ejabberd for Elixir Developers	0
The ejabberd Developer Livebook	0
Internationalization and Localization	0
ejabberd Modules Development	0
MucSub: Multi-User Chat Subscriptions	0
ejabberd Test Suites	0
Developing ejabberd with VSCode	0
Getting Started with XMPPFramework	0
API	0
ejabberd ReST API	0
API Reference	0
API Tags	0

Simple ejabberd Rest API Configuration	0
API Permissions	0
OAuth Support	0
ejabberd commands	0
API Versioning	0
Archive	0
ChangeLog	0
Roadmap	0
ejabberd Roadmap	0

Overview

Getting started 🙌

Meet **ejabberd**, your superpowerful messaging framework

This web site is dedicated to help you use and develop for ejabberd [XMPP](#) messaging server.

ejabberd has been in development since 2002 and is used all over the world to power the largest XMPP deployments. This project is so versatile that you can deploy it and customize it for very large scale, no matter what your use case is.

This incredible power comes with a price. You need to learn how to leverage it. Fortunately, the goal of this website is to get you started on your path to mastery. Whether you are a sysadmin, an architect, a developer planning to extend it, or even a simple XMPP user, we have something for you here.

Overview

ejabberd is the de facto XMPP server in the world. The fact that it is used to power the largest deployments in the world should not intimidate you. ejabberd is equally suitable for small instances.

ejabberd has been designed from the ground-up, since 2002 for robust, enterprise deployment. The goal has always been to shoot for the moon and this is what made it a long-lasting success.

ejabberd is specifically designed for enterprise purposes: it is fault-tolerant, can utilise the resources of multiple clustered machines, and can easily scale when more capacity is required (by just adding a box/VM).

Designed at a moment when clients were mostly desktops that only supported a kind of HTTP polling call BOSH, the project managed to adapt to recent changes by introducing support for WebSockets, BOSH improvements, and a solid mobile stack.

It was developed at a time when XMPP was still known as "Jabber", but quickly adopted an evolution process in order to support the various versions of XMPP RFCs. It now encourages innovation and experimentation by supporting most, if not all, extensions produced by the XSF.

ejabberd relies on a dynamic community all over the world. To get an idea of existing contributions, you can check [ejabberd main repository](#) or the repository containing a great amount of [contributed extensions](#).

This is possible thanks to a modular architecture based on a core router and an extremely powerful plugin mechanism that is getting richer every day.

Welcome to the beginning of your journey of ejabberd mastery!

Options to use ejabberd

ejabberd can be used in different ways. The most common one is:

- **[ejabberd Community Server](#)**, the public Open Source version that everyone loves: highly scalable and flexible.

Fortunately, if you need more than just the ejabberd platform software, [ProcessOne](#) can help you with a commercial offering. Commercial offering come in two type of packaging:

- **[ejabberd Business Edition](#)** including features for large companies (enhanced geodistributed companies and mobile support to develop own, rich clients) and world-class support, that can please even the most demanding businesses, with 24/7 options.
- **[Fluux.io](#)** being a way to access and benefit of all the features of ejabberd Business Edition at an attractive and scalable price. Fluux.io allows you to keep control of your data thanks to integration API you can implement on your backend to become a data provider for ejabberd SaaS.

Whatever approach you choose, you can hardly make the wrong choice with ejabberd! In every case you can easily integrate ejabberd with your existing application using:

- [REST API](#) and ejabberdctl command-line tool
- Mobile libraries for iOS: [XMPPFramework](#), [Jayme REST API](#)
- Mobile libraries for Android: [Smack](#), [Retrofit](#)
- Web library with WebSocket support and fallback to BOSH: [Strophe](#)

Architecture of an ejabberd service

ejabberd brings configurability, scalability and fault-tolerance to the core feature of XMPP – routing messages.

Its architecture is based on a set of pluggable modules that enable different features, including:

- One-to-one messaging
- Store-and-forward (offline messages)
- Contact list (roster) and presence
- Groupchat: MUC (Multi-User Chat)
- Messaging archiving with Message Archive Management (MAM)
- User presence extension: Personal Event Protocol (PEP) and typing indicator
- Privacy settings, through privacy list and simple blocking extensions
- User profile with vCards
- Full feature web support, with BOSH and websockets
- Stream management for message reliability on mobile (aka XEP-0198)
- Message Delivery Receipts (aka XEP-184)
- Last activity
- Metrics and full command-line administration
- and many many more.

The full list of supported protocol and extensions in the [Specification & Reference Sheets](#) page.

This modular architecture allows high customisability and easy access to the required features.

ejabberd enables authenticating users using external or internal databases (Mnesia, SQL), LDAP or external scripts. It also allows connecting anonymous users, when required.

For storing persistent data, ejabberd uses Mnesia (the distributed internal Erlang database), but you can opt for SQL database like MySQL or PostgreSQL

And of course, thanks to its API, ejabberd can be customised to work with a database chosen by the customer.

Deploying and managing an ejabberd service

ejabberd can be deployed for a number of scenarios fitting end-user / developer / customer needs. The default installation setup consists of a single ejabberd node using Mnesia, so it does not require any additional configuration. This primary system is sufficient for fast deployment and connecting XMPP clients. It should be good enough for most of the small deployments (and even medium ones).

A more scalable solution would be deploying ejabberd with an external database for persistent data. As Mnesia is caching part of its data in ejabberd memory (actually in Erlang VM node), this kind of setup make your system more scalable and typically easier to integrate with your usual database. As a sysadmin, yes, you can use your standard backup process.

Those larger setup can run as a cluster of ejabberd nodes. This is a clustering mode where all nodes are active, so it can be use for fault-tolerance, but also to increase the capacity of your ejabberd deployment.

With such a deployment you can load balance the traffic to your cluster node using one of the following solution:

- traditional TCP/IP load balancer (beware of the cost of your solution, typical XMPP connections are persistent).
- DNS load balancing.
- Custom approach that requires client cooperation.

If deployed on a 16 GB RAM machine with at least 4 cores, a single ejabberd node can typically handle 200-300 K online users. This setup is suitable for systems with up to 10 nodes.

Note that your mileage may vary depending on your use case, the feature your are using and how clean the architecture design and the client is developed. That's why, if you plan to reach huge volume, it is recommended to start asking advices from day 1 to an [ejabberd expert](#). Initial mistakes in the solution design are harder to fix once the project is in production.

If the service requires a cluster of more than 10 nodes, we recommend not relying on Mnesia clustering mode. Many solutions are available, the easiest and more inexpensive being to rely on ejabberd Software-as-a-Service approach, see [Advanced Messaging Infrastructure & Services](#).

ejabberd also allows connecting different clusters as parts of larger systems. This is a standard XMPP feature call server-to-server (aka s2s in XMPP lingo). It is used in geo-localised services handling massive traffic from all over the world. Special extension are also available from ProcessOne to handle geodistribution in an even more robust way.

To manage the users, rosters, messages and general settings, we provide a command-line tool, ejabberdctl. That command-line allows you to gather metrics from ejabberd to be able to monitor what is happening in your system, understand and even anticipate issues.

The main benefit of ejabberd is the ability to reach a command-line to type Erlang commands. This allows you to fix and troubleshoot most of the tricky situation and even update and reload code without stopping the service. This is a life saver for your uptime.

Welcome to the benefit of Erlang hot-code swapping!

ejabberd is more than XMPP

Thanks to the modular architecture of ejabberd, the platform is becoming a core component for messaging applications.

Messaging applications require to transfer more than text messages. ejabberd has grow a full set of media related features that makes ejabberd a great choice to support voice and video applications, but also to proxy various kind of media transfer (images, audio and video files for example).

As such, ejabberd support:

- Jingle, XMPP based voice protocol
- SIP (Session Initiation Protocol): Yes, you can pass SIP calls using ejabberd :)
- ICE (Interactive Connectivity Establishment: A Protocol for Network Address Translator (NAT) Traversal)
- STUN
- TURN
- Proxy65 media relay

This makes ejabberd the best XMPP server to support SIP and WebRTC based communication tools.

Helping us in the development process

With thousands of more or less official forks, the core ejabberd team, supported by ProcessOne, is constantly monitoring and reviewing improvements. We use our 15 years of experience to filter the best ideas or improvements to make sure ejabberd is always your most solid choice in term of scalability, robustness and manageability.

The best way to start developing for ejabberd is to clone, watch and star the [project](#), to get in touch on our developer chatroom (ejabberd@conference.process-one.net) or to join [ejabberd community on StackOverflow](#).

Features

ejabberd is a *free and open source* instant messaging server written in [Erlang/OTP](#).

ejabberd is *cross-platform*, distributed, fault-tolerant, and based on open standards to achieve real-time communication.

ejabberd is designed to be a *rock-solid and feature rich* [XMPP](#) server.

ejabberd is suitable for small deployments, whether they need to be *scalable* or not, as well as extremely large deployments.

Check also the features in [ejabberd.im](#), [ejabberd at ProcessOne](#), and the list of [supported protocols in ProcessOne](#) and [XMPP.org](#).

Key Features

ejabberd is:

- *Cross-platform*: **ejabberd** runs under Microsoft Windows and Unix-derived systems such as Linux, FreeBSD and NetBSD.
- *Distributed*: You can run **ejabberd** on a cluster of machines all serving the same Jabber domain(s). When you need more capacity you can simply add a new cheap node to your cluster. Accordingly, you do not need to buy an expensive high-end machine to support tens of thousands concurrent users.
- *Fault-tolerant*: You can deploy an **ejabberd** cluster so that all the information required for a properly working service will be replicated permanently on all nodes. This means that if one of the nodes crashes, the others will continue working without disruption. In addition, nodes can be added or replaced on the fly.
- *Administrator Friendly*: **ejabberd** is built on top of the Erlang programming language. As a result, if you wish, you can perform self-contained deployments. You are not required to install an external database, an external web server, amongst others because everything is already included, and ready to run out of the box. Other administrator benefits include:
 - Comprehensive documentation.
 - Straightforward installers for Linux, Mac OS X, and Windows.
 - Web Administration.
 - Shared Roster Groups.
 - Command line administration tool.
 - Can integrate with existing authentication mechanisms.
 - Capability to send announce messages.
- *Internationalized*: **ejabberd** leads in internationalization and is well suited to build services available across the world. Related features are:
 - Translated to 25 languages.
 - Support for [IDNA](#).
- *Open Standards*: **ejabberd** is the first Open Source Jabber server staking a claiming to full compliance to the XMPP standard.
- Fully XMPP compliant.
- XML-based protocol.
- [Many protocols supported](#).

Additional Features

ejabberd also comes with a wide range of other state-of-the-art features:

- Modular
- Load only the modules you want.
- Extend **ejabberd** with your own custom modules.
- Security
- SASL and STARTTLS for c2s and s2s connections.
- STARTTLS and Dialback s2s connections.
- Web Admin accessible via HTTPS secure access.
- Databases
- Internal database for fast deployment (Mnesia).
- Native MySQL support.
- Native PostgreSQL support.
- ODBC data storage support.
- Microsoft SQL Server support.
- SQLite support.
- Authentication
- Internal Authentication.
- PAM, LDAP and SQL.
- External Authentication script.
- Others
- Support for virtual hosting.
- Compressing XML streams with Stream Compression ([XEP-0138](#)).
- Statistics via Statistics Gathering ([XEP-0039](#)).
- IPv6 support both for c2s and s2s connections.
- [Multi-User Chat](#) module with support for clustering and HTML logging.
- Users Directory based on users vCards.
- [Publish-Subscribe](#) component with support for [Personal Eventing via Pubsub](#) .
- Support for web clients: Support for [XMPP subprotocol for WebSocket](#) and [HTTP Binding \(BOSH\)](#) services.
- SIP support.
- Component support: interface with networks such as AIM, ICQ, IRC and MSN installing special transports.

Frequently Asked Questions

Development process

Why is there a commercial version of ejabberd?

Different needs for different users. Corporations and large scale deployments are very different from smaller deployments and community projects.

While we put a huge development effort to have a product that is on the edge of innovation with ejabberd community version, we are requested to provide a stable version with long term commitment and high level of quality, testing, audit, etc.

Maintaining such a version in parallel to the community version, along with extremely strong commitment in terms of availability and 24/7 support has a tangible cost. With ejabberd business edition we commit to a level of scalability and optimize the software until it is performing to the level agreed with the customer.

Covering all those costs, along with all our Research and Development cost on ejabberd community in general is the real reason we need a business edition.

The business edition is also a way for our customers to share the code between our customers only, thus retaining a huge competitive edge for a limited time (See next section).

So, even if you are not using our business edition, this is a great benefit for you as a user of the community edition and the reason you have seen so many improvements since 2002. Thanks to our business edition customers, ejabberd project itself is a [major contributor to Erlang and Elixir community](#).

Does ProcessOne voluntarily hold some code in ejabberd community to push toward the business edition?

No. We never do that and have no plan doing so with the code we produce and we own.

However, when the code is paid by customer, they retain the ownership of the code. Part of our agreement is that the code produced for them will be limited to a restricted user base, ejabberd business edition until an agreed time expires, generally between 6 months and 1 year.

This is extremely important for both the users of the commercial edition and the users of the community edition:

- For the business edition customers, this is a way to keep their business advantage. This is fair as they paid for the development.
- This is also a great incentive for our customers as they benefit from development for other customers (I guess they agree for the reciprocal sharing of their own code with customers).
- This is fair for the community as the community edition users know they will benefit from new extremely advanced features in a relatively near future. For example, [websocket module](#) was contributed to ejabberd community as part of this process.

This is the model we have found to be fair to our broader user base and lets us produce an amazing code base that benefits all our users.

This dual model is the core strength of our approach and our secret sauce to make sure everyone benefits.

Performance

Is ejabberd the most scalable version?

Yes. Definitely. Despite claims that there is small change you can make to make it more scalable, we already performed the changes during the past year, that make those claims unfunded:

- ejabberd reduced memory consumption in 2013 by switching to binary representation of string instead of list. This can reduce given string by 8.
- We have improved the C code performance a lot, using new Erlang NIF. This provides better performance, removes bottlenecks
- We have a different clustering mechanism available to make sure we can scale to large clusters

This is a common misconception, but our feedback for production service on various customer sites shows that ejabberd is the most scalable version. Once it is properly configured, optimized and tuned, you can handle tens of millions of users on ejabberd systems.

... And we are still improving :)

As a reference, you should read the following blog post: [ejabberd Massive Scalability: 1 Node — 2+ Million Concurrent Users](#)

What are the tips to optimize performance?

Optimisation of XMPP servers performance, including ejabberd, is highly dependent on the use case. You really need to find your bottleneck(s) by monitoring the process queues, finding out what is your limiting factor, tune that and then move to the next one.

The first step is to make sure you run the latest ejabberd. Each new release comes with a bunch of optimisations and a specific bottleneck you are facing may have gone away in the latest version.

For perspective, ejabberd 15.07 is 2 to 3 times more efficient in memory, latency and CPU compared to ejabberd 2.1.

You should also make sure that you are using the latest Erlang version. Each release of Erlang comes with more optimisation regarding locks, especially on SMP servers, and using the latest Erlang version can also help tremendously.

Erlang support

Is ejabberd conforming to the best Erlang practices?

Yes. Our build system is primarily based on rebar. However, as we are multiplatform and need to run in many various environments, we rely on a toolchain that can detect required library dependencies using autotools.

This gives developers and admins the best of both worlds. A very flexible and very versatile build chain, that is very adequate to make sure ejabberd can be used in most operating systems and even integrated in Linux distributions.

ejabberd Use Cases

ejabberd is very versatile and is a solid choice to build messaging services across a large number of industries:

ejabberd

Mobile messaging

ejabberd's massive scalability makes it the most solid choice as the backbone for a very large number of mobile messaging services.

This includes:

- Chaatz
- [Libon](#)
- [Nokia OVI Chat](#)
- Roo Kids : Safe & fun instant messaging app for kids with minimum yet critical parental controls.
- Swisscom IO
- Versapp
- [Whatsapp](#)

Gaming

- [Electronic Arts](#)
- [FACEIT](#)
- [Kixeye](#)
- [Machine Zone \(Game of War\)](#)
- [Nokia nGage](#)
- [Riot Games \(League of Legends\)](#)

Voice and video messaging

- [Nimbuzz](#)
- ooVoo
- [Sipphone](#)
- WowApp

Internet of Things

- AeroFS
- IMA Teleassistance
- [Nabaztag](#) (Violet, Mindscape, then Aldebaran Robotics)

Telecom / Hosting

- [Fastmail](#)
- GMX
- [Mailfence](#)
- Orange

- [SAPO - Portugal Telecom](#)

Customer chat / CRM

- CoBrowser.net: [Coder Interview](#).
- iAdvize
- LiveHelperChat: [Blog post: Full XMPP chat support for ejabberd](#)

Media

- [AFP](#)
- [BBC](#)

Social media

- [Facebook](#)
- Nasza Klasa (NKTalk messenger)
- [StudiVZ](#)
- [Sify](#)
- [Tuenti](#)

Sport

- [Major League of Baseball \(MLB\)](#)

Education

- Apollo group
- Laureate

Push alerts

- [Nokia push notifications](#)
- [Notify.me](#))

Dating

- Grindr
- [Meetic](#)

Community sites

- Jabber.at
- Talkr.im

XMPP Use Cases

[XMPP](#) is a very versatile protocol designed to address many use cases of modern real-time messaging needs. However, it is also a very large protocol and it is difficult to understand at first sight all the use cases that XMPP adequately addresses.

This page is gathering XMPP specifications that make XMPP a good fit for a given use case of industry.

Realtime web

XMPP was designed before the advent of realtime web. However, it managed to adapt thanks to the following specifications:

- XMPP PubSub is defined in [XEP-0060](#). This is a very powerful mechanism that defines channel based communication on top of the XMPP protocol itself. A server can handle millions of channels, called Pubsub nodes. Users interested in specific channels can subscribe to nodes. When data needs to be send to a given channel, authorized publishers can send data to that node. The XMPP server will then broadcast the content to all subscribers. This is very adequate for realtime web as it allows you to broadcast relevant events to web pages.
- WebSocket: XMPP over WebSocket is defined in [RFC 7395](#). It is more efficient and more scalable than XMPP for web's previous specifications called [BOSH](#). WebSocket being a true bidirectional channel, it allows lower latency messaging and is very reliable. Note that BOSH can still be used transparently along with WebSocket to support old web browsers.

Use cases: News, interactive web page, web chat, web games.

Supported by ejabberd: Yes.

As a special exception, the authors give permission to link this program with the OpenSSL library and distribute the resulting binary.

GNU GENERAL PUBLIC LICENSE

Version 2, June 1991

Copyright (C) 1989, 1991 Free Software Foundation, Inc.
51 Franklin Street, Fifth Floor, Boston, MA 02110-1301, USA

Everyone is permitted to copy and distribute verbatim copies
of this license document, but changing it is not allowed.

Preamble

The licenses for most software are designed to take away your freedom to share and change it. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change free software--to make sure the software is free for all its users. This General Public License applies to most of the Free Software Foundation's software and to any other program whose authors commit to using it. (Some other Free Software Foundation software is covered by the GNU Lesser General Public License instead.) You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for this service if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs; and that you know you can do these things.

To protect your rights, we need to make restrictions that forbid anyone to deny you these rights or to ask you to surrender the rights. These restrictions translate to certain responsibilities for you if you distribute copies of the software, or if you modify it.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must give the recipients all the rights that you have. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

We protect your rights with two steps: (1) copyright the software, and (2) offer you this license which gives you legal permission to copy, distribute and/or modify the software.

Also, for each author's protection and ours, we want to make certain that everyone understands that there is no warranty for this free software. If the software is modified by someone else and passed on, we want its recipients to know that what they have is not the original, so that any problems introduced by others will not reflect on the original authors' reputations.

Finally, any free program is threatened constantly by software patents. We wish to avoid the danger that redistributors of a free program will individually obtain patent licenses, in effect making the program proprietary. To prevent this, we have made it clear that any patent must be licensed for everyone's free use or not licensed at all.

The precise terms and conditions for copying, distribution and modification follow.

TERMS AND CONDITIONS FOR COPYING, DISTRIBUTION AND MODIFICATION

0. This License applies to any program or other work which contains a notice placed by the copyright holder saying it may be distributed under the terms of this General Public License. The "Program", below, refers to any such program or work, and a "work based on the Program" means either the Program or any derivative work under copyright law: that is to say, a work containing the Program or a portion of it, either verbatim or with modifications and/or translated into another language. (Hereinafter, translation is included without limitation in the term "modification".) Each licensee is addressed as "you".

Activities other than copying, distribution and modification are not covered by this License; they are outside its scope. The act of running the Program is not restricted, and the output from the Program is covered only if its contents constitute a work based on the Program (independent of having been made by running the Program). Whether that is true depends on what the Program does.

1. You may copy and distribute verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice and disclaimer of warranty; keep intact all the notices that refer to this License and to the absence of any warranty; and give any other recipients of the Program a copy of this License along with the Program.

You may charge a fee for the physical act of transferring a copy, and you may at your option offer warranty protection in exchange for a fee.

2. You may modify your copy or copies of the Program or any portion of it, thus forming a work based on the Program, and copy and distribute such modifications or work under the terms of Section 1 above, provided that you also meet all of these conditions:

- a)** You must cause the modified files to carry prominent notices stating that you changed the files and the date of any change.
- b)** You must cause any work that you distribute or publish, that in whole or in part contains or is derived from the Program or any part thereof, to be licensed as a whole at no charge to all third parties under the terms of this License.
- c)** If the modified program normally reads commands interactively when run, you must cause it, when started running for such interactive use in the most ordinary way, to print or display an announcement including an appropriate copyright notice and a notice that there is no warranty (or else, saying that you provide a warranty) and that users may redistribute the program under these conditions, and telling the user how to view a copy of this License. (Exception: if the Program itself is interactive but does not normally print such an announcement, your work based on the Program is not required to print an announcement.)

These requirements apply to the modified work as a whole. If identifiable sections of that work are not derived from the Program, and can be reasonably considered independent and separate works in themselves, then this License, and its terms, do not apply to those sections when you distribute them as separate works. But when you distribute the same sections as part of a whole which is a work based on the Program, the distribution of the whole must be on the terms of this License, whose permissions for other licensees extend to the entire whole, and thus to each and every part regardless of who wrote it.

Thus, it is not the intent of this section to claim rights or contest your rights to work written entirely by you; rather, the intent is to exercise the right to control the distribution of derivative or collective works based on the Program.

In addition, mere aggregation of another work not based on the Program with the Program (or with a work based on the Program) on a volume of a storage or distribution medium does not bring the other work under the scope of this License.

3. You may copy and distribute the Program (or a work based on it, under Section 2) in object code or executable form under the terms of Sections 1 and 2 above provided that you also do one of the following:

- a)** Accompany it with the complete corresponding machine-readable source code, which must be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
- b)** Accompany it with a written offer, valid for at least three years, to give any third party, for a charge no more than your cost of physically performing source distribution, a complete machine-readable copy of the corresponding source code, to be distributed under the terms of Sections 1 and 2 above on a medium customarily used for software interchange; or,
- c)** Accompany it with the information you received as to the offer to distribute corresponding source code. (This alternative is allowed only for noncommercial distribution and only if you received the program in object code or executable form with such an offer, in accord with Subsection b above.)

The source code for a work means the preferred form of the work for making modifications to it. For an executable work, complete source code means all the source code for all modules it contains, plus any associated interface definition files, plus the scripts used to control compilation and installation of the executable. However, as a special exception, the source code distributed need not include anything that is normally distributed (in either source or binary form) with the major components (compiler, kernel, and so on) of the operating system on which the executable runs, unless that component itself accompanies the executable.

If distribution of executable or object code is made by offering access to copy from a designated place, then offering equivalent access to copy the source code from the same place counts as distribution of the source code, even though third parties are not compelled to copy the source along with the object code.

4. You may not copy, modify, sublicense, or distribute the Program except as expressly provided under this License. Any attempt otherwise to copy, modify, sublicense or distribute the Program is void, and will automatically terminate your rights under this

License. However, parties who have received copies, or rights, from you under this License will not have their licenses terminated so long as such parties remain in full compliance.

5. You are not required to accept this License, since you have not signed it. However, nothing else grants you permission to modify or distribute the Program or its derivative works. These actions are prohibited by law if you do not accept this License. Therefore, by modifying or distributing the Program (or any work based on the Program), you indicate your acceptance of this License to do so, and all its terms and conditions for copying, distributing or modifying the Program or works based on it.

6. Each time you redistribute the Program (or any work based on the Program), the recipient automatically receives a license from the original licensor to copy, distribute or modify the Program subject to these terms and conditions. You may not impose any further restrictions on the recipients' exercise of the rights granted herein. You are not responsible for enforcing compliance by third parties to this License.

7. If, as a consequence of a court judgment or allegation of patent infringement or for any other reason (not limited to patent issues), conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot distribute so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not distribute the Program at all. For example, if a patent license would not permit royalty-free redistribution of the Program by all those who receive copies directly or indirectly through you, then the only way you could satisfy both it and this License would be to refrain entirely from distribution of the Program.

If any portion of this section is held invalid or unenforceable under any particular circumstance, the balance of the section is intended to apply and the section as a whole is intended to apply in other circumstances.

It is not the purpose of this section to induce you to infringe any patents or other property right claims or to contest validity of any such claims; this section has the sole purpose of protecting the integrity of the free software distribution system, which is implemented by public license practices. Many people have made generous contributions to the wide range of software distributed through that system in reliance on consistent application of that system; it is up to the author/donor to decide if he or she is willing to distribute software through any other system and a licensee cannot impose that choice.

This section is intended to make thoroughly clear what is believed to be a consequence of the rest of this License.

8. If the distribution and/or use of the Program is restricted in certain countries either by patents or by copyrighted interfaces, the original copyright holder who places the Program under this License may add an explicit geographical distribution limitation excluding those countries, so that distribution is permitted only in or among countries not thus excluded. In such case, this License incorporates the limitation as if written in the body of this License.

9. The Free Software Foundation may publish revised and/or new versions of the General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies a version number of this License which applies to it and "any later version", you have the option of following the terms and conditions either of that version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of this License, you may choose any version ever published by the Free Software Foundation.

10. If you wish to incorporate parts of the Program into other free programs whose distribution conditions are different, write to the author to ask for permission. For software which is copyrighted by the Free Software Foundation, write to the Free Software Foundation; we sometimes make exceptions for this. Our decision will be guided by the two goals of preserving the free status of all derivatives of our free software and of promoting the sharing and reuse of software generally.

NO WARRANTY

11. BECAUSE THE PROGRAM IS LICENSED FREE OF CHARGE, THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

12. IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MAY MODIFY AND/OR REDISTRIBUTE THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

END OF TERMS AND CONDITIONS

How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively convey the exclusion of warranty; and each file should have at least the "copyright" line and a pointer to where the full notice is found.

```
one line to give the program's name and an idea of what it does.
Copyright (C) yyyy  name of author

This program is free software; you can redistribute it and/or
modify it under the terms of the GNU General Public License
as published by the Free Software Foundation; either version 2
of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful,
but WITHOUT ANY WARRANTY; without even the implied warranty of
MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.  See the
GNU General Public License for more details.

You should have received a copy of the GNU General Public License
along with this program; if not, write to the Free Software
Foundation, Inc., 51 Franklin Street, Fifth Floor, Boston, MA 02110-1301, USA.
```

Also add information on how to contact you by electronic and paper mail.

If the program is interactive, make it output a short notice like this when it starts in an interactive mode:

```
Gnomovision version 69, Copyright (C) year name of author
Gnomovision comes with ABSOLUTELY NO WARRANTY; for details
type `show w'. This is free software, and you are welcome
to redistribute it under certain conditions; type `show c'
for details.
```

The hypothetical commands `show w' and `show c' should show the appropriate parts of the General Public License. Of course, the commands you use may be called something other than `show w' and `show c'; they could even be mouse-clicks or menu items--whatever suits your program.

You should also get your employer (if you work as a programmer) or your school, if any, to sign a "copyright disclaimer" for the program, if necessary. Here is a sample; alter the names:

```
Yoyodyne, Inc., hereby disclaims all copyright
interest in the program `Gnomovision'
(which makes passes at compilers) written
by James Hacker.

signature of Ty Coon, 1 April 1989
Ty Coon, President of Vice
```

This General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the [GNU Lesser General Public License](#) instead of this License.

Security Policy

Supported Versions

We recommend that all users always use the latest version of ejabberd.

To ensure the best experience and security, upgrade to the latest version available on [this repo](#).

Reporting a Vulnerability

Private Reporting

Preferred Method: Use GitHub's private vulnerability reporting system by clicking the "Report a Vulnerability" button in the [Security tab of this repository](#). This ensures your report is securely transmitted and tracked.

Alternative: If you cannot use the GitHub system, send an email to contact@process-one.net with the following details:

- A clear description of the vulnerability.
- Steps to reproduce the issue.
- Any potential impact or exploitation scenarios.

Response Time

We aim to acknowledge receipt of your report within 72 hours. You can expect regular updates on the status of your report.

Resolution

If the vulnerability is confirmed, we will work on a patch or mitigation strategy.

We will notify you once the issue is resolved and coordinate a public disclosure if needed.

Acknowledgements

We value and appreciate the contributions of security researchers and community members.

If you wish, we are happy to acknowledge your efforts publicly by listing your name (or alias) below in this document.

Please let us know if you would like to be recognized when reporting the vulnerability.

Public Discussion

For general inquiries or discussions about the project's security, feel free to chat with us here:

- XMPP room: [ejabberd@conference.process-one.net](xmpp:ejabberd@conference.process-one.net)
- [GitHub Discussions](#)

However, please note that if the issue is **critical** or potentially exploitable, **do not share it publicly**. Instead, we strongly recommend you contact the maintainers directly via the private reporting methods outlined above to ensure a secure and timely response.

Thank you for helping us improve the security of ejabberd!

Glossary

- [ad-hoc command](#)
- [admin account](#)
- [api command](#)
- [bosh](#)
- [c2s](#)
- [command](#)
- [convert](#)
- [cookie](#)
- [eBE](#)
- [eCS](#)
- [ejabberdctl](#)
- [epmd](#)
- [erlang](#)
- [erlang node](#)
- [iq](#)
- [jid](#)
- [mnesia](#)
- [muc](#)
- [node](#)
- [omemo](#)
- [otp](#)
- [presence](#)
- [pubsub](#)
- [roster](#)
- [s2s](#)
- [stanza](#)
- [update](#)
- [upgrade](#)
- [vCard](#)
- [vhost](#)
- [xmpp](#)



Bidirectional-streams Over Synchronous HTTP: an extension for [XMPP](#) defined in [XEP-0124](#) that allows long lived connections for XMPP over HTTP. See [mod_bosh](#).

 Command

May refer to:

- [ad-hoc command](#)
- [API command](#)
- shell command

 E

ejabberd Business Edition. See [Options to use ejabberd](#)

 S

ejabberd Community Server, see [Options to use ejabberd](#)

 Erlang

[Erlang](#) is a programming language used to build massively scalable soft real-time systems with requirements on high availability. Some of its uses are in telecoms, banking, e-commerce, computer telephony and instant messaging. Erlang's runtime system has built-in support for concurrency, distribution and fault tolerance.

 Stanza

[Stanza](#) that provides a structured request-response mechanism.

 Jabber ID

[Jabber ID](#) is a unique identifier for an entity in the XMPP network (like a server, a service, an account, a chat room, ...). It consists of three parts: a local part, a domain part, and an optional resource part, formatted as `localpart@domain/resource`.

 Mnesia

[Mnesia](#) is a distributed database specifically designed for industrial-grade telecommunications applications written in Erlang. It is used by default in ejabberd as it requires no initial configuration. However, for large volumes of data, it is recommended to [setup a SQL database](#).

 Multi-User Chat

[Multi-User Chat](#), an extension for [XMPP](#) defined in [XEP-0045](#) for multi-party information exchange whereby multiple users can exchange messages in the context of a chat room. See [mod_muc](#).

**node**

May refer to:

- disco node, see [mod_disco](#)
- [erlang node](#)
- pubsub node, see [mod_pubsub](#)

**omemo**

OMEMO is an extension for [XMPP](#) defined in [XEP-0384](#) that allows end-to-end encryption between clients.

**otp**

[Open Telecom Platform](#) is a set of Erlang libraries and design principles providing middle-ware to develop these systems. It includes its own distributed database, applications to interface towards other languages, debugging and release handling tools.

**presence**

[Stanza](#) that provides information about network availability of an account (online, away, offline, ...).

**pubsub**

[Publish-Subscribe](#), an extension for [XMPP](#) defined in [XEP-0060](#) for for generic publish-subscribe functionality. See [mod_pubsub](#).

**roster**

A list of contacts of an account, which is stored on the server, used to manage [presence](#).

**stanza**

[XMPP](#) packet of type `message`, `presence` or [IQ](#).

**vcard**

Virtual Card, an extension for [XMPP](#) defined in [XEP-0054](#) that allows to publish contact information like name, email, phone number, address, etc. See [mod_vcard](#).

**host**

Virtual Host, a domain name served by ejabberd that corresponds to a [XMPP Domain](#).

**xmpp**

[Extensible Messaging and Presence Protocol](#): an open protocol based in XML for real-time communications.

Readme



 v26.09

 v26.7.0

 v26.07

 v26.07

 v26.09

 CI

 passing

 coverage 36%

 translated 62%

 docs v26.07

ejabberd is an open-source, robust, scalable and extensible realtime platform built using [Erlang/OTP](#), that includes [XMPP](#) Server, [MQTT](#) Broker and [SIP](#) Service.

Check the features in [ejabberd.im](#), [ejabberd Docs](#), [ejabberd at ProcessOne](#), and the list of [supported protocols in ProcessOne](#) and [XMPP.org](#).

Installation

There are several ways to install ejabberd:

- Source code: compile yourself, see [COMPILE](#)
- Installers:
 - [ProcessOne Download Page](#) or [GitHub Releases](#) for releases.
 - [GitHub Actions](#) for master branch (`run / deb / rpm` for `x64` and `arm64`)
- Docker Containers:
 - `ecs` container image: [Docker Hub](#) and [Github Packages](#), see [ecs README](#) (for `x64`)
 - `ejabberd` container image: [Github Packages](#) for releases and master branch, see [CONTAINER](#) (for `x64` and `arm64`)
- Using your [Operating System](#) package
- Using the [Homebrew](#) package manager

More info can be found in the `Installation` part of [ejabberd Docs](#).

Documentation

Please check the [ejabberd Docs](#) website.

When compiling from source code, you can get some help with:

```
./configure --help
make help
```

Once ejabberd is installed, try:

```
ejabberdctl help
man ejabberd.yml
```

Development

Bug reports and features are tracked using [GitHub Issues](#), please check [CONTRIBUTING](#) for details.

Translations can be improved online using [Weblate](#) or in your local machine as explained in [Localization](#).

Documentation for developers is available in [ejabberd docs: Developers](#).

There are nightly builds of ejabberd, both for `master` branch and for Pull Requests:

- Installers: go to [GitHub Actions: Installers](#), open the most recent commit, on the bottom of that commit page, download the `ejabberd-packages.zip` artifact.
- ejabberd container image: go to [ejabberd Github Packages](#)

Security reports or concerns should preferably be reported privately, please send an email to the address: contact at process-one dot net or some other method from [ProcessOne Contact](#).

For commercial offering and support, including [ejabberd Business Edition](#) and [Fluux \(ejabberd in the Cloud\)](#), please check [ProcessOne ejabberd page](#).

Security

For information on how to report security vulnerabilities, please refer to the [SECURITY.md](#) file. It contains guidelines on how to report vulnerabilities privately and securely, ensuring that any issues are addressed in a timely and confidential manner.

Community

There are several places to get in touch with other ejabberd developers and administrators:

- ejabberd XMPP chatroom: ejabberd@conference.process-one.net
- [GitHub Discussions](#)
- [Stack Overflow](#)

License

- ejabberd is released under the **GNU General Public License v2** (see [COPYING](#))
- [ejabberd translations](#) under **MIT License**.

Install

Installation

There are several ways to install ejabberd Community Server, depending on your needs and your infrastructure.

Self-hosted

Container Images

- [ejabberd and ecs Container Images](#) - Ideal for Windows, macOS, Linux, ...

Binary Installers

- [Linux RUN Installer](#) - Suitable for various Linux distributions
- [Linux DEB and RPM Installers](#) - Specifically for DEB and RPM based Linux

Linux and *BSD

- [Operating System Package](#) - Tailored for System Operators

MacOS

- [Homebrew](#) - Optimized for MacOS

Source Code

- [Source Code](#) - Geared towards developers and advanced administrators

On-Premise (eBE)

- [ejabberd Business Edition](#) - Explore professional support and managed services on your infrastructure

Cloud Hosting (Fluux)

- [Fluux.io](#) - Opt for ejabberd hosting with a user-friendly web dashboard



ejabberd Container Images

[ejabberd](#) is an open-source, robust, scalable and extensible realtime platform built using [Erlang/OTP](#), that includes [XMPP](#) Server, [MQTT](#) Broker and [SIP](#) Service.

This page documents those container images ([images comparison](#)):

- 

published in [ghcr.io/processone/ejabberd](#), built using [ejabberd](#) repository, both for stable ejabberd releases and the `master` branch, in x64 and arm64 architectures.
- 

published in [docker.io/ejabberd/ecs](#), built using [docker-ejabberd/ecs](#) repository for ejabberd stable releases in x64 architectures.

For Microsoft Windows, see [Docker Desktop for Windows 10](#), and [Docker Toolbox for Windows 7](#).

For Kubernetes Helm, see [help-ejabberd](#).

Start ejabberd

daemon

Start ejabberd in a new container:

```
docker run --name ejabberd -d -p 5222:5222 ghcr.io/processone/ejabberd
```

That runs the container as a daemon, using ejabberd default configuration file and XMPP domain `localhost`.

Restart the stopped ejabberd container:

```
docker restart ejabberd
```

Stop the running container:

```
docker stop ejabberd
```

Remove the ejabberd container:

```
docker rm ejabberd
```

with Erlang console

Start ejabberd with an interactive Erlang console attached using the `live` command:

```
docker run --name ejabberd -it -p 5222:5222 ghcr.io/processone/ejabberd live
```

That uses the default configuration file and XMPP domain `localhost`.

with your data

Pass a configuration file as a volume and share the local directory to store database:

```
mkdir conf && cp ejabberd.yml.example conf/ejabberd.yml
mkdir database && chown ejabberd database
```

```
docker run --name ejabberd -it \
-v $(pwd)/conf/ejabberd.yml:/opt/ejabberd/conf/ejabberd.yml \
-v $(pwd)/database:/opt/ejabberd/database \
-p 5222:5222 ghcr.io/processone/ejabberd live
```

Notice that ejabberd runs in the container with an account named `ejabberd` with UID 9000 and group `ejabberd` with GID 9000, and the volumes you mount must grant proper rights to that account.

Next steps

Register admin account



If you set the `REGISTER_ADMIN_PASSWORD` environment variable, an account is automatically registered with that password, and admin privileges are granted to it. The account created depends on what variables you have set:

- `EJABBERD_MACRO_ADMIN=juliet@example.org -> juliet@example.org`
- `EJABBERD_MACRO_HOST=example.org -> admin@example.org`
- None of those variables are set -> `admin@localhost`

The account registration is shown in the container log:

```
$ podman run -it \
--env EJABBERD_MACRO_HOST=example.org \
--env EJABBERD_MACRO_ADMIN=juliet@example.org \
--env REGISTER_ADMIN_PASSWORD=somePassw0rd \
ghcr.io/processone/ejabberd

:> ejabberdctl register juliet example.org somePassw0rd
User juliet@example.org successfully registered
```

This is implemented internally by using [Commands on start](#).

Alternatively, you can register the account manually yourself and edit `conf/ejabberd.yml` and add the ACL as explained in [ejabberd Docs: Administration Account](#).



The default ejabberd configuration has already granted admin privilege to an account that would be called `admin@localhost`, so you just need to register it, for example:

```
docker exec -it ejabberd ejabberdctl register admin localhost passw0rd
```

Check ejabberd log

Check the content of the log files inside the container, even if you do not put it on a shared persistent drive:

```
docker exec -it ejabberd tail -f logs/ejabberd.log
```

Inspect container files

The container uses Alpine Linux. Start a shell inside the container:

```
docker exec -it ejabberd sh
```

Open debug console

Open an interactive debug Erlang console attached to a running ejabberd in a running container:

```
docker exec -it ejabberd ejabberdctl debug
```

CAPTCHA

ejabberd includes two example CAPTCHA scripts. If you want to use any of them, first install some additional required libraries:

```
docker exec --user root ejabberd apk add imagemagick ghostscript-fonts bash
```

Now update your ejabberd configuration file, for example:

```
docker exec -it ejabberd vi conf/ejabberd.yml
```

and add this option:

```
captcha_cmd: "$HOME/bin/captcha.sh"
```

Finally, reload the configuration file or restart the container:

```
docker exec ejabberd ejabberdctl reload_config
```

If the CAPTCHA image is not visible, there may be a problem generating it (the ejabberd log file may show some error message); or the image URL may not be correctly detected by ejabberd, in that case you can set the correct URL manually, for example:

```
captcha_url: https://localhost:5443/captcha
```

For more details about CAPTCHA options, please check the [CAPTCHA](#) documentation section.

Advanced

Ports

The container image exposes several ports (check also [Docs: Firewall Settings](#)):

- 5222 : The default port for XMPP clients.
- 5269 : For XMPP federation. Only needed if you want to communicate with users on other servers.
- 5280 : For admin interface (URL is `admin/`).
- 5443 : With encryption, used for admin interface, API, CAPTCHA, OAuth, Websockets and XMPP BOSH.
- 1880 : For admin interface (URL is `/`, useful for [podman-desktop](#) and [docker-desktop](#)) ☀️
- 1883 : Used for MQTT
- 5478 UDP: STUN service 🟡
- 50000-50099 UDP: TURN service 🟡
- 7777 : SOCKS5 file transfer proxy 🟡
- 5210 : Erlang connectivity when `ERL_DIST_PORT` is set, alternative to EPMD ☀️
- 4369-4399 : EPMD and Erlang connectivity, used for `ejabberdctl` and clustering

Volumes

ejabberd produces two types of data: log files and database spool files (Mnesia). This is the kind of data you probably want to store on a persistent or local drive (at least the database).

The volumes you may want to map:

- `/opt/ejabberd/conf/` : Directory containing configuration and certificates
- `/opt/ejabberd/database/` : Directory containing Mnesia database. You should back up or export the content of the directory to persistent storage (host storage, local storage, any storage plugin)
- `/opt/ejabberd/logs/` : Directory containing log files
- `/opt/ejabberd/upload/` : Directory containing uploaded files. This should also be backed up.

All these files are owned by an account named `ejabberd` with group `ejabberd` in the container. Its corresponding `UID:GID` is `9000:9000`. If you prefer bind mounts instead of volumes, then you need to map this to valid `UID:GID` on your host to get read/write access on mounted directories.

If using Docker, try:

```
mkdir database
sudo chown 9000:9000 database
```

If using Podman, try:

```
mkdir database
podman unshare chown 9000:9000 database
```

It's possible to install additional ejabberd modules using volumes, check [this Docs tutorial](#).

Commands on start

The `ejabberdctl` script reads the `CTL_ON_CREATE` environment variable the first time the container is started, and reads `CTL_ON_START` every time the container is started. Those variables can contain one ejabberdctl command, or several commands separated with the blankspace and `;` characters.

If any of those commands returns a failure, the container starting gets aborted. If there is a command with a result that can be ignored, prefix that command with `!`.

All this works when starting ejabberd with the default method `foreground`, not when using `live`, `ixlive`, ...

This example registers an `admin@localhost` account when the container is first created. Everytime the container starts, it shows the list of registered accounts, checks that the admin account exists and password is valid, changes the password of an account if it exists (ignoring any failure), and shows the ejabberd starts (check also the [full example](#)):

```
environment:
- CTL_ON_CREATE=register admin localhost asd
- CTL_ON_START=stats registeredusers ; \
                check_password admin localhost asd ; \
                ! change_password bot123 localhost qqq ; \
                status
```

Same example using Podman:

```
$ podman run -it \
--env CTL_ON_CREATE="register admin localhost asd" \
--env CTL_ON_START="stats registeredusers ; \
                    check_password admin localhost asd ; \
                    ! change_password bot123 localhost qqq ; \
                    status" \
ghcr.io/processone/ejabberd

...

:> ejabberdctl register admin localhost asd
User admin@localhost successfully registered

:> ejabberdctl stats registeredusers
1

:> ejabberdctl check_password admin localhost asd

:> ejabberdctl change_password bot123 localhost qqq
{not_found,"unknown_user"}
:> FAILURE in command 'change_password bot123 localhost qqq' !!! Ignoring result

:> ejabberdctl status
The node ejabberd@localhost is started. Status: started
ejabberd 25.10.0 is running in that node
```

Macros in environment

ejabberd reads `EJABBERD_MACRO_*` environment variables and uses them to define the corresponding [macros](#), overwriting the corresponding macro definition if it was set in the configuration file. This is supported since ejabberd 24.12.

For example, if you configure this in `ejabberd.yml`:

```
acl:
  admin:
    user: ADMIN
```

now you can define the admin account JID using an environment variable:

```
environment:
  - EJABBERD_MACRO_ADMIN=admin@localhost
```

Check the [full example](#) for other example.

ejabberd-contrib

This section addresses those topics related to [ejabberd-contrib](#):

- [Download source code](#)
- [Install a module](#)
- [Install git for dependencies](#)
- [Install your module](#)

DOWNLOAD SOURCE CODE

The `ejabberd` container image includes the `ejabberd-contrib` git repository source code, but `ecs` does not, so first download it:

```
docker exec ejabberd ejabberdctl modules_update_specs
```

INSTALL A MODULE

Compile and install any of the contributed modules, for example:

```
docker exec ejabberd ejabberdctl module_install mod_statsdx

Module mod_statsdx has been installed and started.
It's configured in the file:
  /opt/ejabberd/.ejabberd-modules/mod_statsdx/conf/mod_statsdx.yml
Configure the module in that file, or remove it
and configure in your main ejabberd.yml
```

INSTALL GIT FOR DEPENDENCIES

Some modules depend on erlang libraries, but the container images do not include `git` or `mix` to download them. Consequently, when you attempt to install such a module, there will be error messages like:

```
docker exec ejabberd ejabberdctl module_install ejabberd_observer_cli

I'll download "recon" using git because I can't use Mix to fetch from hex.pm:
/bin/sh: mix: not found
Fetching dependency observer_cli:
/bin/sh: git: not found
...
```

the solution is to install `git` in the container image:

```
docker exec --user root ejabberd apk add git

fetch https://dl-cdn.alpinelinux.org/alpine/v3.21/main/x86_64/APKINDEX.tar.gz
fetch https://dl-cdn.alpinelinux.org/alpine/v3.21/community/x86_64/APKINDEX.tar.gz
(1/3) Installing pcre2 (10.43-r0)
(2/3) Installing git (2.47.2-r0)
(3/3) Installing git-init-template (2.47.2-r0)
Executing busybox-1.37.0-r12.trigger
OK: 27 MiB in 42 packages
```

and now you can upgrade the module:

```
docker exec ejabberd ejabberdctl module_upgrade ejabberd_observer_cli

I'll download "recon" using git because I can't use Mix to fetch from hex.pm:
```

```

/bin/sh: mix: not found
Fetching dependency observer_cli: Cloning into 'observer_cli'...
Fetching dependency os_stats: Cloning into 'os_stats'...
Fetching dependency recon: Cloning into 'recon'...
Inlining: inline_size=24 inline_effort=150
Old inliner: threshold=0 functions=[{insert,2},{merge,2}]
Module ejabberd_observer_cli has been installed.
Now you can configure it in your ejabberd.yml
I'll download "recon" using git because I can't use Mix to fetch from hex.pm:
/bin/sh: mix: not found

```

INSTALL YOUR MODULE

If you [developed an ejabberd module](#), you can install it in your container image:

1. Create a local directory for `ejabberd-modules` :

```
mkdir docker-modules
```

2. Then create the directory structure for your custom module:

```

cd docker-modules

mkdir -p sources/mod_hello_world/
touch sources/mod_hello_world/mod_hello_world.spec

mkdir sources/mod_hello_world/src/
mv mod_hello_world.erl sources/mod_hello_world/src/

mkdir sources/mod_hello_world/conf/
echo -e "modules:\n  mod_hello_world: {}" > sources/mod_hello_world/conf/mod_hello_world.yml

cd ..

```

3. Grant ownership of that directory to the UID that ejabberd will use inside the Docker image:

```
sudo chown 9000 -R docker-modules/
```

4. Start ejabberd in the container:

```

sudo docker run \
  --name hellotest \
  -d \
  --volume "$(pwd)/docker-modules:/home/ejabberd/.ejabberd-modules/" \
  -p 5222:5222 \
  -p 5280:5280 \
  ejabberd/ecs

```

5. Check the module is available for installing, and then install it:

```

sudo docker exec -it hellotest ejabberdctl modules_available
mod_hello_world []

sudo docker exec -it hellotest ejabberdctl module_install mod_hello_world

```

6. If the module works correctly, you will see `Hello` in the ejabberd logs when it starts:

```

sudo docker exec -it hellotest grep Hello logs/ejabberd.log
2020-10-06 13:40:13.154335+00:00 [info]
<0.492.0>@mod_hello_world:start/2:15 Hello, ejabberd world!

```

ejabberdapi

When the container is running (and thus ejabberd), you can exec commands inside the container using `ejabberdctl` or any other of the available interfaces, see [Understanding ejabberd "commands"](#)

Additionally, the container image includes the `ejabberdapi` executable. Please check the [ejabberd-api homepage](#) for configuration and usage details.

For example, if you configure ejabberd like this:

```

listen:
-
  port: 5282
  module: ejabberd_http
  request_handlers:

```

```

    "/api": mod_http_api

acl:
  loopback:
    ip:
      - 127.0.0.0/8
      - ::1/128
      - ::FFFF:127.0.0.1/128

api_permissions:
  "admin access":
    who:
      access:
        allow:
          acl: loopback
    what:
      - "register"

```

Then you could register new accounts with this query:

```
docker exec -it ejabberd ejabberdapi register --endpoint=http://127.0.0.1:5282/ --jid=admin@localhost --password=passw0rd
```

Clustering

When setting several containers to form a [cluster of ejabberd nodes](#), each one must have a different [Erlang Node Name](#) and the same [Erlang Cookie](#).

For this you can either:

- edit `conf/ejabberdctl.cfg` and set variables `ERLANG_NODE` and `ERLANG_COOKIE`
- set the environment variables `ERLANG_NODE_ARG` and `ERLANG_COOKIE`

Example to connect a local `ejabberdctl` to a containerized ejabberd:

1. When creating the container, export port 5210, and set `ERLANG_COOKIE` :

```
docker run --name ejabberd -it \
-e ERLANG_COOKIE=`cat $HOME/.erlang.cookie` \
-p 5210:5210 -p 5222:5222 \
ghcr.io/processone/ejabberd
```

2. Set `ERL_DIST_PORT=5210` in `ejabberdctl.cfg` of container and local ejabberd
3. Restart the container
4. Now use `ejabberdctl` in your local ejabberd deployment

To connect using a local `ejabberd` script:

```
ERL_DIST_PORT=5210 _build/dev/rel/ejabberd/bin/ejabberd ping
```

Example using environment variables (see full example [docker-compose.yml](#)):

```

environment:
  - ERLANG_NODE_ARG=ejabberd@node7
  - ERLANG_COOKIE=dummycookie123

```

Once you have the ejabberd nodes properly set and running, you can tell the secondary nodes to join the master node using the [join_cluster](#) API call.

Example using environment variables (see the full [docker-compose.yml clustering example](#)):

```

environment:
  - ERLANG_NODE_ARG=ejabberd@replica
  - ERLANG_COOKIE=dummycookie123
  - CTL_ON_CREATE=join_cluster ejabberd@main

```

Change Mnesia Node Name

To use the same Mnesia database in a container with a different hostname, it is necessary to change the old hostname stored in Mnesia.

This section is equivalent to the ejabberd Documentation [Change Computer Hostname](#), but particularized to containers that use this `ecs` container image from ejabberd 23.01 or older.

SETUP OLD CONTAINER

Let's assume a container running ejabberd 23.01 (or older) from this `ecs` container image, with the database directory binded and one registered account. This can be produced with:

```
OLDCONTAINER=ejaold
NEWCONTAINER=ejanew

mkdir database
sudo chown 9000:9000 database
docker run -d --name $OLDCONTAINER -p 5222:5222 \
  -v $(pwd)/database:/opt/ejabberd/database \
  ghcr.io/processone/ejabberd:23.01
docker exec -it $OLDCONTAINER ejabberdctl started
docker exec -it $OLDCONTAINER ejabberdctl register user1 localhost somepass
docker exec -it $OLDCONTAINER ejabberdctl registered_users localhost
```

Methods to know the Erlang node name:

```
ls database/ | grep ejabberd@
docker exec -it $OLDCONTAINER ejabberdctl status
docker exec -it $OLDCONTAINER grep "started in the node" logs/ejabberd.log
```

CHANGE MNESIA NODE

First of all let's store the Erlang node names and paths in variables. In this example they would be:

```
OLDCONTAINER=ejaold
NEWCONTAINER=ejanew
OLDNODE=ejabberd@95145dde27c
NEWNODE=ejabberd@localhost
OLDFILE=/opt/ejabberd/database/old.backup
NEWFILE=/opt/ejabberd/database/new.backup
```

1. Start your old container that can still read the Mnesia database correctly. If you have the Mnesia spool files, but don't have access to the old container anymore, go to [Create Temporary Container](#) and later come back here.
2. Generate a backup file and check it was created:

```
docker exec -it $OLDCONTAINER ejabberdctl backup $OLDFILE
ls -l database/*.backup
```

3. Stop ejabberd:

```
docker stop $OLDCONTAINER
```

4. Create the new container. For example:

```
docker run \
  --name $NEWCONTAINER \
  -d \
  -p 5222:5222 \
  -v $(pwd)/database:/opt/ejabberd/database \
  ghcr.io/processone/ejabberd:latest
```

5. Convert the backup file to new node name:

```
docker exec -it $NEWCONTAINER ejabberdctl mnesia_change_nodename $OLDNODE $NEWNODE $OLDFILE $NEWFILE
```

6. Install the backup file as a fallback:

```
docker exec -it $NEWCONTAINER ejabberdctl install_fallback $NEWFILE
```

7. Restart the container:

```
docker restart $NEWCONTAINER
```

8. Check that the information of the old database is available.

In this example, it should show that the account `user1` is registered:

```
docker exec -it $NEWCONTAINER ejabberdctl registered_users localhost
```

9. When the new container is working perfectly with the converted Mnesia database, you may want to remove the unneeded files: the old container, the old Mnesia spool files, and the backup files.

CREATE TEMPORARY CONTAINER

In case the old container that used the Mnesia database is not available anymore, a temporary container can be created just to read the Mnesia database and make a backup of it, as explained in the previous section.

This method uses `--hostname` command line argument for docker, and `ERLANG_NODE_ARG` environment variable for ejabberd. Their values must be the hostname of your old container and the Erlang node name of your old ejabberd node. To know the Erlang node name please check [Setup Old Container](#).

Command line example:

```
OLDHOST=${OLDNODE#*@}
docker run \
  -d \
  --name $OLDCONTAINER \
  --hostname $OLDHOST \
  -p 5222:5222 \
  -v $(pwd)/database:/opt/ejabberd/database \
  -e ERLANG_NODE_ARG=$OLDNODE \
  ghcr.io/processone/ejabberd:latest
```

Check the old database content is available:

```
docker exec -it $OLDCONTAINER ejabberdctl registered_users localhost
```

Now that you have ejabberd running with access to the Mnesia database, you can continue with step 2 of previous section [Change Mnesia Node](#).

Build Container Image

The container image includes ejabberd as a standalone OTP release built using Elixir.

Build ejabberd

The ejabberd Erlang/OTP release is configured with:

- `mix.exs`: Customize ejabberd release
- `vars.config`: ejabberd compilation configuration options
- `config/runtime.exs`: Customize ejabberd paths
- `ejabberd.yml.template`: ejabberd default config file

DIRECT BUILD

Build ejabberd Community Server container image from ejabberd master git repository:

```
docker buildx build \
  -t personal/ejabberd \
  -f .github/container/Dockerfile \
  .
```

PODMAN BUILD

Some minor remarks:

- When building, it mentions that `healthcheck` is not supported by the Open Container Initiative image format
- To start with command `live`, you may want to add environment variable `EJABBERD_BYPASS_WARNINGS=true`

```
podman build \
  -t ejabberd \
  -f .github/container/Dockerfile \
  .

podman run --name ejal -d -p 5222:5222 localhost/ejabberd

podman exec ejal ejabberdctl status

podman exec -it ejal sh

podman stop ejal

podman run --name ejal -it -e EJABBERD_BYPASS_WARNINGS=true -p 5222:5222 localhost/ejabberd live
```

Build ecs

The ejabberd Erlang/OTP release is configured with:

- `rel/config.exs`: Customize ejabberd release
- `rel/dev.exs`: ejabberd environment configuration for development release
- `rel/prod.exs`: ejabberd environment configuration for production release
- `vars.config`: ejabberd compilation configuration options
- `conf/ejabberd.yml`: ejabberd default config file

Build ejabberd Community Server base image from ejabberd master on Github:

```
docker build -t personal/ejabberd .
```

Build ejabberd Community Server base image for a given ejabberd version:

```
./build.sh 18.03
```

Composer Examples

Minimal Example

This is the barely minimal file to get a usable ejabberd.

If using Docker, write this `docker-compose.yml` file and start it with `docker-compose up`:

```
services:
  main:
    image: ghcr.io/processone/ejabberd
    container_name: ejabberd
    ports:
      - "5222:5222"
      - "5269:5269"
      - "5280:5280"
      - "5443:5443"
```

If using Podman, write this `minimal.yml` file and start it with `podman kube play minimal.yml`:

```
apiVersion: v1

kind: Pod

metadata:
  name: ejabberd

spec:
  containers:
  - name: ejabberd
    image: ghcr.io/processone/ejabberd
    ports:
      - containerPort: 5222
        hostPort: 5222
      - containerPort: 5269
        hostPort: 5269
      - containerPort: 5280
        hostPort: 5280
      - containerPort: 5443
        hostPort: 5443
```

Customized Example

This example shows the usage of several customizations: it uses a local configuration file, defines a configuration macro using an environment variable, stores the mnesia database in a local path, registers an account when it's created, and checks the number of registered accounts every time it's started.

Prepare an ejabberd configuration file:

```
mkdir conf && cp ejabberd.yml.example conf/ejabberd.yml
```

Create the database directory and allow the container access to it:

- Docker:

```
mkdir database && sudo chown 9000:9000 database
```

- Podman:

```
mkdir database && podman unshare chown 9000:9000 database
```

If using Docker, write this `docker-compose.yml` file and start it with `docker-compose up`:

```
version: '3.7'

services:
  main:
    image: ghcr.io/processone/ejabberd
    container_name: ejabberd
    environment:
      - EJABBERD_MACRO_HOST=example.com
      - EJABBERD_MACRO_ADMIN=admin@example.com
      - REGISTER_ADMIN_PASSWORD=somePassword
      - CTL_ON_START=registered_users example.com ;
```

```

      status
ports:
  - "5222:5222"
  - "5269:5269"
  - "5280:5280"
  - "5443:5443"
volumes:
  - ./conf/ejabberd.yml:/opt/ejabberd/conf/ejabberd.yml:ro
  - ./database:/opt/ejabberd/database

```

If using Podman, write this `custom.yml` file and start it with `podman kube play custom.yml`:

```

apiVersion: v1

kind: Pod

metadata:
  name: ejabberd

spec:
  containers:

  - name: ejabberd
    image: ghcr.io/processone/ejabberd
    env:
      - name: EJABBERD_MACRO_HOST
        value: example.com
      - name: EJABBERD_MACRO_ADMIN
        value: admin@example.com
      - name: REGISTER_ADMIN_PASSWORD
        value: somePassw0rd
      - name: CTL_ON_START
        value: registered_users example.com ;
          status
    ports:
      - containerPort: 5222
        hostPort: 5222
      - containerPort: 5269
        hostPort: 5269
      - containerPort: 5280
        hostPort: 5280
      - containerPort: 5443
        hostPort: 5443
    volumeMounts:
      - mountPath: /opt/ejabberd/conf/ejabberd.yml
        name: config
        readOnly: true
      - mountPath: /opt/ejabberd/database
        name: db

  volumes:
  - name: config
    hostPath:
      path: ./conf/ejabberd.yml
      type: File
  - name: db
    hostPath:
      path: ./database
      type: DirectoryOrCreate

```

Clustering Example

In this example, the main container is created first. Once it is fully started and healthy, a second container is created, and once ejabberd is started in it, it joins the first one.

An account is registered in the first node when created (and we ignore errors that can happen when doing that - for example when account already exists), and it should exist in the second node after join.

Notice that in this example the main container does not have access to the exterior; the replica exports the ports and can be accessed.

If using Docker, write this `docker-compose.yml` file and start it with `docker-compose up`:

```

version: '3.7'

services:

  main:
    image: ghcr.io/processone/ejabberd
    container_name: main
    environment:
      - ERLANG_NODE_ARG=ejabberd@main
      - ERLANG_COOKIE=dummycookie123
      - CTL_ON_CREATE=! register admin localhost asd

```

```

healthcheck:
  test: netstat -nl | grep -q 5222
  start_period: 5s
  interval: 5s
  timeout: 5s
  retries: 120

replica:
  image: ghcr.io/processone/ejabberd
  container_name: replica
  depends_on:
    main:
      condition: service_healthy
  environment:
    - ERLANG_NODE_ARG=ejabberd@replica
    - ERLANG_COOKIE=dummycookie123
    - CTL_ON_CREATE=join_cluster ejabberd@main
    - CTL_ON_START=registered_users localhost ;
      status
  ports:
    - "5222:5222"
    - "5269:5269"
    - "5280:5280"
    - "5443:5443"

```

If using Podman, write this `cluster.yml` file and start it with `podman kube play cluster.yml`:

```

apiVersion: v1

kind: Pod

metadata:
  name: cluster

spec:
  containers:

  - name: first
    image: ghcr.io/processone/ejabberd
    env:
      - name: ERLANG_NODE_ARG
        value: main@cluster
      - name: ERLANG_COOKIE
        value: dummycookie123
      - name: CTL_ON_CREATE
        value: register admin localhost asd
      - name: CTL_ON_START
        value: stats registeredusers ;
          status
      - name: EJABBERD_MACRO_PORT_C2S
        value: 6222
      - name: EJABBERD_MACRO_PORT_C2S_TLS
        value: 6223
      - name: EJABBERD_MACRO_PORT_S2S
        value: 6269
      - name: EJABBERD_MACRO_PORT_HTTP_TLS
        value: 6443
      - name: EJABBERD_MACRO_PORT_HTTP
        value: 6280
      - name: EJABBERD_MACRO_PORT_MQTT
        value: 6883
      - name: EJABBERD_MACRO_PORT_PROXY65
        value: 6777
    volumeMounts:
      - mountPath: /opt/ejabberd/conf/ejabberd.yml
        name: config
        readOnly: true

  - name: second
    image: ghcr.io/processone/ejabberd
    env:
      - name: ERLANG_NODE_ARG
        value: replica@cluster
      - name: ERLANG_COOKIE
        value: dummycookie123
      - name: CTL_ON_CREATE
        value: join_cluster main@cluster ;
          started ;
          list_cluster
      - name: CTL_ON_START
        value: stats registeredusers ;
          check_password admin localhost asd ;
          status
    ports:
      - containerPort: 5222
        hostPort: 5222
      - containerPort: 5280
        hostPort: 5280
    volumeMounts:
      - mountPath: /opt/ejabberd/conf/ejabberd.yml
        name: config
        readOnly: true

```

```
volumes:  
- name: config  
  hostPath:  
    path: ./conf/ejabberd.yml  
    type: File
```

Images Comparison

Let's summarize the differences between both container images. Legend:

- 🌿: is the recommended alternative
- 🟡: changed in ejabberd 26.01
- 🌞: changed in ...
- ⚙️: changed in ejabberd 25.03

	 ejabberd	 ecs
Source code	ejabberd/.github/container	docker-ejabberd/ecs
Generated by	container.yml	tests.yml
Built for	stable releases master branch	stable releases master branch zip
Architectures	linux/amd64 linux/arm64	linux/amd64
Software	Erlang/OTP 28.5.0.7-alpine  Elixir 1.19.6 	Alpine 3.22 Erlang/OTP 27.3.4.18 Elixir 1.19.6
Published in	ghcr.io/processone/ejabberd	docker.io/ejabberd/ecs ghcr.io/processone/ecs
 Additional content		
ejabberd-contrib	included	not included
ejabberdapi	included 	included
 Ports		
1880 for WebAdmin	yes 	yes 
5210 for ERL_DIST_PORT	supported	supported 
 Paths		
\$HOME	/opt/ejabberd/	/home/ejabberd/
User data	\$HOME  /home/ejabberd/ 	\$HOME /opt/ejabberd/ 
ejabberdctl	ejabberdctl  bin/ejabberdctl 	bin/ejabberdctl ejabberdctl 
captcha.sh	\$HOME/bin/captcha.sh 	\$HOME/bin/captcha.sh 
.sql files	\$HOME/sql/.sql   \$HOME/database/*.sql 	\$HOME/database/*.sql \$HOME/sql/*.sql  
Mnesia spool files	\$HOME/database/  \$HOME/database/NODENAME/ 	\$HOME/database/NODENAME/ \$HOME/database/  
 Variables		
EJABBERD_MACRO_*	supported	supported
Macros used in ejabberd.yml	yes 	yes 
EJABBERD_MACRO_ADMIN	Grant admin rights  (default admin@localhost)	Hardcoded admin@localhost
REGISTER_ADMIN_PASSWORD	Register admin account 	unsupported
CTL_OVER_HTTP	enabled 	unsupported

Binary Installers

Linux RUN Installer

The `*.run` binary installer will deploy and configure a full featured ejabberd server and does not require any extra dependencies. It includes a stripped down version of Erlang. As such, when using ejabberd installer, you do not need to install Erlang separately.

Those instructions assume installation on `localhost` for development purposes. In this document, when mentioning `ejabberd-YY.MM`, we assume `YY.MM` is the release number, for example 18.01.

Installation using the `*.run` binary installer:

1. Go to [ejabberd GitHub Releases](#).
2. Download the `run` package for your architecture
3. Right-click on the downloaded file and select "Properties", click on the "Permissions" tab and tick the box that says "Allow executing file as program". Alternatively, you can set the installer as executable using the command line:

```
chmod +x ejabberd-YY.MM-1-linux-x64.run
```

4. If the installer runs as superuser (by `root` or using `sudo`), it installs ejabberd binaries in `/opt/ejabberd-XX.YY/`; installs your configuration, Mnesia database and logs in `/opt/ejabberd/`, and setups an ejabberd service unit in `systemd`:

```
sudo ./ejabberd-YY.MM-1-linux-x64.run
```

5. If the installer runs as a regular user, it asks the base path where ejabberd should be installed. In that case, the ejabberd service unit is not set in `systemd`, and `systemctl` cannot be used to start ejabberd; start it manually.
6. After successful installation by `root`, ejabberd is automatically started. Check its status with

```
systemctl status ejabberd
```

7. Now that ejabberd is installed and running with the default configuration, it's time to do some basic setup: edit `/opt/ejabberd/conf/ejabberd.yml` and setup in the `hosts` option the domain that you want ejabberd to serve. By default it's set to the name of your computer on the local network.
8. Restart ejabberd completely using `systemctl`, or using `ejabberdctl`, or simply tell it to reload the configuration file:

```
sudo systemctl restart ejabberd
sudo /opt/ejabberd-22.05/bin/ejabberdctl restart
sudo /opt/ejabberd-22.05/bin/ejabberdctl reload_config
```

9. Quite probably you will want to register an account and grant it admin rights, please check [Next Steps: Administration Account](#).
10. Now you can go to the web dashboard at `http://localhost:5280/admin/` and fill the username field with the full account JID, for example `admin@domain` (or `admin@localhost` as above). Then fill the password field with that account's `password`. The next step is to get to know [how to configure ejabberd](#).
11. If something goes wrong during the installation and you would like to start from scratch, you will find the steps to uninstall in the file `/opt/ejabberd-22.05/uninstall.txt`.

Linux DEB and RPM Installers

ProcessOne provides DEB and RPM all-in-one binary installers with the same content that the `*.run` binary installer mentioned in the previous section.

Those are self-sufficient packages that contain a minimal Erlang distribution, this ensures that it does not interfere with your existing Erlang version and is also a good way to make sure ejabberd will run with the latest Erlang version.

Those packages install ejabberd in `/opt/ejabberd-XX.YY/`. Your configuration, Mnesia database and logs are available in `/opt/ejabberd/`.

You can download directly the DEB and RPM packages from [ejabberd GitHub Releases](#).

If you prefer, you can also get those packages from our official [ejabberd packages repository](#).

Operating System Packages

Many operating systems provide specific ejabberd packages adapted to the system architecture and libraries. They usually also check dependencies and perform basic configuration tasks like creating the initial administrator account.

List of known ejabberd packages:

- [Alpine Linux](#)
- [Arch Linux](#)
- [Debian](#)
- [Fedora](#)
- [FreeBSD](#)
- [Gentoo](#)
- [OpenSUSE](#)
- [NetBSD](#)
- [Ubuntu](#)

Consult the resources provided by your Operating System for more information.

There's also an [ejabberd snap](#) to install ejabberd on several operating systems using `snap` package manager.

Install ejabberd from Source Code

The canonical distribution form of ejabberd stable releases is the source code package. Compiling ejabberd from source code is quite easy in *nix systems, as long as your system have all the dependencies.

Requirements

To compile ejabberd you need:

- GNU Make
- GCC
- Libexpat ≥ 1.95
- Libyaml $\geq 0.1.4$
- [Erlang/OTP](#) ≥ 25.0 . It is recommended Erlang/OTP 27.3, which is the version used in the binary installers and container images.
- OpenSSL $\geq 1.0.0$
- Curl. Optional, for ejabberdctl feature [CTL_OVER_HTTP](#)

Other optional libraries are:

- Zlib $\geq 1.2.3$, For [Zlib Stream Compression](#)
- PAM library, for [PAM Authentication](#)
- ImageMagick's Convert program and Ghostscript fonts, for [CAPTCHA challenges](#).
- [Elixir](#) $\geq 1.10.3$, for [Elixir Development](#). It is recommended Elixir 1.14.0 or higher.

If your system splits packages in libraries and development headers, install the development packages too.

For example, in Debian:

```
apt-get install libexpat1-dev libgd-dev libpam0g-dev \
    libsqlite3-dev libwebp-dev libyaml-dev \
    autoconf automake erlang elixir rebar3
```

Download

There are several ways to obtain the ejabberd source code:

- Source code package from [ProcessOne Downloads](#) or [GitHub Releases](#)
- Latest development code from [ejabberd Git repository](#) using the commands:

```
git clone https://github.com/processone/ejabberd.git
cd ejabberd
```

Compile

The generic instructions to compile ejabberd are:

```
./autogen.sh
./configure
make
```

Let's view them in detail.

`./configure`

The build configuration script supports many options. Get the full list:

```
./configure --help
```

In this example, `./configure` prepares the installed program to run with a user called `ejabberd` that should exist in the system (it isn't recommended to run `ejabberd` with `root` user):

```
./configure --enable-user=ejabberd --enable-mysql
```

If you get `Error loading module rebar3`, please consult how to use [rebar with old Erlang](#).

GENERAL OPTIONS:

- `--bindir=` : Specify the path to the user executables (where `epmd` and `iex` are available).
- `--prefix=` : Specify the path prefix where the files will be copied when running the `make install` command.
- `--with-erlang=` : Specify the directory where [Erlang/OTP is installed](#).
- `--with-min-erlang=9.0.5` : Allow to compile ejabberd with Erlang/OTP 20.0, bypassing the verification of lower supported version. Use only if you know what you are doing.
-  added in [25.07](#)
- `--with-rebar=` : Specify the path to `rebar`, `rebar3` or [mix](#)
-  added in [20.12](#) and improved in [24.02](#)
- `--enable-user[=USER]` : Allow this normal system user to execute the `ejabberdctl` script (see section [ejabberdctl](#)), read the configuration files, read and write in the spool directory, read and write in the log directory. The account user and group must exist in the machine before running `make install`. This account needs a HOME directory, because the [Erlang cookie file](#) will be created and read there.
- `--enable-group[=GROUP]` : Use this option additionally to `--enable-user` when that account is in a group that doesn't coincide with its username.

OPTIONS THAT ENABLE/DISABLE FEATURES:

- `--enable-all` : Enable many of the database and dependencies options described here: `debug`, `elixir`, `lua`, `mssql`, `mysql`, `odbc`, `pam`, `pgsql`, `redis`, `sip`, `sqlite`, `stun`, `tools`, `zlib`. This is specially useful for Dialyzer checks.
- `--disable-bootstrap` : Don't include `bootstrap` dependencies for `mod_invites`.
-  added in [26.04](#)
- `--disable-debug` : Compile without `+debug_info`.
- `--enable-elixir` : Build ejabberd with Elixir extension support. Works only with `rebar3`, not `rebar2`. Requires to have Elixir installed. If interested in Elixir development, you may prefer to use `--with-rebar=mix`
-  improved in [24.02](#)
- `--disable-erlang-version-check` : Don't check Erlang/OTP version.
- `--enable-full-xml` : Use XML features in XMPP stream (ex: CDATA). This requires XML compliant clients).
- `--enable-hipe` : Compile natively with HiPE. This is an experimental feature, and not recommended.
- `--enable-lager` : Use lager Erlang logging tool instead of standard error logger.
- `--enable-latest-deps` : Makes rebar use latest versions of dependencies developed alongside ejabberd instead of version specified in `rebar.config`. Should be only used when developing ejabberd.
- `--enable-lua` : Enable Lua support, to import from Prosody.
-  added in [21.04](#)

- `--enable-mssql` : Enable Microsoft SQL Server support, this option requires `--enable-odbc` (see [Supported storages](#)).
- `--enable-mysql` : Enable MySQL support (see [Supported storages](#)).
- `--enable-odbc` : Enable pure ODBC support.
- `--enable-pam` : Enable the PAM authentication method (see [PAM Authentication](#) section).
- `--enable-pgsql` : Enable PostgreSQL support (see [Supported storages](#)).
- `--enable-redis` : Enable Redis support to use for external session storage.
- `--enable-roster-gateway-workaround` : Turn on workaround for processing gateway subscriptions.
- `--enable-sip` : Enable SIP support.
- `--enable-sql-schema-multihost` : Use multihost SQL schema by default.
- `--enable-sqlite` : Enable SQLite support (see [Supported storages](#)).
- `--disable-stun` : Disable STUN/TURN support.
- `--enable-system-deps` : Makes rebar use locally installed dependencies instead of downloading them.
- `--enable-tools` : Enable the use of development tools.



changed in [21.04](#)

- `--disable-zlib` : Disable Stream Compression (XEP-0138) using zlib.

make

This manages many tasks:

- Download and compile [erlang dependencies](#)
- Compile ejabberd
- [System install](#), uninstall
- Build OTP [production](#) / [development](#) releases
- Development: edoc, [options](#), [translations](#), tags
- Testing: dialyzer, [hooks](#), [test](#), xref

Get the full task list:

```
make help
```

Install

There are several ways to install and run ejabberd after it's compiled from source code:

- [system install](#)
- [system install a release](#)
- building a [production](#) OTP release
- building a [development](#) OTP release
- start directly with [relivectl](#)
- start directly with [relive](#) which uses Rebar3/Mix

Check the [Install Comparison](#) table to find the differences.

System Install

To install ejabberd in the destination directories, run:

```
make install
```

Note that you probably need administrative privileges in the system to install ejabberd.

The created files and directories depend on the options provided to `./configure`, by default they are:

- `/etc/ejabberd/` : Configuration directory:
- `ejabberd.yml` : ejabberd configuration file (see [File Format](#))
- `ejabberdctl.cfg` : Configuration file of the administration script (see [Erlang Runtime System](#))
- `inetrc` : Network DNS configuration file for Erlang
- `/lib/ejabberd/` :
- `ebin/` : Erlang binary files (*.beam)
- `include/` : Erlang header files (*.hrl)
- `priv/` : Additional files required at runtime
- `bin/` : Executable programs
- `lib/` : Binary system libraries (*.so)
- `msgs/` : Translation files (*.msgs) (see [Default Language](#))
- `/sbin/ejabberdctl` : Administration script (see [ejabberdctl](#))
- `/share/doc/ejabberd/` : Documentation of ejabberd
- `/var/lib/ejabberd/` : Spool directory:
- `.erlang.cookie` : The [Erlang cookie file](#)
- `acl.DCD, ...` : Mnesia database spool files (*.DCD, *.DCL, *.DAT)
- `/var/log/ejabberd/` : Log directory (see [Logging](#)):
- `ejabberd.log` : ejabberd service log
- `erlang.log` : Erlang/OTP system log

System Install Release

 added in [24.02](#)

This builds a [production release](#), and then performs a [system install](#) of that release, obtaining a result similar to the one mentioned in the previous section.

Simply run:

```
make install-rel
```

The benefits of `install-rel` over `install`:

- this uses OTP release code from `rebar/rebar3/mix`, and consequently requires less code in our `Makefile.in` file
- `uninstall-rel` correctly deletes all the library files
- the `*.beam` files are smaller as debug information is stripped

Production Release

 improved in [21.07](#)

You can build an OTP release that includes ejabberd, Erlang/OTP and all the required erlang dependencies in a single tar.gz file. Then you can copy that file to another machine that has the same machine architecture, and run ejabberd without installing anything else.

To build that production release, run:

```
make prod
```

If you provided to `./configure` the option `--with-rebar` to use rebar3 or mix, this will directly produce a tar.gz that you can copy.

This example uses rebar3 to manage the compilation, builds an OTP production release, copies the resulting package to a temporary path, and starts ejabberd there:

```
./autogen.sh
./configure --with-rebar=rebar3
make
make prod
mkdir $HOME/eja-release
tar -xvzf _build/prod/ejabberd-*.tar.gz -C $HOME/eja-release
$HOME/eja-release/bin/ejabberdctl live
```

Development Release



new in 21.07

If you provided to `./configure` the option `--with-rebar` to use rebar3 or mix, you can build an OTP development release.

This is designed to run ejabberd in the local machine for development, manual testing... without installing in the system.

This development release has some customizations: uses a dummy certificate file, if you register the account `admin@localhost` it has admin rights...

This example uses Elixir's mix to manage the compilation, builds an OTP development release, and starts ejabberd there:

```
./autogen.sh
./configure --with-rebar=mix
make
make dev
_build/dev/rel/ejabberd/bin/ejabberdctl live
```

Relivectl



new in 26.01

`make relivectl` starts an interactive ejabberd without requiring installation or building OTP release. It uses the `ejabberdctl` script, and stores data in local path `_build/relivectl/`.

Example usage:

```
./autogen.sh
./configure
make relivectl
```

The benefit over `make dev` is that `relivectl` doesn't build an OTP release, so it's faster to start.

Relive



new in 22.05

`make relive` starts an interactive ejabberd without requiring installation or building OTP release. It uses `rebar3 shell` or `mix run`, and stores data in local path `_build/relive/`.

Example usage:

```
./autogen.sh
./configure
make relive
```

As it uses Rebar3/Mix tools, it automatically compiles code at start and [recompiles changed code at runtime](#).

However, as it doesn't use the `ejabberdctl` script, it doesn't read `ejabberdctl.cfg`.

Specific notes

asdf

When Erlang/OTP (and/or Elixir) is installed using [asdf](#) (multiple runtime version manager), it is available only for your account, in `$HOME/.asdf/shims/erl`. In that case, you cannot install ejabberd globally in the system, and you cannot use the `root` account to start it, because that account doesn't have access to erlang.

In that scenario, there are several ways to run/install ejabberd:

- Run a [development release](#) locally without installing
- Copy a [production release](#) locally
- Use [system install](#), but install it locally:

```
./autogen.sh
./configure --prefix=$HOME/eja-install --enable-user
make
make install
$HOME/eja-install/sbin/ejabberdctl live
```

BSD

The command to compile ejabberd in BSD systems is `gmake`.

You may want to check [pkgsrc.se for ejabberd](#).

Up to ejabberd [23.04](#), some old scripts were included in ejabberd source for NetBSD compilation, and you can take a look to those files for reference in ejabberd [23.04/examples/mtr/](#) path.

Erlang Configuration

Usually ejabberd is started using the `ejabberdctl` script, which takes care to setup the [Erlang runtime system options](#) following your `ejabberdctl.cfg` configuration file.

However, there are other methods to start ejabberd, and they use specific methods to configure those options:

- When using `rebar3` `make relive`:
system is set in `rebar.config`, apps in `rel/relive.config`
- When using `mix` `make relive`:
system is set in `Makefile`, apps in `config/runtime.exs`
- When using `rebar3` `make dev && ejabberd`:
system is set in `rel/vm.args`, apps in `rel/sys.config`
- When using `mix` `make dev|prod && ejabberd`:
system is set in `rel/vm.args.eex` and `rel/env.sh.eex`, apps in `config/runtime.exs`

Erlang Local Install

If you compile Erlang/OTP from source code and install in a local folder, for example in `$HOME/erl-local` like this:

```
cd otp_src
./configure --prefix=$HOME/erl-local
make
make install
```

then those are the steps to compile ejabberd:

```
cd ejabberd-master
export PATH=$HOME/erl-local/bin:$PATH
./autogen.sh
./configure --with-erlang=$HOME/erl-local --with-rebar=./rebar3
make
```

Also notice that Erlang and ejabberd can be used only from that system account, see section [asdf](#) for details.

Install Comparison

Let's summarize all the [make](#) targets related to installation to determine their usage differences:

make ...	install	install-rel	prod	dev	relivectl	rebuild
Writes files in path	/	/	_build/ prod/	_build/ dev/	_build/ relivectl/	_build/ rebuild/
Installs	✓	✓	manually uncompress *.tar.gz	-	-	-
Uninstall with	uninstall ⚠ incomplete	uninstall-rel ✓	manual remove	-	-	-
Start tool	ejabberdctl	ejabberdctl	ejabberdctl	ejabberdctl	ejabberdctl	rebuild
Reads ejabberdctl.cfg	✓	✓	✓	✓	✓	-
Recompiles	-	✓	✓	✓	✗	-
Starts ejabberd	-	-	-	-	✓	-
Recompiles at runtime	-	-	-	-	✗	-
Execution time (s.)	13	40	57	35	4	-

macOS

If compiling from sources on Mac OS X, you must configure ejabberd to use custom OpenSSL, Yaml, iconv. The best approach is to use [Homebrew](#) to install your dependencies, then exports your custom path to let configure and make be aware of them.

```
brew install git erlang elixir openssl expat libyaml libiconv libgd sqlite rebar rebar3 automake autoconf
export LDFLAGS="-L/usr/local/opt/openssl/lib -L/usr/local/lib -L/usr/local/opt/expat/lib"
export CFLAGS="-I/usr/local/opt/openssl/include -I/usr/local/include -I/usr/local/opt/expat/include"
export CPPFLAGS="-I/usr/local/opt/openssl/include/ -I/usr/local/include -I/usr/local/opt/expat/include"
./configure
make
```

Check also the guide for [Installing ejabberd development environment on OSX](#)

man

ejabberd includes a man page which documents the toplevel and modules options, the same information that is published in the [Top-Level Options](#) and [Modules Options](#) sections.

The man file can be read locally with:

```
man -l man/ejabberd.yml.5
```

rebar with old Erlang

The ejabberd source code package includes `rebar` and `rebar3` binaries that work with Erlang/OTP 25.0 up to 28.

To compile ejabberd using rebar/rebar3 and Erlang 20.0 up to 23.3, you can install it from your operating system, or compile yourself from the rebar source code, or download the old binary from ejabberd 21.12:

```
wget https://github.com/processone/ejabberd/raw/21.12/rebar
wget https://github.com/processone/ejabberd/raw/21.12/rebar3
```

To compile ejabberd using rebar/rebar3 and Erlang 24.0 up to 24.3, you can install it from your operating system, or compile yourself from the rebar source code, or download the old binary from ejabberd 24.12:

```
wget https://github.com/processone/ejabberd/raw/24.12/rebar
wget https://github.com/processone/ejabberd/raw/24.12/rebar3
```

Start

You can use the `ejabberdctl` command line administration script to start and stop ejabberd. Some examples, depending on your installation method:

- When [installed in the system](#):

```
ejabberdctl start
/sbin/ejabberdctl start
```

- When [built an OTP production release](#):

```
_build/prod/rel/ejabberd/bin/ejabberdctl start
_build/prod/rel/ejabberd/bin/ejabberdctl live
```

- Start interactively without installing or building OTP release:

```
make relive
```

Install ejabberd on macOS

Homebrew

[Homebrew](#) is a package manager for macOS that aims to port the many Unix & Linux software that is not easily available or compatible. Homebrew installation is simple and the instruction is available on its website.

Check also the guide for [Installing ejabberd development environment on OSX](#)

The ejabberd configuration included in Homebrew's ejabberd has as default domain `localhost`, and has already granted administrative privileges to the account `admin@localhost`.

1. Once you have Homebrew installed, open Terminal. Run

```
brew install ejabberd
```

This should install the latest or at most the one-before-latest version of ejabberd. The installation directory should be reported at the end of this process, but usually the main executable is stored at `/usr/local/sbin/ejabberdctl`.

2. Start ejabberd in interactive mode, which prints useful messages in the Terminal.

```
/usr/local/sbin/ejabberdctl live
```

3. Create the account `admin@localhost` with password set as `password`:

```
/usr/local/sbin/ejabberdctl register admin localhost password
```

4. Now you can go to the web dashboard at `http://localhost:5280/admin/` and fill the username field with the full account JID, for example `admin@localhost`, then fill the password field with that account's `password`.
5. Without configuration there's not much to see here, therefore the next step is to get to know [how to configure ejabberd](#).

Installing ejabberd development environment on OSX

This short guide will show you how to compile ejabberd from source code on Mac OS X, and get users chatting right away.

Before you start

ejabberd is supported on Mac OS X 10.6.8 and later. Before you can compile and run ejabberd, you also need the following to be installed on your system:

- Gnu Make and GCC (the GNU Compiler Collection). To ensure that these are installed, you can install the Command Line Tools for Xcode, available via Xcode or from the Apple Developer website.
- Git
- Erlang/OTP 19.1 or higher. We recommend using Erlang 21.2.
- Autotools

Homebrew

An easy way to install some of the dependencies is by using a package manager, such as [Homebrew](#) – the Homebrew commands are provided here:

- Git: `brew install git`
- Erlang /OTP: `brew install erlang`
- Elixir: `brew install elixir`
- Autoconf: `brew install autoconf`
- Automake: `brew install automake`
- Openssl: `brew install openssl`
- Expat: `brew install expat`
- Libyaml: `brew install libyaml`
- Libiconv: `brew install libiconv`
- Sqlite: `brew install sqlite`
- GD: `brew install gd`
- Rebar: `brew install rebar rebar3`

You can install everything with a single command:

```
brew install erlang elixir openssl expat libyaml libiconv libgd sqlite rebar rebar3 automake autoconf
```

Installation

To build and install ejabberd from source code, do the following:

1. Clone the Git repository: `git clone git@github.com:processone/ejabberd.git`
2. Go to your ejabberd build directory: `cd ejabberd`
3. Run the following commands, assuming you want to install your ejabberd deployment into your home directory:

```
chmod +x autogen.sh
./autogen.sh
export LDFLAGS="-L/usr/local/opt/openssl/lib -L/usr/local/lib -L/usr/local/opt/expat/lib"
export CFLAGS="-I/usr/local/opt/openssl/include/ -I/usr/local/include -I/usr/local/opt/expat/include"
export CPPFLAGS="-I/usr/local/opt/openssl/include/ -I/usr/local/include -I/usr/local/opt/expat/include"
./configure --prefix=$HOME/my-ejabberd --enable-sqlite
make && make install
```

Note that the previous command reference the previously installed dependencies from [Homebrew](#).

Running ejabberd

- From your ejabberd build directory, go to the installation directory: `cd $HOME/my-ejabberd`
- To start the ejabberd server, run the following command: `sbin/ejabberdctl start`
- To verify that ejabberd is running, enter the following: `sbin/ejabberdctl status` If the server is running, response should be as follow:

```
$ sbin/ejabberdctl status
The node ejabberd@localhost is started with status: started
ejabberd 14.12.40 is running in that node
```

- To connect to the ejabberd console after starting the server: `sbin/ejabberdctl debug`
- Alternatively, you can also run the server in interactive mode: `sbin/ejabberdctl live`

Registering a user

The default XMPP domain served by ejabberd right after the build is `localhost`. This is different from the IP address, DNS name of the server. It means remote users can connect to ejabberd even if it is running on your machine with `localhost` XMPP domain, by using your computer IP address or DNS name. This can prove handy in development phase to get more testers.

Adium

Adium is a popular XMPP client on OSX. You can use it

1. Launch Adium. If the Adium Setup Assistant opens, close it.
2. In the **Adium** menu, select **Preferences**, and then select the **Accounts** tab.
3. Click the **+** button and select **XMPP (Jabber)**.
4. Enter a Jabber ID (for example, "user1@localhost") and password, and then click **Register New Account**.
5. In the **Server** field, enter the following:
6. Users registering on the computer on which ejabberd is running: `localhost`
7. Users registering from a different computer: the ejabberd server's IP address
8. Click **Request New Account**.

After registration, the user will connect automatically.

Registered users wishing to add an existing account to Adium should enter the ejabberd server's IP address in the **Connect Server** field on the **Options** tab.

Command line

You can register a user with the `ejabberdctl` utility: `ejabberdctl register user domain password`

For example: `ejabberdctl register user1 localhost myp4ssw0rd`

Domains

To use your system's domain name instead of localhost, edit the following ejabberd configuration file: `$HOME/my-ejabberd/etc/ejabberd.yml` (point to the place of your real installation).

Note: You may find example `ejabberd.cfg` files. This is the old obsolete format for configuration file. You can ignore the and focus on the new and more user friendly Yaml format.

Find the line listing the hosts:

```
hosts:  
- "localhost"
```

Replace `localhost` with your XMPP domain name, for example:

```
hosts:  
- "example.org"
```

Save the configuration file and restart the ejabberd server. A user's Jabber ID will then use the domain instead of localhost, for example: `user1@example.org`

You can also configure multiple (virtual) domains for one server:

```
hosts:  
- "example1.org"  
- "example2.org"
```

Get chatting

Users that are registered on your server can now add their accounts in a chat application like Adium (specifying either the server's IP address or domain name), add each other as contacts, and start chatting.

Next Steps

Starting ejabberd

Depending on how you installed ejabberd, it may be started automatically by the operating system at system boot time.

You can use the `ejabberdctl` command line administration script to start and stop ejabberd, check its status and many other administrative tasks.

If you provided the configure option `--enable-user=USER` (see compilation [options](#)), you can execute `ejabberdctl` with either that system account or root.

Usage example:

```
prompt> ejabberdctl start

prompt> ejabberdctl status
The node ejabberd@localhost is started with status: started
ejabberd is running in that node

prompt> ejabberdctl stop
```

If ejabberd doesn't start correctly and a crash dump file is generated, there was a severe problem. You can try to start ejabberd in interactive mode with the command `bin/ejabberdctl live` to see the error messages provided by Erlang and identify the exact the problem.

The `ejabberdctl` administration script is included in the `bin` directory in the Linux Installers and Docker image.

Please refer to the section [ejabberdctl](#) for details about `ejabberdctl`, and configurable options to fine tune the Erlang runtime system.

Autostart on Linux

If you compiled ejabberd from source code or some other method that doesn't setup autostarting ejabberd, you can try this method.

On a *nix system, create a system user called 'ejabberd', give it write access to the directories `database/` and `logs/`, and set that as home.

If you want ejabberd to be started as daemon at boot time with that user, copy `ejabberd.init` from the `bin` directory to something like `/etc/init.d/ejabberd`. Then you can call `/etc/init.d/ejabberd start` to start the server.

Or if you have a `systemd` distribution:

1. copy `ejabberd.service` to `/etc/systemd/system/`
2. run `systemctl daemon-reload`
3. run `systemctl enable ejabberd.service`
4. To start the server, you can run `systemctl start ejabberd`

When ejabberd is started, the processes that are started in the system are `beam` or `beam.smp`, and also `epmd`. For more information regarding `epmd` consult the section relating to [epmd](#).

Administration Account

admin account

Account registered in ejabberd with administrative privileges granted in the ejabberd configuration file. Features that take into consideration those privileges: [api_permissions](#), [ejabberd_web_admin](#), [mod_announce](#), [mod_configure](#), [mod_http_api](#), [mod_muc](#), ...

Some ejabberd installation methods ask you details for the first account, and take care to register that account and grant it administrative rights; in that case you can skip this section.

After installing ejabberd from source code or other methods, you may want to register the first XMPP account and grant it administrative rights:

1. Register an XMPP account on your ejabberd server. For example, if `example.org` is configured in the [hosts](#) section in your ejabberd configuration file, then you may want to register an account with JID `admin1@example.org`.

There are two ways to register an XMPP account in ejabberd:

- Using an XMPP client and [In-Band Registration](#).
- Using `ejabberdctl`:

```
ejabberdctl register admin1 example.org password
```

2. Edit the ejabberd configuration file to give administration rights to the XMPP account you registered:

```
acl:
  admin:
    user: admin1@example.org

access_rules:
  configure:
    allow: admin
```

You can grant administrative privileges to many XMPP accounts, and also to accounts in other XMPP servers.

3. Restart ejabberd to load the new configuration, or run the `reload_config` command.
4. Open the Web Admin page in your favourite browser. The exact address depends on your ejabberd configuration, and may be:
 - <http://localhost:5280/admin/> on [binary installers](#)
 - <https://localhost:5443/admin/> on [binary installers](#)
 - <https://localhost:5280/admin/> on [Debian package](#)
5. Your web browser shows a login window. Introduce the **full JID**, in this example `admin1@example.org`, and the account password. If the web address hostname is the same that the account JID, you can provide simply the username instead of the full JID: `admin1`. See [Web Admin](#) for details.

Configuring ejabberd

Now that you got ejabberd installed and running, it's time to configure it to your needs. You can follow on the [Configuration](#) section and take also a look at the [Tutorials](#).

Configure

Configuring ejabberd

Here are the main entry points to learn more about ejabberd configuration. ejabberd is extremely powerful and can be configured in many ways with many options.

Do not let this complexity scare you. Most of you will be fine with default config file (or light changes).

Tutorials for first-time users:

- [How to move to ejabberd XMPP server](#)
- [How to set up ejabberd video & voice calling \(STUN/TURN\)](#)
- [How to configure ejabberd to get 100% in XMPP compliance test](#)

Detailed documentation in sections:

- [File Format](#)
- [Basic Configuration](#): hosts, acl, logging...
- [Authentication](#): auth_method
- [Databases](#)
- [LDAP](#)
- [Listen Modules](#): c2s, s2s, http, sip, stun...
- [Listen Options](#)
- [Top-Level Options](#)
- [Modules Options](#)

There's also a [copy of the old configuration](#) document which was used up to ejabberd 20.03.

File format

Yaml File Format

ejabberd loads its configuration file during startup. This configuration file is written in [YAML](#) format, and its file name MUST have “.yaml” or “.yml” extension. This helps ejabberd to differentiate between this new format and the [legacy configuration file](#) format.

Please, consult `ejabberd.log` for configuration errors. ejabberd will report syntax related errors, as well as complains about unknown options and invalid values. Make sure you respect indentation (YAML is sensitive to this) or you will get pretty cryptic errors.

Note that ejabberd never edits the configuration file. If you are changing parameters at runtime from web admin interface, you will need to apply them to configuration file manually. This is to prevent messing up with your config file comments, syntax, etc.

Reload at Runtime

You can modify the ejabberd configuration file and reload it at runtime: the changes you made are applied immediately, no need to restart ejabberd. This applies to adding, changing or removing vhosts, listened ports, modules, ACLs or any other options.

How to do this?

1. Let's assume your ejabberd server is already running
2. Modify the configuration file
3. Run the [reload_config](#) command
4. ejabberd will read that file, check its YAML syntax is valid, check the options are valid and known...
5. If there's any problem in the configuration file, the reload is aborted and an error message is logged with details, so you can fix the problem.
6. If the file is right, it detects the changed options, and applies them immediately (add/remove hosts, add/remove modules, ...)

Legacy Configuration File

In previous ejabberd version the configuration file should be written in Erlang terms. The format is still supported, but it is highly recommended to convert it to the new YAML format with the [convert_to_yaml](#) API command using [ejabberdctl](#).

If you want to specify some options using the old Erlang format, you can set them in an additional cfg file, and include it using the `include_config_file` option, see [Include Additional Files](#).

Include Additional Files

The option [include_config_file](#) in a configuration file instructs ejabberd to include other configuration files immediately.

This is a basic example:

```
include_config_file: /etc/ejabberd/additional.yaml
```

In this example, the included file is not allowed to contain a `listen` option. If such an option is present, the option will not be accepted. The file is in a subdirectory from where the main configuration file is.

```
include_config_file:
  ./example.org/additional_not_listen.yaml:
    disallow: [listen]
```

Please notice that options already defined in the main configuration file cannot be redefined in the included configuration files. But you can use [host_config](#) and [append_host_config](#) as usual (see [Virtual Hosting](#)).

In this example, `ejabberd.yml` defines some ACL for the whole ejabberd server, and later includes another file:

```
acl:
  admin:
    user:
      - admin@localhost
include_config_file:
  /etc/ejabberd/acl.yml
```

The file `acl.yml` can add additional administrators to one of the virtual hosts:

```
append_host_config:
  localhost:
    acl:
      admin:
        user:
          - bob@localhost
          - jan@localhost
```

Macros and Keywords



improved in [25.03](#)

In the ejabberd configuration file, you can define a macro or keyword for a value (atom, integer, string...) and later use it when configuring an ejabberd option.

Macros is a feature implemented internally by the `yconf` library and are replaced early and transparently to ejabberd. However, macros **cannot** be defined inside `host_config`.

Keywords is a feature similar to macros, implemented by ejabberd itself, and are replaced after macro replacement. Keywords can be defined inside `host_config` for module options, but not for toplevel options. Keywords cannot be used in those toplevel options: `hosts`, `loglevel`, `version`.

First `define_macro` and `define_keyword` and then use them like this:

```
define_macro:
  NAME1: value1

define_keyword:
  NAME2: "value2"

some_option1: NAME1
other_option1: "I am @NAME1@"

some_option2: NAME2
other_option2: "I am @NAME2@"
```

where:

- **NAME**: should be specified in capital letters for convenience. Duplicated macro/keyword names are not allowed. If a macro is defined with the same name than a keyword, the macro is used.
- **value**: for all options, the value can be any valid YAML element. It is also possible to use as value the name of another macro.
- **use** a macro/keyword when configuring the option: simply set `NAME` instead of option value. Macros are processed after additional configuration files have been included, so it is possible to use macros that are defined in configuration files included before the usage.
- **use inside a string**: surround its name with `@` characters

Let's see examples of all this in detail:

Atom

```
define_macro:
  ANON: both

define_keyword:
  TLS: optional

anonymous_protocol: ANON
s2s_use_starttls: TLS
```

The resulting configuration is:

```
anonymous_protocol: both
s2s_use_starttls: optional
```

Integer

```
define_macro:
  LOG_LEVEL_NUMBER: 5
  NUMBER_PORT_C2S: 5222

define_keyword:
  NUMBER_PORT_HTTP: 5280

loglevel: LOG_LEVEL_NUMBER

listen:
  -
    port: NUMBER_PORT_C2S
    module: ejabberd_c2s
  -
    port: NUMBER_PORT_HTTP
    module: ejabberd_http
```

The resulting configuration is:

```
loglevel: 5

listen:
  -
    port: 5222
    module: ejabberd_c2s
  -
    port: 5280
    module: ejabberd_http
```

Map

Option values can be any arbitrary YAML value:

```
define_macro:
  USERBOB:
    user:
      - bob@localhost

define_keyword:
  USERJAN:
    user:
      - jan@localhost

acl:
  admin: USERBOB
  moderator: USERJAN
```

The resulting configuration is:

```
acl:
  admin:
    user:
      - bob@localhost
  moderator:
    user:
      - jan@localhost
```

String

```
define_macro:
  NAME: "MUC Service"
  PERSISTENT: true

define_keyword:
  TITLE: "Example Room"

modules:
  mod_muc:
    name: NAME
    default_room_options:
      persistent: PERSISTENT
    title: TITLE
```

The resulting configuration is:

```
modules:
  mod_muc:
    name: "MUC Service"
    default_room_options:
      persistent: true
      title: "Example Room"
```

Inside string

A macro or keyword can be used inside an option string:

```
define_keyword:
  CMD: "captcha"

captcha_cmd: "tools/@CMD@.sh"
```

is equivalent to:

```
define_keyword:
  CMD: "tools/captcha.sh"

captcha_cmd: "@CMD@"
```

is equivalent to:

```
define_keyword:
  CMD: "tools/captcha.sh"

captcha_cmd: CMD
```

The resulting configuration in all the cases is:

```
captcha_cmd: tools/captcha.sh
```

Macro over keyword

If a macro and a keyword are defined with the same name, the macro definition takes precedence and the keyword definition is ignored:

```
define_macro:
  LANGUAGE: "bg"

define_keyword:
  LANGUAGE: "pt"

language: LANGUAGE
```

The resulting configuration is:

```
language: "bg"
```

Keyword inside macro

A macro definition can use a keyword:

```
define_macro:
  MACRO: "tools/@KEYWORD@"

define_keyword:
  KEYWORD: "captcha.sh"

captcha_cmd: MACRO
```

The resulting configuration is:

```
captcha_cmd: "tools/captcha.sh"
```

Predefined keywords

Several keywords are predefined automatically by ejabberd, so you can use them without need to define them explicitly:

- **HOST**: the virtual [host name](#), for example `"example.org"`. That keyword is only predefined for module options, not toplevel options.
- **HOME**: path of the home directory of the user running ejabberd, for example `"/home/ejabberd"`
- **CONFIG_PATH**: path of the directory with ejabberd configuration files, for example `"/home/ejabberd/opt/ejabberd/conf"`
- **DATABASE_PATH**:  added in [26.04](#) . This keyword is the path of the Mnesia spool directory, that contains the Mnesia spool files, SQLite files, Erlang cookie... for example `"/home/ejabberd/opt/ejabberd/database"`
- **LOG_PATH**: path of the directory with ejabberd log files, for example `"/home/ejabberd/opt/ejabberd/logs"`
- **SEMVER**: ejabberd version number in [semver format](#) when compiled with Elixir's mix (`"24.5"`), otherwise it's in `XX.YY` format (`"24.05"`)
- **VERSION**: ejabberd version number in `XX.YY` format, for example `"24.05"`

It is possible to overwrite predefined keywords, global or for a vhost like in this example:

```
host_config:
  localhost:
    define_keyword:
      VERSION: "1.2.3"

ext_api_url: "http://localhost:@VERSION@/api"
```

The resulting behaviour is equivalent to a configuration like:

```
host_config:
  localhost:
    ext_api_url: "http://localhost/1.2.3/api"

ext_api_url: "http://localhost/25.03/api"
```

Macro and host_config

Macros can be **used** inside [host_config](#):

```
define_macro:
  MYSQL_PORT: 1234
  PGSQL_PORT: 4567

host_config:
  mysql.localhost:
    sql_port: MYSQL_PORT
  pgsql.localhost:
    sql_port: MYSQL_PORT
```

The resulting configuration is:

```
host_config:
  mysql.localhost:
    sql_port: 1234
  pgsql.localhost:
    sql_port: 4567
```

**Don't use macro defined in host_config**

Macros can not be **defined** inside `host_config`. Use the previous method instead. That problematic macro is not replaced:

```
host_config:
  mysql.localhost:
    define_macro:
      SQL_PORT: 1234
  pgsql.localhost:
    define_macro:
      SQL_PORT: 4567

sql_port: SQL_PORT

# [critical] Failed to start ejabberd application:
#   Invalid value of option sql_port:
#     Expected integer, got string instead
```

Keyword and host_config

Keywords can be **used** and **defined** inside `host_config`:

```
hosts:
- localhost
- example.org

define_keyword:
  HOSTNAME: "Generic Name"

host_config:
  example.org:
    define_keyword:
      HOSTNAME: "Example Host"

modules:
  mod_vcard:
    name: "vJUD of @HOSTNAME@"
```

The resulting configuration is:

```
host_config:
  localhost:
    modules:
      mod_vcard:
        name: "vJUD of Generic Name"
  example.org:
    modules:
      mod_vcard:
        name: "vJUD of Example Host"
```

**Don't use in toplevel a keyword defined in host_config**

Keywords can be defined inside `host_config`, but only if they are being used in module options, not in toplevel options. That problematic keyword is not replaced:

```
host_config:
  mysql.localhost:
    define_keyword:
      SQL_PORT: 1234
  pgsql.localhost:
    define_keyword:
      SQL_PORT: 4567

sql_port: SQL_PORT

# [critical] Failed to start ejabberd application:
#   Invalid value of option sql_port:
#     Expected integer, got string instead
```

Basic Configuration

XMPP Domains

Host Names

ejabberd supports managing several independent XMPP domains on a single ejabberd instance, using a feature called virtual hosting.

The option `hosts` defines a list containing one or more domains that ejabberd will serve.

Of course, the `hosts` list can contain just one domain if you do not want to host multiple XMPP domains on the same instance.

Examples:

- Serving one domain:

```
hosts: [example.org]
```

- Serving three domains:

```
hosts:  
- example.net  
- example.com  
- jabber.somesite.org
```

Notice that Unicode is used natively by XMPP, Erlang/OTP, YAML... and ejabberd. You can setup non-latin host names using Unicode/UTF8. Since ejabberd 25.08 the host can be expressed in IDNA/punycode. For example:

```
hosts:  
- localhost1  
- localhost2  
- 000  
- xn--localhost4-x2a
```

Virtual Hosting

When managing several XMPP domains in a single instance, those domains are truly independent. It means they can even have different configuration parameters.

Options can be defined separately for every virtual host using the `host_config` option.

Examples:

- Domain `example.net` is using the internal authentication method while domain `example.com` is using the LDAP server running on the domain `localhost` to perform authentication:

```
host_config:
example.net:
  auth_method: internal
example.com:
  auth_method: ldap
  ldap_servers:
    - localhost
  ldap_uids:
    - uid
  ldap_rootdn: "dc=localdomain"
  ldap_password: ""
```

- Domain `example.net` is using SQL to perform authentication while domain `example.com` is using the LDAP servers running on the domains `localhost` and `otherhost`:

```
host_config:
example.net:
  auth_method: sql
  sql_type: odbc
  sql_server: "DSN=ejabberd;UID=ejabberd;PWD=ejabberd"
example.com:
  auth_method: ldap
  ldap_servers:
    - localhost
    - otherhost
  ldap_uids:
    - uid
  ldap_rootdn: "dc=example,dc=com"
  ldap_password: ""
```

To define specific ejabberd modules in a virtual host, you can define the global `modules` option with the common modules, and later add specific modules to certain virtual hosts. To accomplish that, instead of defining each option in `host_config` use `append_host_config` with the same syntax.

In this example three virtual hosts have some similar modules, but there are also other different modules for some specific virtual hosts:

```
## This ejabberd server has three vhosts:
hosts:
  - one.example.org
  - two.example.org
  - three.example.org

## Configuration of modules that are common to all vhosts
modules:
  mod_roster: {}
  mod_configure: {}
  mod_disco: {}
  mod_private: {}
  mod_time: {}
  mod_last: {}
  mod_version: {}

append_host_config:
  ## Add some modules to vhost one:
  one.example.org:
    modules:
      mod_muc:
        host: conference.one.example.org
        mod_ping: {}
  ## Add a module just to vhost two:
  two.example.org:
    modules:
      mod_muc:
        host: conference.two.example.org
```

Logging

ejabberd configuration can help a lot by having the right amount of logging set up.

There are several toplevel options to configure logging:

- `loglevel`: Verbosity of log files generated by ejabberd.
- `hide_sensitive_log_data`: Privacy option to disable logging of IP address or sensitive data.
- `log_modules_fully`: Modules that will log everything independently from the general `loglevel` option.
- `log_rotate_size`
- `log_rotate_count`: Setting count to N keeps N rotated logs. Setting count to 0 does not disable rotation, it instead rotates the file and keeps no previous versions around. Setting size to X rotate log when it reaches X bytes.
- `log_burst_limit_count`
- `log_burst_limit_window_time`

The values in default configuration file are:

```
log_rotate_size: 10485760
log_rotate_count: 1
```

For example, log warning and higher messages, but all c2s messages, and hide sensitive data:

```
loglevel: warning
hide_sensitive_log_data: true
log_modules_fully: [ejabberd_c2s]
```

Default Language

The `language` option defines the default language of server strings that can be seen by XMPP clients. If a XMPP client does not support `xml:lang`, ejabberd uses the language specified in this option.

The option syntax is:

language: `Language`: The default value is `en`. In order to take effect there must be a translation file `Language.msg` in ejabberd's `msgs` directory.

For example, to set Russian as default language:

```
language: ru
```

The page [Internationalization and Localization](#) provides more details.

CAPTCHA

Some ejabberd modules can be configured to require a CAPTCHA challenge on certain actions, for instance `mod_block_strangers`, `mod_muc`, `mod_register`, and `mod_register_web`. If the client does not support CAPTCHA Forms ([XEP-0158](#)), a web link is provided so the user can fill the challenge in a web browser.

Example scripts are provided that generate the image using [ImageMagick's](#) Convert program and [Ghostschrift](#) fonts. Remember to install those dependencies: in Debian install the `imagemagick` and `gsfonts` packages; in container images check their documentation for details.

The relevant top-level options are:

- `captcha_cmd`: `Path` | `Module`: Full path to a script that generates the image, or name of a module that supports generating CAPTCHA images (`mod_ecaptcha`, `mod_captcha_rust`). The default value disables the feature: `undefined`
- `captcha_url`: `URL` | `auto`: An URL where CAPTCHA requests should be sent, or `auto` to determine the URL automatically. The default value is `auto`.

And finally, configure `request_handlers` for the `ejabberd_http` listener with a path handled by `ejabberd_captcha`, where the CAPTCHA images will be served.

Example configuration:

```
hosts: [example.org]

captcha_cmd: /lib/ejabberd-24.12/priv/bin/captcha.sh # make install or install-rel
# captcha_cmd: /opt/ejabberd-24.12/lib/captcha.sh # Binary installers and container images
# captcha_cmd: tools/captcha.sh # if ejabberd started with: make relive
# captcha_cmd: mod_ecaptcha # if installed: ejabberdctl module_install mod_ecaptcha

captcha_url: auto
## captcha_url: http://example.org:5280/captcha
## captcha_url: https://example.org:443/captcha
## captcha_url: http://example.com/captcha

listen:
-
  port: 5280
  module: ejabberd_http
  request_handlers:
    /captcha: ejabberd_captcha
```

ACME

[ACME](#) is used to automatically obtain SSL certificates for the domains served by ejabberd, which means that certificate requests and renewals are performed to some CA server (aka "ACME server") in a fully automated mode.

Setting up ACME

In ejabberd, ACME is configured using the [acme](#) top-level option, check there the available options and example configuration.

The automated mode is enabled by default. However, some configuration of ejabberd is still required, because ACME requires HTTP challenges: an ACME remote server will connect to your ejabberd server on HTTP port 80 during certificate issuance.

For that reason you must have an [ejabberd_http](#) listener with TLS disabled handling an "ACME well known" path. For example:

```
listen:
-
  module: ejabberd_http
  port: 5280
  tls: false
  request_handlers:
    /.well-known/acme-challenge: ejabberd_acme
```

Note that the ACME protocol **requires** challenges to be sent on port 80. Since this is a privileged port, ejabberd cannot listen on it directly without root privileges. Thus you need some mechanism to forward port 80 to the port defined by the listener (port 5280 in the example above). There are several ways to do this: using NAT, setcap (Linux only), or HTTP front-ends (e.g. [sslt](#), [nginx](#), [haproxy](#) and so on). Pick one that fits your installation the best, but **DON'T** run ejabberd as root.

If you see errors in the logs with ACME server problem reports, it's **highly** recommended to change `ca_url` option in the [acme](#) top-level option to the URL pointing to some staging ACME environment, fix the problems until you obtain a certificate, and then change the URL back and retry using `request_certificate ejabberdctl` command (see below). This is needed because ACME servers typically have rate limits, preventing you from requesting certificates too rapidly and you can get stuck for several hours or even days. By default, ejabberd uses [Let's Encrypt](#) authority. Thus, the default value of `ca_url` option is `https://acme-v02.api.letsencrypt.org/directory` and the staging URL will be `https://acme-staging-v02.api.letsencrypt.org/directory`:

```
acme:
  ## Staging environment
  ca_url: https://acme-staging-v02.api.letsencrypt.org/directory
  ## Production environment (the default):
  # ca_url: https://acme-v02.api.letsencrypt.org/directory
```

The automated mode can be disabled by setting `auto` option to `false` in the [acme](#) top-level option:

```
acme:
  auto: false
```

In this case automated renewals are still enabled, however, in order to request a new certificate, you need to run `request_certificate` API command:

```
ejabberdctl request_certificate all
```

If you only want to request certificates for a subset of the domains, run:

```
ejabberdctl request-certificate domain.tld,pubsub.domain.tld,server.com,conference.server.com,...
```

You can view the certificates obtained using ACME and [list_certificates](#):

```
$ ejabberdctl list-certificates
domain.tld /path/to/cert/file1 true
server.com /path/to/cert/file2 false
```

The output is mostly self-explained: every line contains the domain, the corresponding certificate file, and whether this certificate file is used or not. A certificate might not be used for several reasons: mostly because ejabberd detects a better certificate (i.e. not expired, or having a longer lifetime). It's recommended to revoke unused certificates if they are not yet expired (see below).

At any point you can revoke a certificate using [revoke_certificate](#): pick the certificate file from the listing above and run:

```
ejabberdctl revoke-certificate /path/to/cert/file
```

If the commands return errors, consult the log files for details.

ACME implementation details

In nutshell, certification requests are performed in two phases. Firstly, ejabberd creates an account at the ACME server. That is an EC private key. Secondly, a certificate is requested. In the case of a revocation, no account is used - only a certificate in question is needed. All information is stored under `acme` directory inside spool directory of ejabberd (typically `/var/lib/ejabberd`). An example content of the directory is the following:

```
$ tree /var/lib/ejabberd
/var/lib/ejabberd
├── acme
│   ├── account.key
│   └── live
│       ├── 251ce180d964e98a2f18b65504df2ab7c55943e2
│       └── 93816a8429ebbaa75574eb3f59d4a806b67d6917
└── ...
```

Here, `account.key` is the EC private key used to identify the ACME account. You can inspect its content using `openssl` command:

```
openssl ec -text -noout -in /var/lib/ejabberd/acme/account.key
```

Obtained certificates are stored under `acme/live` directory. You can inspect any of the certificates using `openssl` command as well:

```
openssl x509 -text -noout -in /var/lib/ejabberd/acme/live/251ce180d964e98a2f18b65504df2ab7c55943e2
```

In the case of errors, you can delete the whole `acme` directory - ejabberd will recreate its content on next certification request. However, don't delete it too frequently - usually there is a rate limit on the number of accounts and certificates an ACME server creates. In particular, for [Let's Encrypt](#) the limits are described [here](#).

Access Rights

This section describes new ACL syntax introduced in ejabberd 16.06. For old access rule and ACL syntax documentation, please refer to [configuration document archive](#)

ACL

Access control in ejabberd is performed via Access Control Lists (ACLs), using the `acl` option. The declarations of ACLs in the configuration file have the following syntax:

```
acl:
  ACLName:
    ACLType: ACLValue
```

`ACLName` is any name you may want to give, and later can use it in other options to reference that ACL. The following `ACLName` are pre-defined:

- `all` : Matches any JID.
- `none` : Matches no JID.

ACLType: ACLValue can be one of the following:

- **all** : Matches all JIDs. Example:

```
acl:
  world: all
```

- **user: Username** : Matches the user with the name `Username` on any of the local virtual host. Example:

```
acl:
  admin:
    user: yozhik
```

- **user: {Username: Server} | Jid** : Matches the user with the JID `Username@Server` and any resource. Example:

```
acl:
  admin:
    - user: peter@example.org
    - user:
        yozhik@example.org
    - user: {susan: example.org}
    - user:
        jan: example.org
```

- **server: Server** : Matches any JID from server `Server` . Example:

```
acl:
  exampleorg:
    server: example.org
```

- **resource: Resource** : Matches any JID with a resource `Resource` . Example:

```
acl:
  mucklres:
    resource: muckl
```

- **shared_group: Groupname** : Matches any member of a Shared Roster Group with name `Groupname` in the virtual host. Example:

```
acl:
  techgroupmembers:
    shared_group: techteam
```

- **shared_group: {Groupname: Server}** : Matches any member of a Shared Roster Group with name `Groupname` in the virtual host `Server` . Example:

```
acl:
  techgroupmembers:
    - shared_group: {techteam: example.org}
    - shared_group:
        secteam: example.org
```

- **ip: Network** : Matches any IP address from the `Network` . Example:

```
acl:
  loopback:
    ip:
      - 127.0.0.0/8
      - ":::1"
```

- **user_regexp: Regexp** : Matches any local user with a name that matches `Regexp` on local virtual hosts. Example:

```
acl:
  tests:
    user_regexp: "^test[0-9]*$"
```

- **user_regexp: {Regexp: Server} | JidRegexp** : Matches any user with a name that matches `Regexp` at server `Server` . Example:

```
acl:
  tests:
    user_regexp:
      - "^test1": example.org
      - "^test2@example.org"
```

- **server_regexp: Regexp** : Matches any JID from the server that matches `Regexp` . Example:

```
acl:
  icq:
    server_regexp: "^icq\\."
```

- **resource_regexp: Regexp** : Matches any JID with a resource that matches `Regexp` . Example:

```
acl:
icq:
  resource_regexp: "^laptop\\."
```

- **node_regexp: {UserRegexp: ServerRegexp}** : Matches any user with a name that matches `UserRegexp` at any server that matches `ServerRegexp` . Example:

```
acl:
yozhik:
  node_regexp:
    "yozhik$": "^example.(com|org)$"
```

- **user_glob: Glob** :
- **user_glob: {Glob: Server}** :
- **server_glob: Glob** :
- **resource_glob: Glob** :
- **node_glob: {UserGlob: ServerGlob}** : This is the same as above. However, it uses shell glob patterns instead of regexp. These patterns can have the following special characters:
 - `*` : matches any string including the null string.
 - `?` : matches any single character.
 - `[...]` : matches any of the enclosed characters. Character ranges are specified by a pair of characters separated by a `-` . If the first character after `[` is a `!` , any character not enclosed is matched.

Access Rules

The `access_rules` option is used to allow or deny access to different services. The syntax is:

```
access_rules:
  AccessName:
    - allow|deny: ACLName|ACLDefinition
```

Each definition may contain arbitrary number of `- allow` or `- deny` sections, and each section can contain any number of `acl` rules (as defined in [previous section](#), it recognizes one additional rule `acl: ACLName` that matches when the ACL named `ACLName` matches). If no ACL name or definition is defined, the ACL `all` is applied.

Definition's `- allow` and `- deny` sections are processed in top to bottom order, and first one for which all listed `acl` rules matches is returned as result of access rule. If no rule matches, `deny` is returned.

To simplify configuration two shortcut version are available: `- allow: acl` and `- allow`, example below shows equivalent definitions where short or long version are used:

```
access_rules:
  a_short: admin
  a_long:
    - acl: admin
  b_short:
    - deny: banned
    - allow
  b_long:
    - deny:
      - acl: banned
    - allow:
      - all
```

If you define specific Access rights in a virtual host, remember that the globally defined Access rights have precedence over those. This means that, in case of conflict, the Access granted or denied in the global server is used and the Access of a virtual host doesn't have effect.

Example:

```
access_rules:
  configure:
```

```

- allow: admin
something:
- deny: someone
- allow
s2s_banned:
- deny: problematic_hosts
- deny:
- acl: banned_forever
- deny:
- ip: 222.111.222.111/32
- deny:
- ip: 111.222.111.222/32
- allow
xmlrpc_access:
- allow:
- user: peter@example.com
- allow:
- user: ivone@example.com
- allow:
- user: bot@example.com
- ip: 10.0.0.0/24

```

The following `AccessName` are pre-defined:

- `all`: Always returns the value 'allow'.
- `none`: Always returns the value 'deny'.

Shaper Rules

The `shaper_rules` top-level option declares shapers to use for matching user/hosts. The syntax is:

```

shaper_rules:
  ShaperRuleName:
    - Number|ShaperName: ACLName|ACLDefinition

```

Semantic is similar to that described in [Access Rights](#) section, only difference is that instead using `- allow` or `- deny`, name of shaper or number should be used.

Examples:

```

shaper_rules:
  connections_limit:
    - 10:
      - user: peter@example.com
    - 100: admin
    - 5
  download_speed:
    - fast: admin
    - slow: anonymous_users
    - normal
  log_days: 30

```

Limiting Opened Sessions

The special access `max_user_sessions` specifies the maximum number of sessions (authenticated connections) per user. If a user tries to open more sessions by using different resources, the first opened session will be disconnected. The error `session replaced` will be sent to the disconnected session. The value for this option can be either a number, or `infinity`. The default value is `infinity`.

The syntax is:

```

shaper_rules:
  max_user_sessions:
    - Number: ACLName|ACLDefinition

```

This example limits the number of sessions per user to 5 for all users, and to 10 for admins:

```

shaper_rules:
  max_user_sessions:
    - 10: admin
    - 5

```

Connections to Remote Server

The special access `max_s2s_connections` specifies how many simultaneous S2S connections can be established to a specific remote XMPP server. The default value is `1`. There's also available the access `max_s2s_connections_per_node`.

The syntax is:

```
shaper_rules:
  max_s2s_connections: MaxNumber
```

For example, let's allow up to 3 connections with each remote server:

```
shaper_rules:
  max_s2s_connections: 3
```

Shapers

The `shaper` top-level option defines limitations in the connection traffic. The basic syntax is:

```
shaper:
  ShaperName: Rate
```

where `Rate` stands for the maximum allowed incoming rate in bytes per second. When a connection exceeds this limit, ejabberd stops reading from the socket until the average rate is again below the allowed maximum.

This example defines a shaper with name `normal` that limits traffic speed to 1,000bytes/second, and another shaper with name `fast` that limits traffic speed to 50,000bytes/second:

```
shaper:
  normal: 1000
  fast: 50000
```

You can use the full syntax to set the `BurstSize` too:

```
shaper:
  ShaperName:
    rate: Rate
    burst_size: BurstSize
```

With `BurstSize` you can allow client to send more data, but its amount can be clamped reasonably. Each connection is allowed to send `BurstSize` of data before processing is delayed, and that amount is replenished by `Rate` each second, but never more than what `BurstSize` allows. This allows the client to send quite a bit of data at once, but still have limited amount of data to send on constant basis.

In this example, the `normal` shaper has `Rate` set to `1000` and the `BurstSize` takes that same value. The `not_normal` shaper has the same `Rate` that before, and sets a higher `BurstSize`:

```
shaper:
  normal: 1000
  not_normal:
    rate: 1000
    burst_size: 20000
```

Authentication

Supported Methods

The authentication methods supported by ejabberd are:

- `internal` — See section [Internal](#).
- `external` — See section [External Script](#).
- `ldap` — See section [LDAP](#).
- `sql` — See section [Relational Databases](#).
- `anonymous` — See section [Anonymous Login and SASL Anonymous](#).
- `pam` — See section [PAM Authentication](#).
- `jwt` — See section [JWT Authentication](#).

The top-level option `auth_method` defines the authentication methods that are used for user authentication. The option syntax is:

```
auth_method: [Method1, Method2, ...]
```

When the `auth_method` option is omitted, ejabberd relies on the default database which is configured in `default_db` option. If this option is not set neither, then the default authentication method will be `internal`.

Account creation is only supported by `internal`, `external` and `sql` auth methods.

General Options

The top-level option `auth_password_format` allows to store the passwords in SCRAM format, see the [SCRAM](#) section.

Other top-level options that are relevant to the authentication configuration: [disable_sasl_mechanisms](#), [fqdn](#).

Authentication caching is enabled by default, and can be disabled in a specific vhost with the option `auth_use_cache`. The global authentication cache can be configured for all the authentication methods with the global top-level options: `auth_cache_missed`, `auth_cache_size`, `auth_cache_life_time`. For example:

```
auth_cache_size: 1500
auth_cache_life_time: 10 minutes

host_config:
  example.org:
    auth_method: [internal]
  example.net:
    auth_method: [ldap]
    auth_use_cache: false
```

Internal

ejabberd uses its internal Mnesia database as the default authentication method. The value `internal` will enable the internal authentication method.

To store the passwords in SCRAM format instead of plaintext, see the [SCRAM](#) section.

Examples:

- To use internal authentication on `example.org` and LDAP authentication on `example.net` :

```
host_config:
  example.org:
    auth_method: [internal]
  example.net:
    auth_method: [ldap]
```

- To use internal authentication with hashed passwords on all virtual hosts:

```
auth_method: internal
auth_password_format: scram
```

External Script

In the `external` authentication method, ejabberd uses a custom script to perform authentication tasks. The server administrator can write that external authentication script in any programming language.

Please check some example scripts, and the details on the interface between ejabberd and the script in the [Developers > Internals > External Authentication](#) section.

Options:

- [extauth_pool_name](#)
- [extauth_pool_size](#)
- [extauth_program](#)

Please note that caching interferes with the ability to maintain multiple passwords per account. So if your authentication mechanism supports application-specific passwords, caching must be disabled in the host that uses this authentication method with the option [auth_use_cache](#).

This example sets external authentication, specifies the extauth script, disables caching, and starts three instances of the script for each virtual host defined in ejabberd:

```
auth_method: [external]
extauth_program: /etc/ejabberd/JabberAuth.class.php
extauth_pool_size: 3
auth_use_cache: false
```

Anonymous Login and SASL Anonymous

The `anonymous` authentication method enables two modes for anonymous authentication:

Anonymous login : This is a standard login, that use the classical login and password mechanisms, but where password is accepted or preconfigured for all anonymous users. This login is compliant with SASL authentication, password and digest non-SASL authentication, so this option will work with almost all XMPP clients

SASL Anonymous : This is a special SASL authentication mechanism that allows to login without providing username or password (see [XEP-0175](#)). The main advantage of SASL Anonymous is that the protocol was designed to give the user a login. This is useful to avoid in some case, where the server has many users already logged or registered and when it is hard to find a free username. The main disadvantage is that you need a client that specifically supports the SASL Anonymous protocol.

The anonymous authentication method can be configured with the following options. Remember that you can use the [host_config](#) option to set virtual host specific options (see section [Virtual Hosting](#)):

- [allow_multiple_connections](#)
- [anonymous_protocol](#)

Examples:

- To enable anonymous login on all virtual hosts:

```
auth_method: [anonymous]
anonymous_protocol: login_anon
```

- Similar as previous example, but limited to `public.example.org`:

```
host_config:
  public.example.org:
    auth_method: [anonymous]
    anonymous_protocol: login_anon
```

- To enable anonymous login and internal authentication on a virtual host:

```
host_config:
  public.example.org:
    auth_method:
      - internal
      - anonymous
    anonymous_protocol: login_anon
```

- To enable SASL Anonymous on a virtual host:

```
host_config:
  public.example.org:
    auth_method: [anonymous]
    anonymous_protocol: sasl_anon
```

- To enable SASL Anonymous and anonymous login on a virtual host:

```
host_config:
  public.example.org:
    auth_method: [anonymous]
    anonymous_protocol: both
```

- To enable SASL Anonymous, anonymous login, and internal authentication on a virtual host:

```
host_config:
  public.example.org:
    auth_method:
      - internal
      - anonymous
    anonymous_protocol: both
```

There are more configuration examples and XMPP client example stanzas in [Anonymous users support](#).

PAM Authentication

ejabberd supports authentication via Pluggable Authentication Modules (PAM). PAM is currently supported in AIX, FreeBSD, HP-UX, Linux, Mac OS X, NetBSD and Solaris.

If compiling ejabberd from source code, PAM support is disabled by default, so you have to enable PAM support when [configuring](#) the ejabberd compilation: `./configure --enable-pam`

Options:

- [pam_service](#)
- [pam_userinfotype](#)

Example:

```
auth_method: [pam]
pam_service: ejabberd
```

Though it is quite easy to set up PAM support in ejabberd, there are several problems that you may need to solve:

- To perform PAM authentication, ejabberd uses external C-program called `epam`. By default, it is located in `/var/lib/ejabberd/priv/bin/` directory. You have to set it root on execution in the case when your PAM module requires root privileges (`pam_unix.so` for example). Also you have to grant access for ejabberd to this file and remove all other permissions from it. Execute with root privileges:

```
chown root:ejabberd /var/lib/ejabberd/priv/bin/epam
chmod 4750 /var/lib/ejabberd/priv/bin/epam
```

- Make sure you have the latest version of PAM installed on your system. Some old versions of PAM modules cause memory leaks. If you are not able to use the latest version, you can `kill(1) epam` process periodically to reduce its memory consumption: ejabberd will restart this process immediately.
- ejabberd [binary installers](#) include `epam` pointing to module paths that may not work in your system. If authentication doesn't work correctly, check if syslog (example: `journalctl -t epam -f`) reports errors like `PAM unable to dlopen(/home/runner/... No such file or directory`. In that case, create a PAM configuration file (example: `/etc/pam.d/ejabberd`) and provide the real path to that file in your machine:

```
##PAM-1.0
auth    sufficient /usr/lib/x86_64-linux-gnu/security/pam_unix.so audit
account sufficient /usr/lib/x86_64-linux-gnu/security/pam_unix.so audit
```

- `epam` program tries to turn off delays on authentication failures. However, some PAM modules ignore this behavior and rely on their own configuration options. You can create a configuration file (in Debian it would be `/etc/pam.d/ejabberd`). This example shows how to turn off delays in `pam_unix.so` module:

```
##PAM-1.0
auth    sufficient pam_unix.so likeauth nullok nodelay
account sufficient pam_unix.so
```

That is not a ready to use configuration file; you must use it as a hint when building your own PAM configuration instead. Note that if you want to disable delays on authentication failures in the PAM configuration file, you have to restrict access to this file, so a malicious user can't use your configuration to perform brute-force attacks.

- You may want to allow login access only for certain users. `pam_listfile.so` module provides such functionality.
- If you use `pam_winbind` to authorize against a Windows Active Directory, then `/etc/nsswitch.conf` must be configured to use `winbind` as well.

JWT Authentication

ejabberd supports authentication using JSON Web Token (JWT). When enabled, clients send signed tokens instead of passwords, which are checked using a public key (for asymmetric algorithms) or secret key (for symmetric algorithms) specified in the `jwt_key` option. JWT payload must look like this:

```
{
  "jid": "test@example.org",
  "exp": 1564436511
}
```

Options:

- [jwt_key](#)
- [jwt_auth_only_rule](#)
- [jwt_jid_field](#)

Example:

```
auth_method: jwt
jwt_key: /path/to/jwt/key
```

In this example, admins can use both JWT and plain passwords, while the rest of users can use only JWT.

```
# the order is important here, don't use [sql, jwt]
auth_method: [jwt, sql]

access_rules:
  jwt_only:
    deny: admin
    allow: all

jwt_auth_only_rule: jwt_only
```

Please notice that, when using JWT authentication, `mod_offline` will not work. With JWT authentication the accounts do not exist in the database, and there is no way to know if a given account exists or not.

For more information about JWT authentication, you can check a brief tutorial in the [ejabberd 19.08 release notes](#).

SCRAM

The top-level option `auth_password_format` defines in what format the users passwords are stored: SCRAM format or plaintext format.

The top-level option `auth_scram_hash` defines the hash algorithm that will be used to scram the password.

ejabberd supports channel binding to the external channel, allowing the clients to use `-PLUS` authentication mechanisms.

In summary, depending on the configured options, ejabberd supports:

- `SCRAM_SHA-1(-PLUS)`
- `SCRAM_SHA-256(-PLUS)`
- `SCRAM_SHA-512(-PLUS)`

For details about the client-server communication when using SCRAM, refer to [SASL Authentication and SCRAM](#).

Internal storage

When ejabberd starts with internal auth method and SCRAM password format configured:

```
auth_method: internal
auth_password_format: scram
```

and detects that there are plaintext passwords stored, they are automatically converted to SCRAM format:

```
[info] Passwords in Mnesia table 'passwd' will be SCRAM'ed
[info] Transforming table 'passwd', this may take a while
```

SQL Database

Please note that if you use SQL auth method and SCRAM password format, the plaintext passwords already stored in the database are not automatically converted to SCRAM format.

To convert plaintext passwords to SCRAM format in your database, use the `convert_to_scram` command:

```
ejabberdctl convert_to_scram example.org
```

Foreign authentication

Note on SCRAM using and foreign authentication limitations: when using the SCRAM password format, it is not possible to use foreign authentication method in ejabberd, as the real password is not known.

Foreign authentication are use to authenticate through various bridges ejabberd provide. Foreign authentication includes at the moment SIP and TURN auth support and they will not be working with SCRAM.

Database Configuration

ejabberd uses its internal Mnesia database by default. However, it is possible to use a relational database, key-value storage or an LDAP server to store persistent, long-living data.

ejabberd is very flexible: you can configure different authentication methods for different virtual hosts, you can configure different authentication mechanisms for the same virtual host (fallback), you can set different storage systems for modules, and so forth.

Supported storages

The following databases are supported by ejabberd:

- [Mnesia](#). Used by default, nothing to setup to start using it
- [MySQL](#). Check the tutorial [Using ejabberd with MySQL](#)
- [PostgreSQL](#)
- [MS SQL Server/SQL Azure](#). Check the [Microsoft SQL Server](#) section
- [SQLite](#)
- Any [ODBC](#) compatible database
- [Redis](#) (only for transient data). Check the [Redis](#) section
- [LDAP](#) is documented in the [LDAP](#) section

Virtual Hosting

If you define several [host names](#) in the `ejabberd.yml` configuration file, probably you want that each virtual host uses a different configuration of database, authentication and storage, so that usernames do not conflict and mix between different virtual hosts.

For that purpose, the options described in the next sections must be set inside the [host_config](#) top-level option for [each virtual host](#)).

For example:

```
host_config:
  public.example.org:
    sql_type: postgresql
    sql_server: localhost
    sql_database: database-public-example-org
    sql_username: ejabberd
    sql_password: password
    auth_method: [sql]
```

Default database

You can simplify your configuration by setting the default database with the [default_db](#) top-level option:

- it sets the default authentication method when the [auth_method](#) top-level option is not configured
- it defines the database to use in ejabberd modules that support the [db_type](#) option, when that option is not configured.

Database Schema



updated in [24.06](#)

**Update**

Modify the database schema and all its tables to match the installed ejabberd version. Not to be confused with [upgrade ejabberd](#) or [convert schema](#).

The [update_sql_schema](#) top-level option allows ejabberd to create and update the tables automatically in the SQL database when using MySQL, PostgreSQL or SQLite. That option was added in ejabberd [23.10](#), and enabled by default in [24.06](#). If you can use that feature:

1. Create the database in your SQL server
2. Create an account in the SQL server and grant it rights in the database
3. Configure in ejabberd the [SQL Options](#) that allow it to connect
4. Start ejabberd ...
5. and it will take care to create the tables (or update them if they exist from a previous ejabberd version)

If that option is disabled, or you are using a different SQL database, or an older ejabberd release, then you must create the tables in the database manually before starting ejabberd. The SQL database schema files are available:

- If installing ejabberd from sources, sql files are in the installation directory. By default: `/usr/local/lib/ejabberd/priv/sql`
- If installing ejabberd from Process-One installer, sql files are in the ejabberd's installation path under `<base>/lib/ejabberd*/priv/sql`

See [ejabberd SQL Database Schema](#) for details on database schemas.

Singlehost or Multihost



renamed from default/new to singlehost/multihost in [25.10](#)

If using MySQL, PostgreSQL, Microsoft SQL or SQLite, you can choose between two database schemas:

- the *singlehost* schema is preferable when serving one massive domain, or just a few domains,
- the *multihost* schema is preferable when serving many small domains.

The *singlehost* schema stores only one XMPP domain in the database. The [XMPP domain](#) is not stored as this is the same for all the accounts, and this saves space in massive deployments. However, to handle several domains, you have to setup one database per domain and configure each one independently using [host_config](#), so in that case you may prefer the *multihost* schema.

The *multihost* schema stores the XMPP domain in a new column `server_host` in the database entries, so it allows to handle several XMPP domains in a single ejabberd database. Using this schema is preferable when serving several XMPP domains and changing domains from time to time. However, if you have only one massive domain, you may prefer to use the *singlehost* schema.

To use the *multihost* schema, edit the ejabberd configuration file and enable [sql_schema_multihost](#) top-level option:

```
sql_schema_multihost: true
```

When creating the tables, if ejabberd can use the [update_sql_schema](#) top-level option as explained in the [Database Schema](#) section, it will take care to create the tables with the correct schema.

On the other hand, if you are creating the tables manually, remember to use the proper SQL schema! For example, if you are using MySQL and choose the *singlehost* schema, use `mysql.sql`. If you are using PostgreSQL and need the *multihost* schema, use `pg.new.sql`.



Change the database schema and all its tables from `singlehost` schema to `multihost` schema. Not to be confused with [upgrade ejabberd](#) or [update schema](#).

If you already have a MySQL or PostgreSQL database with the *singlehost* schema and contents, you can convert it to the *multihost* schema:

- **MySQL:** Edit the file `sql/mysql.old-to.new.sql` which is included with ejabberd, fill `DEFAULT_HOST` in the first line, and import that SQL file in your database. Then enable the `sql_schema_multihost` top-level option in the ejabberd configuration, and restart ejabberd.
- **PostgreSQL:** First enable the `sql_schema_multihost` top-level option and `mod_admin_update_sql` in your ejabberd configuration:

```
sql_schema_multihost: true
modules:
  mod_admin_update_sql: {}
```

then restart ejabberd, and finally execute the [update_sql](#) command:

```
ejabberdctl update_sql
```

SQL Options

The actual database access is defined in the options with `sql_` prefix. The values are used to define if we want to use ODBC, or one of the two native interface available, PostgreSQL or MySQL.

To configure SQL there are several top-level options:

- [sql_type](#)
- [sql_server](#)
- [sql_port](#)
- [sql_database](#)
- [sql_username](#)
- [sql_password](#)
- [sql_ssl](#), see section [SQL with SSL connection](#)
- [sql_ssl_verify](#)
- [sql_ssl_cafile](#)
- [sql_ssl_certfile](#)
- [sql_pool_size](#)
- [sql_keepalive_interval](#)
- [sql_odbc_driver](#)
- [sql_start_interval](#)
- [sql_prepared_statements](#)
- [update_sql_schema](#), see section [Database Schema](#)
- [sql_schema_multihost](#), see section [Singlehost or Multihost](#)

Example of plain ODBC connection:

```
sql_server: "DSN=database;UID=ejabberd;PWD=password"
```

Example of MySQL connection:

```
sql_type: mysql
sql_server: server.company.com
sql_port: 3306 # the default
sql_database: mydb
sql_username: user1
sql_password: "*****"
sql_pool_size: 5
```

SQL with SSL Connection

The `sql_ssl` top-level option allows SSL encrypted connections to MySQL, PostgreSQL, and Microsoft SQL servers.

Please notice that ejabberd verifies the certificate presented by the SQL server against the CA certificate list. For that reason, if your SQL server uses a self-signed certificate, you need to setup `sql_ssl_verify` and `sql_ssl_cafile`, for example:

```
sql_ssl: true
sql_ssl_verify: false
sql_ssl_cafile: "/path/to/sql_server_cacert.pem"
```

This tells ejabberd to ignore problems from not matching any CA certificate from default list, and instead try to verify using the specified CA certificate.

SQL Authentication

You can authenticate users against an SQL database, see the option `auth_method` in the [Authentication](#) section.

To store the passwords in SCRAM format instead of plaintext, see the [SCRAM](#) section.

SQL Storage

Several ejabberd [modules](#) have options called `db_type`, and can store their tables in an SQL database instead of internal.

In this sense, if you defined your database access using the [SQL Options](#), you can configure a module to use your database by adding the option `db_type: sql` to that module.

Alternatively, if you want all modules to use your SQL database when possible, you may prefer to set SQL as your [default database](#).

Microsoft SQL Server

For now, MS SQL is only supported in Unix-like OS'es. Your Erlang/OTP must be compiled with ODBC support. And ejabberd too: either [configuring](#) compilation with `./configure --enable-mssql --enable-odbc` or `./configure --enable-all`.

Also, in some cases you need to add the machine name to `sql_username`, especially when you have `sql_server` defined as an IP address, e.g.:

```
sql_type: mssql
sql_server: 1.2.3.4
sql_username: user1@host
```

By default, ejabberd will use the [FreeTDS](#) driver. You need to have the driver file `libtdsodbc.so` installed in your library PATH on your system. For example, in Debian, install with `apt-get install tdsodbc`.

If the FreeTDS driver is not installed in a standard location, or if you want to use another ODBC driver, you can specify the path to the driver using the `sql_odbc_driver` option, available in ejabberd [20.12](#) or later. For example, if you want to use Microsoft ODBC Driver 17 for SQL Server:

```
sql_odbc_driver: "/opt/microsoft/msodbcsql17/lib64/libmsodbcsql-17.3.so.1.1"
```

Note that if you use a Microsoft driver, you may have to use an IP address instead of a host name for the `sql_server` option.

If hostname (or IP address) is specified in `sql_server` option, ejabberd will connect using a an ODBC DSN connection string constructed with:

- `SERVER=sql_server`
- `DATABASE=sql_database`
- `UID=sql_username`
- `PWD=sql_password`
- `PORT=sql_port`
- `ENCRYPTION=required` (only if `sql_ssl` is true)
- `CLIENT_CHARSET=UTF-8`

Since ejabberd 23.04, it is possible to use different connection options by putting a full ODBC connection string in `sql_server` (e.g. `DSN=database;UID=ejabberd;PWD=password`). The DSN must be configured in existing system or user `odbc.ini` file, where it can be configured as desired, using a driver from system `odbcinst.ini`. The `sql_odbc_driver` option will have no effect in this case and `unixODBC` must be installed on your machine.

If specifying an ODBC connection string, an ODBC connection string must also be specified for any other hosts using MS SQL DB, otherwise the auto-generated ODBC configuration will interfere.

Redis

`Redis` is an advanced key-value cache and store. You can use it to store transient data, such as records for C2S (client) sessions.

The available top-level options are:

- `redis_server`
- `redis_port`
- `redis_password`
- `redis_db`
- `redis_connect_timeout`

Example configuration:

```
redis_server: redis.server.com
redis_db: 1
```

LDAP Configuration

Supported storages

The following LDAP servers are tested with ejabberd:

- [Active Directory](#) (see section [Active Directory](#))
- [OpenLDAP](#)
- [CommuniGate Pro](#)
- Normally any LDAP compatible server should work; inform us about your success with a not-listed server so that we can list it here.

LDAP

ejabberd has built-in LDAP support. You can authenticate users against LDAP server and use LDAP directory as vCard storage.

Usually ejabberd treats LDAP as a read-only storage: it is possible to consult data, but not possible to create accounts or edit vCard that is stored in LDAP. However, it is possible to change passwords if `mod_register` module is enabled and LDAP server supports [RFC 3062](#).

LDAP Connection

Two connections are established to the LDAP server per vhost, one for authentication and other for regular calls.

To configure the LDAP connection there are these top-level options:

- [ldap_servers](#)
- [ldap_backups](#)
- [ldap_encrypt](#)
- [ldap_tls_verify](#)
- [ldap_tls_certfile](#)
- [ldap_tls_cacertfile](#)
- [ldap_tls_depth](#)
- [ldap_port](#)
- [ldap_rootdn](#)
- [ldap_password](#)
- [ldap_deref_aliases](#)

Example:

```
auth_method: [ldap]
ldap_servers:
- ldap1.example.org
ldap_port: 389
ldap_rootdn: "cn=Manager,dc=domain,dc=org"
ldap_password: "*****"
```

When there are several LDAP servers available as backup, set one in `ldap_servers` and the others in `ldap_backups`. At server start, ejabberd connects to all the servers listed in `ldap_servers`. If a connection is lost, ejabberd connects to the next server in `ldap_backups`. If the connection is lost, the next server in the list is connected, and this repeats infinitely with all the servers in `ldap_servers` and `ldap_backups` until one is successfully connected:

```
ldap_servers:
- ldap1.example.org
```

```
ldap_backups:
- ldap2.example.org
- ldap3.example.org
```

LDAP Authentication

You can authenticate users against an LDAP directory. Note that current LDAP implementation does not support SASL authentication.

To configure LDAP authentication there are these top-level options:

- [ldap_base](#)
- [ldap_uids](#)
- [ldap_filter](#)
- [ldap_dn_filter](#)

LDAP Examples

Common example

Let's say `ldap.example.org` is the name of our LDAP server. We have users with their passwords in `ou=Users,dc=example,dc=org` directory. Also we have addressbook, which contains users emails and their additional infos in `ou=AddressBook,dc=example,dc=org` directory. The connection to the LDAP server is encrypted using TLS, and using the custom port 6123. Corresponding authentication section should look like this:

```
## Authentication method
auth_method: [ldap]
## DNS name of our LDAP server
ldap_servers: [ldap.example.org]
## Bind to LDAP server as "cn=Manager,dc=example,dc=org" with password "secret"
ldap_rootdn: "cn=Manager,dc=example,dc=org"
ldap_password: secret
ldap_encrypt: tls
ldap_port: 6123
## Define the user's base
ldap_base: "ou=Users,dc=example,dc=org"
## We want to authorize users from 'shadowAccount' object class only
ldap_filter: "(objectClass=shadowAccount)"
```

Now we want to use users LDAP-info as their vCards. We have four attributes defined in our LDAP schema: `mail` — email address, `givenName` — first name, `sn` — second name, `birthDay` — birthday. Also we want users to search each other. Let's see how we can set it up:

```
modules:
  mod_vcard:
    db_type: ldap
    ## We use the same server and port, but want to bind anonymously because
    ## our LDAP server accepts anonymous requests to
    ## "ou=AddressBook,dc=example,dc=org" subtree.
    ldap_rootdn: ""
    ldap_password: ""
    ## define the addressbook's base
    ldap_base: "ou=AddressBook,dc=example,dc=org"
    ## uidattr: user's part of JID is located in the "mail" attribute
    ## uidattr_format: common format for our emails
    ldap_uids:
      mail: "%u@mail.example.org"
    ## We have to define empty filter here, because entries in addressbook does not
    ## belong to shadowAccount object class
    ldap_filter: ""
    ## Now we want to define vCard pattern
    ldap_vcard_map:
      NICKNAME: {"%u": []} # just use user's part of JID as their nickname
      GIVEN: {"%s": [givenName]}
      FAMILY: {"%s": [sn]}
      FN: {"%s, %s": [sn, givenName]} # example: "Smith, John"
      EMAIL: {"%s": [mail]}
      BDAY: {"%s": [birthDay]}
    ## Search form
    ldap_search_fields:
      User: "%u"
      Name: givenName
      "Family Name": sn
      Email: mail
      Birthday: birthDay
```

```
## vCard fields to be reported
## Note that JID is always returned with search results
ldap_search_reported:
  "Full Name": FN
  Nickname: NICKNAME
  Birthday: BDAY
```

Note that `mod_vcard` with LDAP backend checks for the existence of the user before searching their information in LDAP.

Active Directory

Active Directory is just an LDAP-server with predefined attributes. A sample configuration is shown below:

```
auth_method: [ldap]
ldap_servers: [office.org] # List of LDAP servers
ldap_base: "DC=office,DC=org" # Search base of LDAP directory
ldap_rootdn: "CN=Administrator,CN=Users,DC=office,DC=org" # LDAP manager
ldap_password: "*****" # Password to LDAP manager
ldap_uids: [sAMAccountName]
ldap_filter: "(memberOf=*)"

modules:
  mod_vcard:
    db_type: ldap
    ldap_vcard_map:
      NICKNAME: {"%u": []}
      GIVEN: {"%s": [givenName]}
      MIDDLE: {"%s": [initials]}
      FAMILY: {"%s": [sn]}
      FN: {"%s": [displayName]}
      EMAIL: {"%s": [mail]}
      ORGNAME: {"%s": [company]}
      ORGUNIT: {"%s": [department]}
      CTRY: {"%s": [c]}
      LOCALITY: {"%s": [l]}
      STREET: {"%s": [streetAddress]}
      REGION: {"%s": [st]}
      PCODE: {"%s": [postalCode]}
      TITLE: {"%s": [title]}
      URL: {"%s": [wwwHomePage]}
      DESC: {"%s": [description]}
      TEL: {"%s": [telephoneNumber]}
    ldap_search_fields:
      User: "%u"
      Name: givenName
      "Family Name": sn
      Email: mail
      Company: company
      Department: department
      Role: title
      Description: description
      Phone: telephoneNumber
    ldap_search_reported:
      "Full Name": FN
      Nickname: NICKNAME
      Email: EMAIL
```

Shared Roster in LDAP

Since `mod_shared_roster_ldap` has a few complex options, some of them are documented with more detail here:

Filters

ldap_ufilter: “User Filter” – used for retrieving the human-readable name of roster entries (usually full names of people in the roster). See also the parameters `ldap_userdesc` and `ldap_userid`. If unspecified, defaults to the top-level parameter of the same name. If that one also is unspecified, then the filter is assembled from values of other parameters as follows (`[ ldap_SOMETHING ]` is used to mean “the value of the configuration parameter `ldap_SOMETHING`”):

```
(&(&([ldap_memberattr]=[ldap_memberattr_format])([ldap_groupattr]=%g))[ldap_filter])
```

Subsequently `%u` and `%g` are replaced with a `*`. This means that given the defaults, the filter sent to the LDAP server would be `(&(memberUid=*)(cn=*))`. If however the `ldap_memberattr_format` is something like `uid=%u,ou=People,o=org`, then the filter will be `(&(memberUid=uid=*,ou=People,o=org)(cn=*))`.

ldap_filter: Additional filter which is AND-ed together with *User Filter* and *Group Filter*. If unspecified, defaults to the top-level parameter of the same name. If that one is also unspecified, then no additional filter is merged with the other filters.

Note that you will probably need to manually define the *User* and *Group Filter* (since the auto-assembled ones will not work) if:

- your `ldap_memberattr_format` is anything other than a simple `%u`,
- **and** the attribute specified with `ldap_memberattr` does not support substring matches.

An example where it is the case is OpenLDAP and *(unique)MemberName* attribute from the *groupOf(Unique)Names* objectClass. A symptom of this problem is that you will see messages such as the following in your `slapd.log`:

```
get_filter: unknown filter type=139
filter="(&(?=undefined)(?=undefined)(something=else))"
```

Control parameters

These parameters control the behaviour of the module.

`ldap_memberattr_format_re`: A regex for extracting user ID from the value of the attribute named by `ldap_memberattr`.

An example value `"CN=(\\w*), (OU=.* ,)*DC=company,DC=com"` works for user IDs such as the following:

- `CN=Romeo,OU=Montague,DC=company,DC=com`
- `CN=Abram,OU=Servants,OU=Montague,DC=company,DC=com`
- `CN=Juliet,OU=Capulet,DC=company,DC=com`
- `CN=Peter,OU=Servants,OU=Capulet,DC=company,DC=com`

In case:

- the option is unset,
- or the `re` module is unavailable in the current Erlang environment,
- or the regular expression does not compile,

then instead of a regular expression, a simple format specified by `ldap_memberattr_format` is used. Also, in the last two cases an error message is logged during the module initialization.

Also, note that in all cases `ldap_memberattr_format` (and **not** the regex version) is used for constructing the default “User/Group Filter” — see section [Filters](#).

Retrieving the roster

When the module is called to retrieve the shared roster for a user, the following algorithm is used:

1. [step:rfilter] A list of names of groups to display is created: the *Roster Filter* is run against the base DN, retrieving the values of the attribute named by `ldap_groupattr`.
2. Unless the group cache is fresh (see the `ldap_group_cache_validity` option), it is refreshed:
 - a. Information for all groups is retrieved using a single query: the *Group Filter* is run against the Base DN, retrieving the values of attributes named by `ldap_groupattr` (group ID), `ldap_groupdesc` (group "Display Name") and `ldap_memberattr` (IDs of group members).
 - b. group "Display Name", read from the attribute named by `ldap_groupdesc`, is stored in the cache for the given group
 - c. the following processing takes place for each retrieved value of attribute named by `ldap_memberattr`:
 - i. the user ID part of it is extracted using `ldap_memberattr_format(_re)`,
 - ii. then (unless `ldap_auth_check` is set to `off`) for each found user ID, the module checks (using the ejabberd authentication subsystem) whether such user exists in the given virtual host. It is skipped if the check is enabled and fails. This step is here for historical reasons. If you have a tidy DIT and properly defined "Roster Filter" and "Group Filter", it is safe to disable it by setting `ldap_auth_check` to `off` — it will speed up the roster retrieval.
 - iii. the user ID is stored in the list of members in the cache for the given group.
3. For each item (group name) in the list of groups retrieved in step [step:rfilter]:
 - a. the display name of a shared roster group is retrieved from the group cache
 - b. for each IDs of users which belong to the group, retrieved from the group cache:
 - i. the ID is skipped if it's the same as the one for which we are retrieving the roster. This is so that the user does not have himself in the roster.
 - ii. the display name of a shared roster user is retrieved:
- A. first, unless the user name cache is fresh (see the `ldap_user_cache_validity` option), it is refreshed by running the *User Filter*, against the Base DN, retrieving the values of attributes named by `ldap_userid` and `ldap_userdesc`.
- B. then, the display name for the given user ID is retrieved from the user name cache.

Multi-Domain

By default, the module option `ldap_userjidattr` is set to the empty string, in that case the JID of the user's contact is formed by compounding UID of the contact @ Host of the user owning the roster.

When the option `ldap_userjidattr` is set to something like "mail", then it uses that field to determine the JID of the contact. This is useful if the `ldap mail` attribute contains the JID of the accounts.

Basically, it allows us to define a `groupOfNames` (e.g. `xmppRosterGroup`) and list any users, anywhere in the ldap directory by specifying the attribute defining the JID of the members.

This allows hosts/domains other than that of the roster owner. It is also more flexible, since the LDAP manager can specify the JID of the users without any assumptions being made. The only down side is that there must be an LDAP attribute (field) filled in for all Jabber/XMPP users.

Below is a sample, a relevant LDAP entry, and ejabberd's module configuration:

```
cn=Example Org Roster,ou=groups,o=Example Organisation,dc=acme,dc=com
objectClass: groupOfNames
objectClass: xmppRosterGroup
objectClass: top
xmppRosterStatus: active
member:
description: Roster group for Example Org
cn: Example Org Roster
uniqueMember: uid=john,ou=people,o=Example Organisation,dc=acme,dc=com
```

```

uniqueMember: uid=pierre,ou=people,o=Example Organisation,dc=acme,dc=com
uniqueMember: uid=jane,ou=people,o=Example Organisation,dc=acme,dc=com

uid=john,ou=people,o=Example Organisation,dc=acme,dc=com
objectClass: top
objectClass: person
objectClass: organizationalPerson
objectClass: inetOrgPerson
objectClass: mailUser
objectClass: sipRoutingObject
uid: john
givenName: John
sn: Doe
cn: John Doe
displayName: John Doe
accountStatus: active
userPassword: secretpass
IMAPURL: imap://imap.example.net:143
mailHost: smtp.example.net
mail: john@example.net
sipLocalAddress: john@example.net

```

Below is the sample ejabberd.yml module configuration to match:

```

mod_shared_roster_ldap:
  ldap_servers:
    - "ldap.acme.com"
  ldap_encrypt: tls
  ldap_port: 636
  ldap_rootdn: "cn=Manager,dc=acme,dc=com"
  ldap_password: "supersecretpass"
  ldap_base: "dc=acme,dc=com"
  ldap_filter: "(objectClass=*)"
  ldap_rfilter: "(&(objectClass=xmppRosterGroup)(xmppRosterStatus=active))"
  ldap_gfilter: "(&(objectClass=xmppRosterGroup)(xmppRosterStatus=active)(cn=%g))"
  ldap_groupattr: "cn"
  ldap_groupdesc: "cn"
  ldap_memberattr: "uniqueMember"
  ldap_memberattr_format_re: "uid=([a-z.]*), (ou=.*,)*(o=.*,)*dc=acme,dc=com"
  ldap_userid: "uid"
  ldap_userdesc: "cn"
  ldap_userjidattr: "mail"
  ldap_auth_check: false
  ldap_user_cache_validity: 86400
  ldap_group_cache_validity: 86400


```

Configuration examples

Since there are many possible [DIT](#) layouts, it will probably be easiest to understand how to configure the module by looking at an example for a given DIT (or one resembling it).

FLAT DIT

This seems to be the kind of DIT for which this module was initially designed. Basically there are just user objects, and group membership is stored in an attribute individually for each user. For example in a layout like this, it's stored in the `ou` attribute:



Such layout has a few downsides, including:


- information duplication – the group name is repeated in every member object
- difficult group management – information about group members is not centralized, but distributed between member objects
- inefficiency – the list of unique group names has to be computed by iterating over all users

This however seems to be a common DIT layout, so the module keeps supporting it. You can use the following configuration...

```


modules:
  mod_shared_roster_ldap:
    ldap_base: "ou=flat,dc=nodomain"
    ldap_rfilter: "(objectClass=inetOrgPerson)"
    ldap_groupattr: ou
    ldap_memberattr: cn
    ldap_filter: "(objectClass=inetOrgPerson)"
    ldap_userdesc: displayName
  
```

...to be provided with a roster upon connecting as user `czesio`, as shown in this figure:



DEEP DIT


This type of DIT contains distinctly typed objects for users and groups – see the next figure. They are shown separated into different subtrees, but it's not a requirement.



If you use the following example module configuration with it:

```
modules:
mod_shared_roster_ldap:
  ldap_base: "ou=deep,dc=nodomain"
  ldap_rfilter: "(objectClass=groupOfUniqueNames)"
  ldap_filter: ""
  ldap_gfilter: "(&(objectClass=groupOfUniqueNames)(cn=%g))"
  ldap_groupdesc: description
  ldap_memberattr: uniqueMember
  ldap_memberattr_format: "cn=%u,ou=people,ou=deep,dc=nodomain"
  ldap_ufilter: "(&(objectClass=inetOrgPerson)(cn=%u))"
  ldap_userdesc: displayName
```

...and connect as user `czesio`, then ejabberd will provide you with the roster shown in this figure:



vCard in LDAP

Since LDAP may be complex to configure in [mod_vcard](#), this section provides more details.

ejabberd can map LDAP attributes to vCard fields. This feature is enabled when the `mod_vcard` module is configured with `db_type:`

`ldap`. Notice that it does not depend on the authentication method (see [LDAP Authentication](#)).

Usually ejabberd treats LDAP as a read-only storage: it is possible to consult data, but not possible to create accounts or edit vCard that is stored in LDAP. However, it is possible to change passwords if `mod_register` module is enabled and LDAP server supports [RFC 3062](#).

This feature has its own optional parameters. The first group of parameters has the same meaning as the top-level LDAP parameters to set the authentication method: `ldap_servers`, `ldap_port`, `ldap_rootdn`, `ldap_password`, `ldap_base`, `ldap_uids`, `ldap_deref_aliases` and `ldap_filter`. See section [LDAP Authentication](#) for detailed information about these options. If one of these options is not set, ejabberd will look for the top-level option with the same name.

Examples:

- Let's say `ldap.example.org` is the name of our LDAP server. We have users with their passwords in `ou=Users,dc=example,dc=org` directory. Also we have addressbook, which contains users emails and their additional infos in `ou=AddressBook,dc=example,dc=org` directory. Corresponding authentication section should look like this:

```
## authentication method
auth_method: ldap
## DNS name of our LDAP server
ldap_servers:
- ldap.example.org
## We want to authorize users from 'shadowAccount' object class only
ldap_filter: "(objectClass=shadowAccount)"
```

- Now we want to use users LDAP-info as their vCards. We have four attributes defined in our LDAP schema: `mail` — email address, `givenName` — first name, `sn` — second name, `birthDay` — birthday. Also we want users to search each other. Let's see how we can set it up:

```
modules:
mod_vcard:
  db_type: ldap
  ## We use the same server and port, but want to bind anonymously because
  ## our LDAP server accepts anonymous requests to
  ## "ou=AddressBook,dc=example,dc=org" subtree.
  ldap_rootdn: ""
  ldap_password: ""
  ## define the addressbook's base
  ldap_base: "ou=AddressBook,dc=example,dc=org"
  ## uidattr: user's part of JID is located in the "mail" attribute
  ## uidattr_format: common format for our emails
  ldap_uids: {"mail": "%u@mail.example.org"}
  ## Now we want to define vCard pattern
  ldap_vcard_map:
    NICKNAME: {"%u": []} # just use user's part of JID as their nickname
    FIRST: {"%s": [givenName]}
    LAST: {"%s": [sn]}
    FN: {"%s, %s": [sn, givenName]} # example: "Smith, John"
    EMAIL: {"%s": [mail]}
    BDAY: {"%s": [birthDay]}
  ## Search form
  ldap_search_fields:
    User: "%u"
    Name: givenName
    "Family Name": sn
    Email: mail
    Birthday: birthDay
  ## vCard fields to be reported
  ## Note that JID is always returned with search results
  ldap_search_reported:
    "Full Name": FN
    Nickname: NICKNAME
    Birthday: BDAY
```

Note that `mod_vcard` with LDAP backend checks an existence of the user before searching their info in LDAP.

- `ldap_vcard_map` example:

```
ldap_vcard_map:
  NICKNAME: {"%u": []} # just use user's part of JID as their nickname
  FN: {"%s": [displayName]}
  CTRY: {Russia: []}
  EMAIL: {"%u@%d": []}
  DESC: {"%s\n%s": [title, description]}
```

- `ldap_search_fields` example:

```
ldap_search_fields:
  User: uid
  "Full Name": displayName
  Email: mail
```

- `ldap_search_reported` example:

```
ldap_search_reported:
  "Full Name": FN
  Email: EMAIL
  Birthday: BDAY
  Nickname: NICKNAME
```

Listen Modules

Please note

This section describes the most recent ejabberd version. If you are using an old ejabberd release, please refer to the corresponding archived version of this page in the [Archive](#).

Listen Option

The [listen](#) top-level option defines for which ports, addresses and network protocols ejabberd will listen and what services will be run on them.

Each element of the list is an associative array of [listen options](#):

- **port** and **module** mandatory options
- **ip**, **transport**, **unix_socket** additional options
- other options for listening module, enumerated later.

For example:

```
listen:
-
  port: 5222
  ip: 127.0.0.1
  module: ejabberd_c2s
  starttls: true
-
  port: 5269
  transport: tcp
  module: ejabberd_s2s_in
```

ejabberd_c2s



Client to Server connection in XMPP. Opposed to [s2s](#).

Handles c2s connections.

General listen options supported: [access](#), [allow_unencrypted_sasl2](#), [cafile](#), [ciphers](#), [dhfile](#), [max_fsm_queue](#), [max_stanza_elements](#), [max_stanza_size](#), [pre_auth_max_stanza_elements](#), [pre_auth_max_stanza_size](#), [protocol_options](#), [send_timeout](#), [shaper](#), [starttls](#), [starttls_required](#), [tls](#), [tls_compression](#), [tls_verify](#), [zlib](#).

ejabberd_s2s_in



Server to Server connection in XMPP. Opposed to [c2s](#).

Handles incoming s2s connections.

General listen options supported: [cafile](#), [ciphers](#), [dhfile](#), [max_fsm_queue](#), [max_stanza_elements](#), [max_stanza_size](#), [pre_auth_max_stanza_elements](#), [pre_auth_max_stanza_size](#), [protocol_options](#), [send_timeout](#), [shaper](#), [tls](#), [tls_compression](#).

ejabberd_service

Interacts with an [external component](#) as defined in [XEP-0114: Jabber Component Protocol](#).

General listen options supported: [access](#), [cafile](#), [certfile](#), [check_from](#), [ciphers](#), [dhfile](#), [global_routes](#), [hosts](#), [max_fsm_queue](#), [max_stanza_size](#), [password](#), [protocol_options](#), [send_timeout](#), [shaper](#), [shaper_rule](#), [tls](#), [tls_compression](#).

mod_mqtt

Support for MQTT requires configuring `mod_mqtt` both in the [listen](#) and the [modules](#) sections. Check the [mod_mqtt module](#) options, and the [MQTT Support](#) section.

General listen options supported: [backlog](#), [max_fsm_queue](#), [max_payload_size](#), [send_timeout](#), [tls](#), [tls_verify](#).

ejabberd_stun

ejabberd can act as a stand-alone STUN/TURN server, and this module handles STUN/TURN requests as defined in [RFC 5389](#) / [RFC 5766](#). In that role ejabberd helps clients with ICE ([RFC 5245](#)) or Jingle ICE ([XEP-0176](#)) support to discover their external addresses and ports and to relay media traffic when it is impossible to establish direct peer-to-peer connection.

General listen options supported: [certfile](#), [send_timeout](#), [shaper](#), [tls](#),

The specific `ejabberd_stun` configurable options are:

- **auth_realm:** *String*

When `auth_type` is set to `user` and you have several virtual hosts configured you should set this option explicitly to the virtual host you want to serve on this particular listening port. Implies `use_turn`.

- **auth_type:** *user|anonymous*

Which authentication type to use for TURN allocation requests. When type `user` is set, ejabberd authentication backend is used. For `anonymous` type no authentication is performed (not recommended for public services). The default is `user`. Implies `use_turn`.

- **shaper:** *Atom*

For `tcp` transports defines shaper to use. The default is `none`.

- **server_name:** *String*

Defines software version to return with every response. The default is the STUN library version.

- **turn_blacklist:** *String | [String,...]*

Specify one or more IP addresses and/or subnet addresses/masks. The TURN server will refuse to relay traffic from/to blacklisted IP addresses. By default, loopback addresses (`127.0.0.0/8` and `::1/128`) are blacklisted.

- **turn_ipv4_address:** *String*

The IPv4 address advertised by your TURN server. The address should not be NAT'ed or firewalled. There is not default, so you should set this option explicitly. Implies `use_turn`.

- **turn_ipv6_address:** *String*

The IPv6 address advertised by your TURN server. The address should not be NAT'ed or firewalled. There is not default, so you should set this option explicitly. Implies `use_turn`.

- **turn_max_allocations:** *Integer|infinity*

Maximum number of TURN allocations available from the particular IP address. The default value is 10. Implies `use_turn`.

- **turn_max_permissions:** *Integer|infinity*

Maximum number of TURN permissions available from the particular IP address. The default value is 10. Implies `use_turn`.

- **turn_max_port:** *Integer*

Together with `turn_min_port` forms port range to allocate from. The default is 65535. Implies `use_turn`.

- **turn_min_port:** *Integer*

Together with `turn_max_port` forms port range to allocate from. The default is 49152. Implies `use_turn`.

- **use_turn:** *true|false*

Enables/disables TURN (media relay) functionality. The default is `false`.

Example configuration with disabled TURN functionality (STUN only):

```
listen:
-
  port: 5478
  transport: udp
  module: ejabberd_stun
-
  port: 5478
  module: ejabberd_stun
-
  port: 5349
  module: ejabberd_stun
  tls: true
  certfile: /etc/ejabberd/server.pem
```

Example configuration with TURN functionality. Note that STUN is always enabled if TURN is enabled. Here, only UDP section is shown:

```
listen:
-
  port: 5478
  transport: udp
```

```
use_turn: true
turn_ipv4_address: 10.20.30.1
module: ejabberd_stun
```

ejabberd_sip

ejabberd has built-in support to handle SIP requests as defined in [RFC 3261](#).

To activate this feature, add the `ejabberd_sip` listen module, enable `mod_sip` module for the desired virtual host, and configure DNS properly.

To add a listener you should configure `ejabberd_sip` listening module as described in [Listen](#) section. If option `tls` is specified, option `certfile` must be specified as well, otherwise incoming TLS connections would fail.

General listen options supported: `certfile`, `send_timeout`, `tls`.

Example configuration with standard ports (as per [RFC 3261](#)):

```
listen:
-
  port: 5060
  transport: udp
  module: ejabberd_sip
-
  port: 5060
  module: ejabberd_sip
-
  port: 5061
  module: ejabberd_sip
  tls: true
  certfile: /etc/ejabberd/server.pem
```

Note that there is no StartTLS support in SIP and [SNI](#) support is somewhat tricky, so for TLS you have to configure different virtual hosts on different ports if you have different certificate files for them.

Next you need to configure DNS SIP records for your virtual domains. Refer to [RFC 3263](#) for the detailed explanation. Simply put, you should add NAPTR and SRV records for your domains. Skip NAPTR configuration if your DNS provider doesn't support this type of records. It's not fatal, however, highly recommended.

Example configuration of NAPTR records:

```
example.com IN NAPTR 10 0 "s" "SIPS+D2T" "" _sips._tcp.example.com.
example.com IN NAPTR 20 0 "s" "SIP+D2T" "" _sip._tcp.example.com.
example.com IN NAPTR 30 0 "s" "SIP+D2U" "" _sip._udp.example.com.
```

Example configuration of SRV records with standard ports (as per [RFC 3261](#)):

```
_sip._udp IN SRV 0 0 5060 sip.example.com.
_sip._tcp IN SRV 0 0 5060 sip.example.com.
_sips._tcp IN SRV 0 0 5061 sip.example.com.
```

Warning

SIP authentication does not support SCRAM. As such, it is not possible to use `mod_sip` to authenticate when ejabberd has been set to encrypt password with SCRAM.

ejabberd_http

Handles incoming HTTP connections.

With the proper request handlers configured, this serves HTTP services like [ACME](#), [API ReST](#), [BOSH](#), [CAPTCHA](#), [Converse.js](#), [Fileserver](#), [Matrix](#), [OAuth](#), [RegisterWeb](#), [Upload](#), [WebAdmin](#), [WebSocket](#), [XML-RPC](#), [XMPP Provider](#), [XMPP host-meta](#).

The [ejabberd-contrib](#) git repository contains other useful HTTP services: [CAPTCHA \(Rust\)](#), [CAPTCHA \(Erlang\)](#), [HTTP Redirect](#), [MUC Logs](#), [Prometheus](#), [ReST \(stanzas and API\)](#), [Unified Push](#), [WebPresence](#).

Options: [cafile](#), [ciphers](#), [custom_headers](#), [dhfile](#), [protocol_options](#), [request_handlers](#), [send_timeout](#), [tag](#), [tls](#), [tls_compression](#), and the [trusted_proxies](#) top-level option.

ejabberd_http_ws

This module enables XMPP communication over WebSocket connection as described in [RFC 7395](#).

WEBSOCKET CONFIG

To enable WebSocket, simply add a handler to the `request_handlers` section of an `ejabberd_http` listener:

```
listen:
-
  port: 5280
  module: ejabberd_http
  request_handlers:
    /websocket: ejabberd_http_ws
```

This module can be configured using those top-level options:

- [websocket_origin](#)
- [websocket_ping_interval](#)
- [websocket_timeout](#)

WEBSOCKET DISCOVERY

With the example configuration previously mentioned, the WebSocket URL would be: `ws://localhost:5280/websocket`

You may want to provide a `host-meta` file so clients can easily discover WebSocket service for your XMPP domain (see [XEP-0156](#)). One easy way to provide that file is using [mod_host_meta](#).

TESTING WEBSOCKET

A test client can be found on Github: [WebSocket test client](#)

There is an example configuration for WebSocket and Converse.js in the ejabberd [21.12](#) release notes.

ejabberd_xmlrpc

Handles XML-RPC requests to execute [API commands](#). It is configured as a request handler in [ejabberd_http](#).

This is the minimum configuration required to enable the feature:

```
listen:
-
  port: 5280
  module: ejabberd_http
  request_handlers:
    /xmlrpc: ejabberd_xmlrpc

api_permissions:
  "public commands":
    who:
      ip: 127.0.0.1/8
    what:
      - connected_users_number
```

Example Python3 script:

```
import xmlrpc.client
server = xmlrpc.client.ServerProxy("http://127.0.0.1:5280/xmlrpc/");
print(server.connected_users_number())
```

By default there is no restriction to who can execute what commands, so it is strongly recommended that you configure restrictions using [API Permissions](#).

This example configuration adds some restrictions (only requests from localhost are accepted, the XML-RPC query must include authentication credentials of a specific account registered in ejabberd, and only two commands are accepted):

```
listen:
-
  port: 5280
  ip: ".*"
  module: ejabberd_http
  request_handlers:
    /xmlrpc: ejabberd_xmlrpc

api_permissions:
  "some XMLRPC commands":
    from: ejabberd_xmlrpc
    who:
      - ip: 127.0.0.1
      - user: user1@localhost
    what:
      - registered_users
      - connected_users_number
```

Example Python3 script for that restricted configuration:

```
import xmlrpc.client
server = xmlrpc.client.ServerProxy("http://127.0.0.1:5280/xmlrpc/");

params = {}
params['host'] = 'localhost'

auth = {'user': 'user1',
        'server': 'localhost',
        'password': 'mypass11',
        'admin': True}

def calling(command, data):
    fn = getattr(server, command)
    return fn(auth, data)

print(calling('registered_users', params))
```

Please notice, when using the old Python2, replace the two first lines with:

```
import xmlrpclib
server = xmlrpclib.Server("http://127.0.0.1:5280/xmlrpc/");
```

It's possible to use OAuth for authentication instead of plain password, see [OAuth Support](#).

In ejabberd 20.03 and older, it was possible to configure `ejabberd_xmlrpc` as a listener.

Just for reference, there's also the old [ejabberd_xmlrpc documentation](#) with example clients in other languages.

Examples

For example, the following simple configuration defines:

- There are three domains. The default certificate file is `server.pem`. However, the c2s and s2s connections to the domain `example.com` use the file `example_com.pem`.
- Port 5222 listens for c2s connections with STARTTLS, and also allows plain connections for old clients.
- Port 5223 listens for c2s connections with the old SSL.
- Port 5269 listens for s2s connections with STARTTLS. The socket is set for IPv6 instead of IPv4.
- Port 5478 listens for STUN requests over UDP.
- Port 5280 listens for HTTP requests, and serves the HTTP-Bind (BOSH) service.
- Port 5281 listens for HTTP requests, using HTTPS to serve HTTP-Bind (BOSH) and the Web Admin as explained in [Managing: Web Admin](#). The socket only listens connections to the IP address 127.0.0.1.

```
hosts:
- example.com
- example.org
- example.net

certfiles:
- /etc/ejabberd/server.pem
- /etc/ejabberd/example_com.pem

listen:
-
```

```

port: 5222
module: ejabberd_c2s
access: c2s
shaper: c2s_shaper
starttls: true
max_stanza_size: 65536
-
port: 5223
module: ejabberd_c2s
access: c2s
shaper: c2s_shaper
tls: true
max_stanza_size: 65536
-
port: 5269
ip: ":::"
module: ejabberd_s2s_in
shaper: s2s_shaper
max_stanza_size: 131072
-
port: 5478
transport: udp
module: ejabberd_stun
-
port: 5280
module: ejabberd_http
request_handlers:
  /bosh: mod_bosh
-
port: 5281
ip: 127.0.0.1
module: ejabberd_http
tls: true
request_handlers:
  /admin: ejabberd_web_admin
  /bosh: mod_bosh

s2s_use_starttls: optional
outgoing_s2s_families:
  - ipv4
  - ipv6
outgoing_s2s_timeout: 10000
trusted_proxies: [127.0.0.1, 192.168.1.11]

```

In this example, the following configuration defines that:

- c2s connections are listened for on port 5222 (all IPv4 addresses) and on port 5223 (SSL, IP 192.168.0.1 and fdca:8ab6:a243:75ef::1) and denied for the user called 'bad'.
- s2s connections are listened for on port 5269 (all IPv4 addresses) with STARTTLS for secured traffic strictly required, and the certificates are verified. Incoming and outgoing connections of remote XMPP servers are denied, only two servers can connect: "jabber.example.org" and "example.com".
- Port 5280 is serving the Web Admin and the HTTP-Bind (BOSH) service in all the IPv4 addresses. Note that it is also possible to serve them on different ports. The second example in section [Managing: Web Admin](#) shows how exactly this can be done. A request handler to serve MQTT over WebSocket is also defined.
- All users except for the administrators have a traffic of limit 1,000Bytes/second
- The [AIM transport](#) aim.example.org is connected to port 5233 on localhost IP addresses (127.0.0.1 and ::1) with password 'aimsecret'.
- The [ICQ transport](#) JIT (icq.example.org and sms.example.org) is connected to port 5234 with password 'jitsecret'.
- The [MSN transport](#) msn.example.org is connected to port 5235 with password 'msnsecret'.
- The [Yahoo! transport](#) yahoo.example.org is connected to port 5236 with password 'yahoosecret'.
- The [Gadu-Gadu transport](#) gg.example.org is connected to port 5237 with password 'ggsecret'.
- The [Jabber Mail Component](#) jmc.example.org is connected to port 5238 with password 'jmcsecret'.
- The service custom has enabled the special option to avoiding checking the from attribute in the packets send by this component. The component can send packets in behalf of any users from the server, or even on behalf of any server.

```

acl:
  blocked:
    user: bad
  trusted_servers:
    server:
      - example.com
      - jabber.example.org
  xmlrpc_bot:
    user:

```

```

- xmllrpc-robot@example.org
shaper:
  normal: 1000
shaper_rules:
  c2s_shaper:
    - none: admin
    - normal
access_rules:
  c2s:
    - deny: blocked
    - allow
xmllrpc_access:
  - allow: xmllrpc_bot
s2s:
  - allow: trusted_servers
certfiles:
  - /path/to/ssl.pem
s2s_access: s2s
s2s_use_starttls: required_trusted
listen:
-
  port: 5222
  module: ejabberd_c2s
  shaper: c2s_shaper
  access: c2s
-
  ip: 192.168.0.1
  port: 5223
  module: ejabberd_c2s
  tls: true
  access: c2s
-
  ip: "FDCA:8AB6:A243:75EF::1"
  port: 5223
  module: ejabberd_c2s
  tls: true
  access: c2s
-
  port: 5269
  module: ejabberd_s2s_in
-
  port: 5280
  module: ejabberd_http
  request_handlers:
    /admin: ejabberd_web_admin
    /bosh: mod_bosh
    /mqtt: mod_mqtt
-
  port: 4560
  module: ejabberd_xmllrpc
  access_commands: {}
-
  ip: 127.0.0.1
  port: 5233
  module: ejabberd_service
  hosts:
    aim.example.org:
      password: aimsecret
-
  ip: "::1"
  port: 5233
  module: ejabberd_service
  hosts:
    aim.example.org:
      password: aimsecret
-
  port: 5234
  module: ejabberd_service
  hosts:
    icq.example.org:
      password: jitsecret
    sms.example.org:
      password: jitsecret
-
  port: 5235
  module: ejabberd_service
  hosts:
    msn.example.org:
      password: msnsecret
-
  port: 5236
  module: ejabberd_service
  password: yahoosecret
-
  port: 5237
  module: ejabberd_service
  hosts:
    gg.example.org:
      password: ggsecret
-
  port: 5238
  module: ejabberd_service
  hosts:
    jmc.example.org:

```

```

    password: jmcsecret
-
  port: 5239
  module: ejabberd_service
  check_from: false
  hosts:
    custom.example.org:
      password: customsecret

```

Note, that for services based in jabberd14 or WPJabber you have to make the transports log and do XDB by themselves:

```

<!--
  You have to add elogger and rlogger entries here when using ejabberd.
  In this case the transport will do the logging.
-->

<log id='logger'>
  <host/>
  <logtype/>
  <format>%d: [%t] (%h): %s</format>
  <file>/var/log/jabber/service.log</file>
</log>

<!--
  Some XMPP server implementations do not provide
  XDB services (for example, jabberd2 and ejabberd).
  xdb_file.so is loaded in to handle all XDB requests.
-->

<xdb id="xdb">
  <host/>
  <load>
    <!-- this is a lib of wpjabber or jabberd14 -->
    <xdb_file>/usr/lib/jabber/xdb_file.so</xdb_file>
  </load>
  <xdb_file xmlns="jabber:config:xdb_file">
    <spool><jabberd:cmdline flag='s'>/var/spool/jabber</jabberd:cmdline></spool>
  </xdb_file>
</xdb>

```

Listen Options

Please note

This section describes the most recent ejabberd version. If you are using an old ejabberd release, please refer to the corresponding archived version of this page in the [Archive](#).

This is a detailed description of each option allowed by the listening modules:

access

AccessName

This option defines access to the port. The default value is `all`.

allow_unencrypted_sasl2

`true` | `false`

As per [XEP-0388](#), ejabberd rejects SASL2 negotiations over non-TLS connections by default. Setting this option to `true` allows SASL2 over plaintext connections, which may be useful in case TLS is terminated by some proxy in front of ejabberd.

backlog

Value

The backlog value defines the maximum length that the queue of pending connections may grow to. This should be increased if the server is going to handle lots of new incoming connections as they may be dropped if there is no space in the queue (and ejabberd was not able to accept them immediately). Default value is 5.

cafile

Path

Path to a file of CA root certificates. The default is to use system defined file if possible.

This option is useful to define the file for a specific port listener. To set a file for all client listeners or for specific vhosts, you can use the `c2s_cafile` top-level option. To set a file for all server connections, you can use the `s2s_cafile` top-level option or the `ca_file` top-level option.

Please note: if this option is set in `ejabberd_c2s` or `ejabberd_s2s_in` and the corresponding top-level option is also set (`c2s_cafile`, `s2s_cafile`), then the top-level option is used, not this one.

certfile

Path

Path to the certificate file. Only makes sense when the `tls` options is set. If this option is not set, you should set the `certfiles` top-level option or configure [ACME](#).

check_from

`true` | `false`

This option can be used with `ejabberd_service` only. [XEP-0114](#) requires that the domain must match the hostname of the component. If this option is set to `false`, ejabberd will allow the component to send stanzas with any arbitrary domain in the 'from' attribute. Only use this option if you are completely sure about it. The default value is `true`, to be compliant with [XEP-0114](#).

ciphers

Ciphers

OpenSSL ciphers list in the same format accepted by 'openssl ciphers' command.

Please note: if this option is set in `ejabberd_c2s` or `ejabberd_s2s_in` and the corresponding top-level option is also set (`c2s_ciphers`, `s2s_ciphers`), then the top-level option is used, not this one.

custom_headers

{Name: Value}

Specify additional HTTP headers to be included in all HTTP responses. Default value is: `[]`

dhfile

Path

Full path to a file containing custom parameters for Diffie-Hellman key exchange. Such a file could be created with the command `openssl dhparam -out dh.pem 2048`. If this option is not specified, default parameters will be used, which might not provide the same level of security as using custom parameters.

Please note: if this option is set in `ejabberd_c2s` or `ejabberd_s2s_in` and the corresponding top-level option is also set (`c2s_dhfile`, `s2s_dhfile`), then the top-level option is used, not this one.

global_routes

true | false

This option emulates legacy behaviour which registers all routes defined in `hosts` on a component connected. This behaviour is considered harmful in the case when it's desired to multiplex different components on the same port, so, to disable it, set `global_routes` to `false`.

The default value is `true`, e.g. legacy behaviour is emulated: the only reason for this is to maintain backward compatibility with existing deployments.

hosts

{Hostname: [HostOption, ...]}

The external Jabber component that connects to this `ejabberd_service` can serve one or more hostnames. As `HostOption` you can define options for the component; currently the only allowed option is the password required to the component when attempt to connect to ejabberd: `password: Secret`. Note that you cannot define in a single `ejabberd_service` components of different services: add an `ejabberd_service` for each service, as seen in an example below. This option may not be necessary if the component already provides the host in its packets; in that case, you can simply provide the password option that will be used for all the hosts (see port 5236 definition in the example below).

ip

string()

The socket will listen only in that network interface. Depending on the type of the IP address, IPv4 or IPv6 will be used.

It is possible to specify a generic address (`"0.0.0.0"` for IPv4 or `:::"` for IPv6), so ejabberd will listen in all addresses. Note that on some operating systems and/or OS configurations, listening on `:::"` will mean listening for IPv4 traffic as well as IPv6 traffic.

Some example values for IP address:

- `"0.0.0.0"` to listen in all IPv4 network interfaces. This is the default value when the option is not specified.
- `:::"` to listen in all IPv6 network interfaces
- `"10.11.12.13"` is the IPv4 address `10.11.12.13`
- `:::FFFF:127.0.0.1"` is the IPv6 address `:::FFFF:127.0.0.1/128`

max_fsm_queue

Size

This option specifies the maximum number of elements in the queue of the FSM (Finite State Machine). Roughly speaking, each message in such queues represents one XML stanza queued to be sent into its relevant outgoing stream. If queue size reaches the limit (because, for example, the receiver of stanzas is too slow), the FSM and the corresponding connection (if any) will be terminated and error message will be logged. The reasonable value for this option depends on your hardware configuration. This option can be specified for `ejabberd_service` and `ejabberd_c2s` listeners, or also globally for `ejabberd_s2s_out`. If the option is not specified for `ejabberd_service` or `ejabberd_c2s` listeners, the globally configured value is used. The allowed values are integers and 'undefined'. Default value: '10000'.

max_payload_size

Size

Specify the maximum payload size in bytes. It can be either an integer or the word `infinity`. The default value is `infinity`.

max_stanza_elements

pos_integer()

 added in 26.04

This option specifies a maximum number of XML elements that individual stanza can have. The default value is `infinity`.

See also `pre_auth_max_stanza_elements` for value used by streams before authentication.

max_stanza_size

pos_integer()

This option specifies an approximate maximum size in bytes of XML stanzas. Approximate, because it is calculated with the precision of one block of read data. For example `65536`.

The default value is `infinity`. Recommended values are `65536` for c2s connections and `131072` for s2s connections. Please notice that the `s2s` max stanza size must be always much higher than `c2s` limit.

Change this value with extreme care as it can cause unwanted disconnect if set too low.

See also `pre_auth_max_stanza_size` for value used by streams before authentication.

module

ModuleName

Mandatory option to define what `listening module` will serve the incoming connections to the port.

password

Secret

Specify the password to verify an external component that connects to the port.

port

pos_integer() | string()

 improved in [20.07](#)

This mandatory option defines which port to listen for incoming connections: it can be a Jabber/XMPP standard port or any other valid port number between `1` and `65535` to listen on TCP or UDP socket,

Alternatively, set the option to a string in form `"unix:/path/to/socket"` to create and listen on a unix domain socket `/path/to/socket`.



Tip

 improved in [25.03](#)

If it's a relative path, then it's created in the mnesia spool directory. For example, if set to `"unix:dir/file.socket"`, then the socket file is created in `/opt/ejabberd/database/dir/file.socket`, or whatever path the Mnesia database is stored in your installation.

File permissions can be set using the [unix_socket](#) option.

pre_auth_max_stanza_elements

pos_integer()

 added in [26.04](#)

Sets limit of maximum number of XML elements that individual stanza can have, this value is only used for streams that haven't authenticated yet. Default value is set to `32`.

See also [max_stanza_elements](#) for value used by streams after authentication completes.

pre_auth_max_stanza_size

pos_integer()

 added in [26.04](#)

Sets limit of maximum size of bytes that XML representation of stanza can have, this value is only used for streams that haven't authenticated yet. Default value is set to `8192`.

See also [max_stanza_size](#) for value used by streams after authentication completes.

protocol_options

ProtocolOpts

List of general options relating to SSL/TLS. These map to [OpenSSL's set_options\(\)](#). The default entry is: `"no_sslv3|cipher_server_preference|no_compression"`

Please note: if this option is set in [ejabberd_c2s](#) or [ejabberd_s2s_in](#) and the corresponding top-level option is also set ([c2s_protocol_options](#), [s2s_protocol_options](#)), then the top-level option is used, not this one.

request_handlers

{Path: Module}

To define one or several handlers that will serve HTTP requests in [ejabberd_http](#). The Path is a string; so the URIs that start with that Path will be served by Module. For example, if you want `mod_foo` to serve the URIs that start with `/a/b/`, and you also want `mod_bosh` to serve the URIs `/bosh/`, use this option:

```
request_handlers:
  /a/b: mod_foo
  /bosh: mod_bosh
  /mqtt: mod_mqtt
```

send_timeout

Integer | infinity

 new in [21.07](#)

Sets the longest time that data can wait to be accepted to sent by OS socket. Triggering this timeout will cause the server to close it. By default it's set to 15 seconds, expressed in milliseconds: 15000

shaper

none | ShaperName

This option defines a shaper for the port (see section [Shapers](#)). The default value is `none`.

shaper_rule

none | ShaperRule

This option defines a shaper rule for [ejabberd_service](#) (see section [Shapers](#)). The recommended value is `fast`.

starttls

true | false

This option specifies that STARTTLS encryption is available on connections to the port. You should also set the [certfiles](#) top-level option or configure [ACME](#).

This option gets implicitly enabled when enabling [starttls_required](#) or [tls_verify](#).

starttls_required

true | false

This option specifies that STARTTLS encryption is required on connections to the port. No unencrypted connections will be allowed. You should also set the [certfiles](#) top-level option or configure [ACME](#).

Enabling this option implicitly enables also the [starttls](#) option.

tag

String

Allow specifying a tag in a `listen` section and later use it to have a special [api_permissions](#) just for it.

For example:

```
listen:
-
  port: 4000
  module: ejabberd_http
  tag: "magic_listener"

api_permissions:
  "magic_access":
    from:
      - tag: "magic_listener"
    who: all
    what: ""
```

The default value is the empty string: `""`.

timeout

Integer

Timeout of the connections, expressed in milliseconds. Default: 5000

tls

true | false

This option specifies that traffic on the port will be encrypted using SSL immediately after connecting. This was the traditional encryption method in the early Jabber software, commonly on port 5223 for client-to-server communications. But this method is nowadays deprecated and not recommended. The preferable encryption method is STARTTLS on port 5222, as defined [RFC 6120: XMPP Core](#), which can be enabled in ejabberd with the option `starttls`.

If this option is set, you should also set the `certfiles` top-level option or configure [ACME](#).

The option `tls` can also be used in `ejabberd_http` to support HTTPS.

Enabling this option implicitly disables the `starttls` option.

tls_compression

true | false

Whether to enable or disable TLS compression. The default value is `false`.

Please note: if this option is set in `ejabberd_c2s` or `ejabberd_s2s_in` and the corresponding top-level option is also set (`c2s_tls_compression`, `s2s_tls_compression`), then the top-level option is used, not this one.

tls_verify

false | true

This option specifies whether to verify the certificate or not when TLS is enabled.

The default value is `false`, which means no checks are performed.

The certificate will be checked against trusted CA roots, either defined at the operation system level or defined in the listener `cafile`. If trusted, it will accept the `jid` that is embedded in the certificate in the `subjectAltName` field of that certificate.

Enabling this option implicitly enables also the `starttls` option.

transport

tcp|udp

Defines the transport protocol. Default is `tcp`.

unix_socket

{mode|owner|group: Value}



added in [23.10](#)

Set the mode, owner and group of the unix domain socket defined in the [port](#) option.

The owner and group must be specified as integers, not as names.

Example:

```
listen:
-
  port: "unix:sockets/ctl_over_http.socket"
  unix_socket:
    mode: '0600'
    owner: 117
    group: 135
  module: ejabberd_http
  request_handlers:
    /ctl: ejabberd_ctl
  tag: "ctl_over_http"
```

use_proxy_protocol

true | false

Is this listener accessed by proxy service that is using proxy protocol for supplying real IP addresses to ejabberd server. You can read about this protocol in [Proxy protocol specification](#). The default value of this option is `false`.

zlib

true | false

This option specifies that Zlib stream compression (as defined in [XEP-0138](#)) is available on connections to the port.

Top-Level Options

Please note

This section describes top level options of ejabberd 26.09. If you are using an old ejabberd release, please refer to the corresponding archived version of this page in the [Archive](#).

The options that changed in this version are marked with .

access_rules

```
{AccessName: {allow|deny: ACLName|ACLDefinition}}
```

This option defines [Access Rules](#). Each access rule is assigned a name that can be referenced from other parts of the configuration file (mostly from `access` options of ejabberd modules). Each rule definition may contain arbitrary number of `allow` or `deny` sections, and each section may contain any number of ACL rules (see [acl](#) option). There are no access rules defined by default.

Example:

```
access_rules:
  configure:
    allow: admin
  something:
    deny: someone
    allow: all
  s2s_banned:
    deny: problematic_hosts
    deny: banned_forever
    deny:
      ip: 222.111.222.111/32
    deny:
      ip: 111.222.111.222/32
    allow: all
  xmlrpc_access:
    allow:
      user: peter@example.com
    allow:
      user: ivone@example.com
    allow:
      user: bot@example.com
      ip: 10.0.0.0/24
```

acl

```
{ACLName: {ACLType: ACLValue}}
```

This option defines [access control lists](#): named sets of rules which are used to match against different targets (such as a JID or an IP address). Every set of rules has name `ACLName`: it can be any string except `all` or `none` (those are predefined names for the rules that match all or nothing respectively). The name `ACLName` can be referenced from other parts of the configuration file, for

example in `access_rules` option. The rules of `ACLName` are represented by mapping `{ACLType: ACLValue}`. These can be one of the following:

- **ip:** `Network` The rule matches any IP address from the `Network`.
- **node_glob:** `Pattern` Same as `node_regexp`, but matching is performed on a specified `Pattern` according to the rules used by the Unix shell.
- **node_regexp:** `user_regexp@server_regexp` The rule matches any JID with node part matching regular expression `user_regexp` and server part matching regular expression `server_regexp`.
- **resource:** `Resource` The rule matches any JID with a resource `Resource`.
- **resource_glob:** `Pattern` Same as `resource_regexp`, but matching is performed on a specified `Pattern` according to the rules used by the Unix shell.
- **resource_regexp:** `Regexp` The rule matches any JID with a resource that matches regular expression `Regexp`.
- **server:** `Server` The rule matches any JID from server `Server`. The value of `Server` must be a valid hostname or an IP address.
- **server_glob:** `Pattern` Same as `server_regexp`, but matching is performed on a specified `Pattern` according to the rules used by the Unix shell.
- **server_regexp:** `Regexp` The rule matches any JID from the server that matches regular expression `Regexp`.
- **user:** `Username` If `Username` is in the form of "user@server", the rule matches a JID against this value. Otherwise, if `Username` is in the form of "user", the rule matches any JID that has `Username` in the node part as long as the server part of this JID is any virtual host served by ejabberd.
- **user_glob:** `Pattern` Same as `user_regexp`, but matching is performed on a specified `Pattern` according to the rules used by the Unix shell.
- **user_regexp:** `Regexp` If `Regexp` is in the form of "regexp@server", the rule matches any JID with node part matching regular expression "regexp" as long as the server part of this JID is equal to "server". If `Regexp` is in the form of "regexp", the rule matches any JID with node part matching regular expression "regexp" as long as the server part of this JID is any virtual host served by ejabberd.

acme

Options

ACME configuration, to automatically obtain SSL certificates for the domains served by ejabberd, which means that certificate requests and renewals are performed to some CA server (aka "ACME server") in a fully automated mode. The `options` are:

- **auto:** `true | false` Whether to automatically request certificates for all configured domains (that yet have no a certificate) on server start or configuration reload. The default is `true`.
- **ca_url:** `URL` The ACME directory URL used as an entry point for the ACME server. The default value is <https://acme-v02.api.letsencrypt.org/directory> - the directory URL of Let's Encrypt authority.
- **cert_type:** `rsa | ec` A type of a certificate key. Available values are `ec` and `rsa` for EC and RSA certificates respectively. It's better to have RSA certificates for the purpose of backward compatibility with legacy clients and servers, thus the default is `rsa`.
- **contact:** `[Contact, ...]` A list of contact addresses (typically emails) where an ACME server will send notifications when problems occur. The value of `Contact` must be in the form of "scheme:address" (e.g. "mailto:user@domain.tld"). The default is an empty list which means an ACME server will send no notices.

Example:

```
acme:
  ca_url: https://acme-v02.api.letsencrypt.org/directory
  contact:
    - mailto:admin@domain.tld
    - mailto:bot@domain.tld
  auto: true
  cert_type: rsa
```

allow_contrib_modules

true | false

Whether to allow installation of third-party modules or not. See [ejabberd-modules](#) documentation section. The default value is `true`.

allow_multiple_connections

true | false

This option is only used when the anonymous mode is enabled. Setting it to `true` means that the same username can be taken multiple times in anonymous login mode if different resource are used to connect. This option is only useful in very special occasions. The default value is `false`.

anonymous_protocol

login_anon | sasl_anon | both

Define what [anonymous](#) protocol will be used:

- `login_anon` means that the anonymous login method will be used.
- `sasl_anon` means that the SASL Anonymous method will be used.
- `both` means that SASL Anonymous and login anonymous are both enabled.

The default value is `sasl_anon`.

api_permissions

[Permission, ...]

Define the permissions for API access. Please consult the ejabberd Docs web → For Developers → ejabberd ReST API → [API Permissions](#).

append_host_config

{Host: Options}

Add a few specific options to a certain [virtual host](#).

auth_cache_life_time

timeout()

Same as [cache_life_time](#), but applied to authentication cache only. If not set, the value from [cache_life_time](#) will be used.

auth_cache_missed

true | false

Same as [cache_missed](#), but applied to authentication cache only. If not set, the value from [cache_missed](#) will be used.

auth_cache_size

pos_integer() | infinity

Same as [cache_size](#), but applied to authentication cache only. If not set, the value from [cache_size](#) will be used.

auth_external_user_exists_check

true | false | force

Note about this option: added in [23.10](#), `force` added in [26.09](#). Supplement check for user existence based on [mod_last](#) data, for authentication methods that don't have a way to reliably tell if a user exists (like is the case for `jwt` and certificate based authentication). This helps with processing offline message for those users. Using `force` values enables this processing even when no external storage method is enabled (like when using certificate authentications together with local stored passwords). The default value is `true`.

auth_method

[mnesia | sql | anonymous | external | jwt | ldap | pam, ...]

A list of [authentication](#) methods to use. If several methods are defined, authentication is considered successful as long as authentication of at least one of the methods succeeds. The default value is `[mnesia]`.

auth_opts

[Option, ...]

This is used by the contributed module `ejabberd_auth_http` that can be installed from the [ejabberd-contrib](#) Git repository. Please refer to that module's README file for details.

auth_password_format

plain | scram

 improved in [20.01](#)

The option defines in what format the users passwords are stored, plain text or in [SCRAM](#) format:

- `plain`: The password is stored as plain text in the database. This is risky because the passwords can be read if your database gets compromised. This is the default value. This format allows clients to authenticate using: the old Jabber Non-SASL (XEP-0078), SASL PLAIN, SASL DIGEST-MD5, and SASL SCRAM-SHA-1/256/512(-PLUS).
- `scram`: The password is not stored, only some information required to verify the hash provided by the client. It is impossible to obtain the original plain password from the stored information; for this reason, when this value is configured it cannot be changed to plain anymore. This format allows clients to authenticate using: SASL PLAIN and SASL SCRAM-SHA-1/256/512(-PLUS). The SCRAM variant depends on the [auth_scram_hash](#) option.

The default value is `plain`.

auth_password_types_hidden_in_sasl1

[plain | scram_sha1 | scram_sha256 | scram_sha512]

 added in [25.07](#)

List of password types that should not be offered in SASL1 authentication. Because SASL1, unlike SASL2, can't have list of available mechanisms tailored to individual user, it's possible that offered mechanisms will not be compatible with stored password, especially if new password type was added recently. This option allows disabling offering some mechanisms in SASL1, to a time until new password type will be available for all users.

auth_scram_hash

sha | sha256 | sha512

Hash algorithm that should be used to store password in [SCRAM](#) format. You shouldn't change this if you already have passwords generated with a different algorithm - users that have such passwords will not be able to authenticate. The default value is `sha`.

auth_stored_password_types

[plain | scram_sha1 | scram_sha256 | scram_sha512]

 added in [25.03](#)

List of password types that should be stored simultaneously for each user in database. When the user sets the account password, database will be updated to store the password in formats compatible with each type listed here. This can be used to migrate user passwords to a more secure format. If this option is set, it will override values set in [auth_scram_hash](#) and [auth_password_format](#) options. The default value is `[]`.

auth_use_cache

true | false

Same as [use_cache](#), but applied to authentication cache only. If not set, the value from [use_cache](#) will be used.

c2s_cafile

Path

Full path to a file containing one or more CA certificates in PEM format. All client certificates should be signed by one of these root CA certificates and should contain the corresponding JID(s) in `subjectAltName` field. There is no default value.

You can use [host_config](#) to specify this option per-vhost.

To set a specific file per listener, use the listener's [cafile](#) option. Please notice that `c2s_cafile` overrides the listener's `cafile` option.

c2s_ciphers

[Cipher, ...]

A list of OpenSSL ciphers to use for c2s connections. The default value is shown in the example below:

Example:

```
c2s_ciphers:
- HIGH
- "!aNULL"
- "!eNULL"
- "!3DES"
- "@STRENGTH"
```

c2s_dhfile

Path

Full path to a file containing custom DH parameters to use for c2s connections. Such a file could be created with the command `"openssl dhparam -out dh.pem 2048"`. If this option is not specified, 2048-bit MODP Group with 256-bit Prime Order Subgroup will be used as defined in RFC5114 Section 2.3.

c2s_protocol_options

[Option, ...]

List of general SSL options to use for c2s connections. These map to OpenSSL's `set_options()`. The default value is shown in the example below:

Example:

```
c2s_protocol_options:
- no_sslv3
- cipher_server_preference
- no_compression
```

c2s_tls_compression

true | false

Whether to enable or disable TLS compression for c2s connections. The default value is `false`.

ca_file

Path

Path to a file of CA root certificates. The default is to use system defined file if possible.

For server connections, this `ca_file` option is overridden by the `s2s_cafile` option.

cache_life_time

timeout()

The time of a cached item to keep in cache. Once it's expired, the corresponding item is erased from cache. The default value is `1 hour`. Several modules have a similar option; and some core ejabberd parts support similar options too, see [auth_cache_life_time](#), [oauth_cache_life_time](#), [router_cache_life_time](#), and [sm_cache_life_time](#).

cache_missed

true | false

Whether or not to cache missed lookups. When there is an attempt to lookup for a value in a database and this value is not found and the option is set to `true`, this attempt will be cached and no attempts will be performed until the cache expires (see [cache_life_time](#)). Usually you don't want to change it. Default is `true`. Several modules have a similar option; and some core ejabberd parts support similar options too, see [auth_cache_missed](#), [oauth_cache_missed](#), [router_cache_missed](#), and [sm_cache_missed](#).

cache_size

pos_integer() | infinity

A maximum number of items (not memory!) in cache. The rule of thumb, for all tables except rosters, you should set it to the number of maximum online users you expect. For roster multiply this number by 20 or so. If the cache size reaches this threshold, it's fully cleared, i.e. all items are deleted, and the corresponding warning is logged. You should avoid frequent cache clearance, because this degrades performance. The default value is `1000`. Several modules have a similar option; and some core ejabberd parts support similar options too, see [auth_cache_size](#), [oauth_cache_size](#), [router_cache_size](#), and [sm_cache_size](#).

captcha_cmd

Path | ModuleName

 improved in [23.01](#)

Full path to a script that generates **CAPTCHA** images. The keyword `@VERSION@` is replaced with ejabberd version number in `XX.YY` format. The keyword `@SEMVER@` is replaced with ejabberd version number in semver format when compiled with Elixir's mix, or `XX.YY` format otherwise. Alternatively, it can be the name of a module that implements ejabberd CAPTCHA support. There is no default value: when this option is not set, CAPTCHA functionality is completely disabled.

Examples:

When using the ejabberd installers or container image, the example captcha scripts can be used like this:

```
captcha_cmd: /opt/ejabberd-@VERSION@/Lib/ejabberd-@SEMVER@/priv/bin/captcha.sh
```

captcha_host

String

Deprecated. Use `captcha_url` instead.

captcha_limit

pos_integer() | infinity

Maximum number of **CAPTCHA** generated images per minute for any given JID. The option is intended to protect the server from CAPTCHA DoS. The default value is `infinity`.

captcha_pow

pos_integer() | false

Offer an XEP-0158 SHA-256 hashcash (proof-of-work) challenge in the **CAPTCHA** form, alongside the image challenge or on its own. The value is the difficulty: the number of hexadecimal digits of the SHA-256 digest a client must match. Unlike the image challenge, this does not require `captcha_cmd`. The default value is `false`.

captcha_url

URL | auto | undefined

 improved in 23.04

An URL where **CAPTCHA** requests should be sent. NOTE: you need to configure `request_handlers` for `ejabberd_http` listener as well. If set to `auto`, it builds the URL using a `request_handler` already enabled, with encryption if available. If set to `undefined`, it builds the URL using the deprecated `captcha_host` + `/captcha`. The default value is `auto`.

certfiles

[Path, ...]

The option accepts a list of file paths (optionally with wildcards) containing either PEM certificates or PEM private keys. At startup or configuration reload, ejabberd reads all certificates from these files, sorts them, removes duplicates, finds matching private keys and then rebuilds full certificate chains for the use in TLS connections. Use this option when TLS is enabled in either of ejabberd listeners: `ejabberd_c2s`, `ejabberd_http` and so on. NOTE: if you modify the certificate files or change the value of the option, run `ejabberdctl reload-config` in order to rebuild and reload the certificate chains.

Examples:

If you use **Let's Encrypt** certificates for your domain "domain.tld", the configuration will look like this:

```
certfiles:
  - /etc/letsencrypt/live/domain.tld/fullchain.pem
  - /etc/letsencrypt/live/domain.tld/privkey.pem
```

cluster_backend

Backend

A database backend to use for storing information about cluster. The only available value so far is `mnesia`.

cluster_nodes

[Node, ...]

A list of Erlang nodes to connect on ejabberd startup. This option is mostly intended for ejabberd customization and sophisticated setups. The default value is an empty list.

default_db

mnesia | sql

[Default database](#) to store persistent data in ejabberd. Some components can be configured with specific toplevel options like [oauth_db_type](#). Many modules can be configured with specific module options, usually named `db_type`. The default value is `mnesia`.

default_ram_db

mnesia | redis | sql

Default volatile (in-memory) storage for ejabberd. Some components can be configured with specific toplevel options like [router_db_type](#) and [sm_db_type](#). Some modules can be configured with specific module options, usually named `ram_db_type`. The default value is `mnesia`.

define_keyword

{NAME: Value}

 added in 25.03

Allows to define configuration [keywords](#).

Example:

```
define_keyword:
  SQL_USERNAME: "eja.global"

host_config:
  localhost:
    define_keyword:
      SQL_USERNAME: "eja.localhost"

sql_username: "prefix.@SQL_USERNAME@"
```

define_macro

{NAME: Value}

 improved in 25.03

Allows to define configuration [macros](#).

Example:

```
define_macro:
  DEBUG: debug
  LOG_LEVEL: DEBUG
  USERBOB:
    user: bob@localhost
```

```
loglevel: LOG_LEVEL
acl:
  admin: USERBOB
```

disable_sasl_mechanisms

[Mechanism, ...]

Specify a list of SASL mechanisms (such as `DIGEST-MD5` or `SCRAM-SHA1`) that should not be offered to the client. For convenience, the value of `Mechanism` is case-insensitive. The default value is an empty list, i.e. no mechanisms are disabled by default.

disable_sasl_scram_downgrade_protection

true | false

Allows to disable sending data required by *XEP-0474: SASL SCRAM Downgrade Protection*. There are known buggy clients (like those that use strophejs 1.6.2) which will not be able to authenticate when servers sends data from that specification. This options allows server to disable it to allow even buggy clients connects, but in exchange decrease MITM protection. The default value of this option is `false` which enables this extension.

domain_balancing

{Domain: Options}

An algorithm to [load-balance](#) the components that are plugged on an ejabberd cluster. It means that you can plug one or several instances of the same component on each ejabberd node and that the traffic will be automatically distributed. The algorithm to deliver messages to the component(s) can be specified by this option. For any component connected as `Domain`, available `Options` are:

- **component_number:** `2..1000` The number of components to balance.
- **type:** `Value` How to deliver stanzas to connected components. The default value is `random`. Possible values:

- bare_destination

by the bare JID (without resource) of the packet's `to` attribute

- bare_source

by the bare JID (without resource) of the packet's `from` attribute is used

- destination

an instance is chosen by the full JID of the packet's `to` attribute

- random

an instance is chosen at random

- source

by the full JID of the packet's `from` attribute

Example:

```
domain_balancing:
  component.domain.tld:
    type: destination
    component_number: 5
  transport.example.org:
    type: bare_source
```

ext_api_headers

Headers

String of headers (separated with commas `,`) that will be provided by ejabberd when sending ReST requests. The default value is an empty string of headers: `""`.

ext_api_http_pool_size

`pos_integer()`

Define the size of the HTTP pool, that is, the maximum number of sessions that the ejabberd ReST service will handle simultaneously. The default value is: `100`.

ext_api_path_oauth

Path

Define the base URI path when performing OAUTH ReST requests. The default value is: `"/oauth"`.

ext_api_url

URL

Define the base URI when performing ReST requests. The default value is: `"http://localhost/api"`.

extauth_pool_name

Name

Define the pool name appendix in [external auth](#), so the full pool name will be `extauth_pool_Name`. The default value is the hostname.

extauth_pool_size

Size

The option defines the number of instances of the same [external auth](#) program to start for better load balancing. The default is the number of available CPU cores.

extauth_program

Path

Indicate in this option the full path to the [external authentication script](#). The script must be executable by ejabberd.

fqdn

Domain

A fully qualified domain name that will be used in SASL DIGEST-MD5 authentication. The default is detected automatically.

hide_sensitive_log_data

`true | false`

A privacy option to not log sensitive data (mostly IP addresses). The default value is `false` for backward compatibility.

host_config

`{Host: Options}`

The option is used to redefine `options` for **virtual host** `Host`. In the example below LDAP authentication method will be used on virtual host `domain.tld` and SQL method will be used on virtual host `example.org`.

Example:

```
hosts:
- domain.tld
- example.org

auth_method:
- sql

host_config:
domain.tld:
  auth_method:
  - ldap
```

hosts

`[Domain1, Domain2, ...]`

List of one or more **host names** (or domains) that ejabberd will serve. This is a **mandatory** option.

hosts_alias

`{Alias: Host}`

 added in 25.07

Define aliases for existing vhosts managed by ejabberd. An alias may be a regexp expression. This option is only consulted by the `ejabberd_http` listener.

Example:

```
hosts:
- domain.tld
- example.org

hosts_alias:
xmpp.domain.tld: domain.tld
jabber.domain.tld: domain.tld
mytest.net: example.org
"exa*": example.org
```

include_config_file

`[Filename, ...] | {Filename: Options}`

Read and **include additional file** from `Filename`. If the value is provided in `{Filename: Options}` format, the `options` must be one of the following:

- **allow_only:** `[OptionName, ...]` Allows only the usage of those options in the included file `Filename`. The options that do not match this criteria are not accepted. The default value is to include all options.
- **disallow:** `[OptionName, ...]` Disallows the usage of those options in the included file `Filename`. The options that match this criteria are not accepted. The default value is an empty list.

install_contrib_modules

`[Module, ...]`

 added in 23.10

Modules from `ejabberd-modules` to install automatically at start time. The default value is an empty list of modules: `[]`.

Example:

```
install_contrib_modules:
- mod_tombstones

modules:
  mod_tombstones: {}
```

jwt_auth_only_rule

AccessName

This ACL rule defines accounts that can use only the [JWT](#) auth method, even if others are also defined in the ejabberd configuration file. In other words: if there are several auth methods enabled for this host (JWT, SQL, ...), users that match this rule can only use JWT. The default value is `none`.

jwt_jid_field

FieldName

By default, the JID is defined in the `"jid"` JWT field. In this option you can specify other [JWT](#) field name where the JID is defined.

jwt_key

FilePath

Path to the file that contains the [JWT](#) key.

Supported file formats:

- **PEM format** - Standard PEM-encoded keys (RSA, EC, EdDSA, etc.)
- **JWK (JSON Web Key)** - JSON format
- **JWK Set** - JSON with a `"keys"` array (but must contain exactly one key)

By default there is no path defined, that is: `undefined`.

language

Language

Define the [default language](#) of server strings that can be seen by XMPP clients. If an XMPP client does not possess `xml:lang` attribute, the specified language is used. The default value is `"en"`.

ldap_backups

[Host, ...]

A list of IP addresses or DNS names of LDAP backup servers (see [LDAP connection](#)). When no servers listed in [ldap_servers](#) option are reachable, ejabberd connects to these backup servers. The default is an empty list, i.e. no backup servers specified. Please notice that ejabberd only connects to the next server when the existing connection is lost; it doesn't detect when a previously-attempted server becomes available again.

ldap_base

Base

LDAP base directory which stores users accounts. There is no default value: you must set the option in order for LDAP connections to work properly.

ldap_deref_aliases

never | always | finding | searching

Whether to dereference aliases or not. The default value is `never`.

ldap_dn_filter

{Filter: FilterAttrs}

This filter is applied on the results returned by the main filter. The filter performs an additional LDAP lookup to make the complete result. This is useful when you are unable to define all filter rules in `ldap_filter`. You can define `"%u"`, `"%d"`, `"%s"` and `"%D"` pattern variables in *Filter*: `"%u"` is replaced by a user's part of the JID, `"%d"` is replaced by the corresponding domain (virtual host), all `"%s"` variables are consecutively replaced by values from the attributes in `FilterAttrs` and `"%D"` is replaced by Distinguished Name from the result set. There is no default value, which means the result is not filtered. **WARNING:** Since this filter makes additional LDAP lookups, use it only as the last resort: try to define all filter rules in `ldap_filter` option if possible.

Example:

```
ldap_dn_filter:
"(&(name=%s)(owner=%D)(user=%u@%d))": [sn]
```

ldap_encrypt

tls | none

Whether to encrypt LDAP connection using TLS or not. The default value is `none`. NOTE: STARTTLS encryption is not supported.

ldap_filter

Filter

An LDAP filter as defined in [RFC4515](#). There is no default value. Example: `"(&(objectClass=shadowAccount)(memberOf=XMPP Users))"`. NOTE: don't forget to close brackets and don't use superfluous whitespaces. Also you must not use `"uid"` attribute in the filter because this attribute will be appended to the filter automatically.

ldap_password

Password

Bind password. The default value is an empty string.

ldap_port

1..65535

Port to connect to your LDAP server. The default port is `389` if encryption is disabled and `636` if encryption is enabled.

ldap_rootdn

RootDN

Bind Distinguished Name. The default value is an empty string, which means "anonymous connection".

ldap_servers

[Host, ...]

A list of IP addresses or DNS names of your LDAP servers (see [LDAP connection](#)). ejabberd connects immediately to all of them, and reconnects infinitely if connection is lost. The default value is `[localhost]`.

ldap_tls_cacertfile

Path

A path to a file containing PEM encoded CA certificates. This option is required when TLS verification is enabled.

ldap_tls_certfile

Path

A path to a file containing PEM encoded certificate along with PEM encoded private key. This certificate will be provided by ejabberd when TLS enabled for LDAP connections. There is no default value, which means no client certificate will be sent.

ldap_tls_depth

Number

Specifies the maximum verification depth when TLS verification is enabled, i.e. how far in a chain of certificates the verification process can proceed before the verification is considered to be failed. Peer certificate = 0, CA certificate = 1, higher level CA certificate = 2, etc. The value `2` thus means that a chain can at most contain peer cert, CA cert, next CA cert, and an additional CA cert. The default value is `1`.

ldap_tls_verify

false | soft | hard

This option specifies whether to verify LDAP server certificate or not when TLS is enabled. When `hard` is set, ejabberd doesn't proceed if the certificate is invalid. When `soft` is set, ejabberd proceeds even if the check has failed. The default is `false`, which means no checks are performed.

ldap_uids

`[Attr] | {Attr: AttrFormat}`

LDAP attributes which hold a list of attributes to use as alternatives for getting the JID, where `Attr` is an LDAP attribute which holds the user's part of the JID and `AttrFormat` must contain one and only one pattern variable `"%u"` which will be replaced by the user's part of the JID. For example, `"%u@example.org"`. If the value is in the form of `[Attr]` then `AttrFormat` is assumed to be `"%u"`.

listen

`[Options, ...]`

The option for listeners configuration. See the [Listen Modules](#) section for details.

log_burst_limit_count

Number

 added in [22.10](#)

The number of messages to accept in `log_burst_limit_window_time` period before starting to drop them. Default `500`

log_burst_limit_window_time

Number

 added in [22.10](#)

The time period to rate-limit log messages by. Defaults to `1` second.

log_modules_fully

[Module, ...]

 added in [23.01](#)

List of modules that will log everything independently from the general loglevel option.

log_rotate_count

Number

The number of rotated log files to keep. The default value is `1`, which means that only keeps `ejabberd.log.0`, `error.log.0` and `crash.log.0`.

log_rotate_size

pos_integer() | infinity

The size (in bytes) of a log file to trigger rotation. If set to `infinity`, log rotation is disabled. The default value is 10 Mb expressed in bytes: `10485760`.

loglevel

none | emergency | alert | critical | error | warning | notice | info | debug

Verbosity of ejabberd [logging](#). The default value is `info`. NOTE: previous versions of ejabberd had log levels defined in numeric format (`0..5`). The numeric values are still accepted for backward compatibility, but are not recommended.

max_fsm_queue

Size

This option specifies the maximum number of elements in the queue of the FSM (Finite State Machine). Roughly speaking, each message in such queues represents one XML stanza queued to be sent into its relevant outgoing stream. If queue size reaches the limit (because, for example, the receiver of stanzas is too slow), the FSM and the corresponding connection (if any) will be terminated and error message will be logged. The reasonable value for this option depends on your hardware configuration. The allowed values are positive integers. The default value is `10000`.

modules

{Module: Options}

Set all the [modules](#) configuration options.

negotiation_timeout

timeout()

Time to wait for an XMPP stream negotiation to complete. When timeout occurs, the corresponding XMPP stream is closed. The default value is `120` seconds.

net_ticktime

`timeout()`

This option can be used to tune tick time parameter of `net_kernel`. It tells Erlang VM how often nodes should check if intra-node communication was not interrupted. This option must have identical value on all nodes, or it will lead to subtle bugs. Usually leaving default value of this option is best, tweak it only if you know what you are doing. The default value is `1 minute`.

new_sql_schema

`true | false`

 `obsolete` in [25.10](#)

This option was renamed to [sql_schema_multihost](#) in ejabberd [25.10](#). Please update your configuration to use the new option name

oauth_access

`AccessName`

By default creating OAuth tokens is not allowed. To define which users can create OAuth tokens, you can refer to an ejabberd access rule in the `oauth_access` option. Use `all` to allow everyone to create tokens.

oauth_cache_life_time

`timeout()`

Same as [cache_life_time](#), but applied to OAuth cache only. If not set, the value from [cache_life_time](#) will be used.

oauth_cache_missed

`true | false`

Same as [cache_missed](#), but applied to OAuth cache only. If not set, the value from [cache_missed](#) will be used.

oauth_cache_rest_failure_life_time

`timeout()`

 `added` in [21.01](#)

The time that a failure in OAuth ReST is cached. The default value is `infinity`.

oauth_cache_size

`pos_integer() | infinity`

Same as [cache_size](#), but applied to OAuth cache only. If not set, the value from [cache_size](#) will be used.

oauth_client_id_check

`allow | db | deny`

Define whether the client authentication is always allowed, denied, or it will depend if the client ID is present in the database. The default value is `db`.

oauth_db_type

`mnesia` | `sql`

Database backend to use for OAuth authentication. The default value is picked from `default_db` option, or if it's not set, `mnesia` will be used.

oauth_expire

`timeout()`

Time during which the OAuth token is valid, in seconds. After that amount of time, the token expires and the delegated credential cannot be used and is removed from the database. The default is `4294967` seconds.

oauth_use_cache

`true` | `false`

Same as `use_cache`, but applied to OAuth cache only. If not set, the value from `use_cache` will be used.

oom_killer

`true` | `false`

Enable or disable OOM (out-of-memory) killer. When system memory raises above the limit defined in `oom_watermark` option, ejabberd triggers OOM killer to terminate most memory consuming Erlang processes. Note that in order to maintain functionality, ejabberd only attempts to kill transient processes, such as those managing client sessions, s2s or database connections. The default value is `true`.

oom_queue

`Size`

Trigger OOM killer when some of the running Erlang processes have messages queue above this `Size`. Note that such processes won't be killed if `oom_killer` option is set to `false` or if `oom_watermark` is not reached yet.

oom_watermark

`Percent`

A percent of total system memory consumed at which OOM killer should be activated with some of the processes possibly be killed (see `oom_killer` option). Later, when memory drops below this `Percent`, OOM killer is deactivated. The default value is `80` percents.

outgoing_s2s_families

`[ipv6 | ipv4, ...]`

 changed in 23.01

Specify which address families to try, in what order. The default is `[ipv6, ipv4]` which means it first tries connecting with IPv6, if that fails it tries using IPv4. This option is obsolete and irrelevant when using ejabberd 23.01 and Erlang/OTP 22, or newer versions of them.

outgoing_s2s_ipv4_address

Address

 added in 20.12

Specify the IPv4 address that will be used when establishing an outgoing S2S IPv4 connection, for example `"127.0.0.1"`. The default value is `undefined`.

outgoing_s2s_ipv6_address

Address

 added in 20.12

Specify the IPv6 address that will be used when establishing an outgoing S2S IPv6 connection, for example `"::FFFF:127.0.0.1"`. The default value is `undefined`.

outgoing_s2s_port

1..65535

A port number to use for outgoing s2s connections when the target server doesn't have an SRV record. The default value is `5269`.

outgoing_s2s_timeout

timeout()

The timeout in seconds for outgoing S2S connection attempts. The default value is `10` seconds.

pam_service

Name

This option defines the [PAM](#) service name. Refer to the PAM documentation of your operation system for more information. The default value is `ejabberd`.

pam_userinfotype

username | jid

This option defines what type of information about the user ejabberd provides to the [PAM](#) service: only the username, or the user's JID. Default is `username`.

pgsql_users_number_estimate

true | false

Whether to use PostgreSQL estimation when counting registered users. The default value is `false`.

queue_dir

Directory

If [queue_type](#) option is set to `file`, use this `Directory` to store file queues. The default is to keep queues inside Mnesia directory.

queue_type

ram | file

Default type of queues in ejabberd. Modules may have its own value of the option. The value of `ram` means that queues will be kept in memory. If value `file` is set, you may also specify directory in `queue_dir` option where file queues will be placed. The default value is `ram`.

redis_connect_timeout

timeout()

A timeout to wait for the connection to be re-established to the [Redis](#) server. The default is `1 second`.

redis_db

Number

[Redis](#) database number. The default is `0`.

redis_password

Password

The password to the [Redis](#) server. The default is an empty string, i.e. no password.

redis_pool_size

Number

The number of simultaneous connections to the [Redis](#) server. The default value is `10`.

redis_port

1..65535

The port where the [Redis](#) server is accepting connections. The default is `6379`.

redis_queue_type

ram | file

The type of request queue for the [Redis](#) server. See description of `queue_type` option for the explanation. The default value is the value defined in `queue_type` or `ram` if the latter is not set.

redis_server

Host | IP Address | Unix Socket Path

 improved in [24.12](#)

A hostname, IP address or unix domain socket file of the [Redis](#) server. Setup the path to unix domain socket like: `"unix:/path/to/socket"`. The default value is `localhost`.

registration_timeout

timeout()

This is a global option for module [mod_register](#). It limits the frequency of registrations from a given IP or username. So, a user that tries to register a new account from the same IP address or JID during this time after their previous registration will receive an error with the corresponding explanation. To disable this limitation, set the value to `infinity`. The default value is `600 seconds`.

replaced_connection_timeout

`timeout()`

 added in [26.01](#)

Maximum time that new session will wait for termination of session that it's replacing. This allows old session to properly send its unavailable presences, and helps with potential race conditions between old and new sessions presences.

resource_conflict

`setresource` | `closeold` | `closenew`

NOTE: this option is deprecated and may be removed anytime in the future versions. The possible values match exactly the three possibilities described in [XMPP Core: section 7.7.2.2](#). The default value is `closeold`. If the client uses old Jabber Non-SASL authentication (XEP-0078), then this option is not respected, and the action performed is `closeold`.

rest_proxy

`Host`

 added in [25.07](#)

Address of a HTTP Connect proxy used by modules issuing rest calls (like `ejabberd_oauth_rest`)

rest_proxy_password

`string()`

 added in [25.07](#)

Password used to authenticate to HTTP Connect proxy used by modules issuing rest calls (like `ejabberd_oauth_rest`)

rest_proxy_port

`1..65535`

 added in [25.07](#)

Port of a HTTP Connect proxy used by modules issuing rest calls (like `ejabberd_oauth_rest`)

rest_proxy_username

`string()`

 added in [25.07](#)

Username used to authenticate to HTTP Connect proxy used by modules issuing rest calls (like `ejabberd_oauth_rest`)

router_cache_life_time

`timeout()`

Same as [cache_life_time](#), but applied to routing table cache only. If not set, the value from [cache_life_time](#) will be used.

router_cache_missed

`true | false`

Same as [cache_missed](#), but applied to routing table cache only. If not set, the value from [cache_missed](#) will be used.

router_cache_size

`pos_integer() | infinity`

Same as [cache_size](#), but applied to routing table cache only. If not set, the value from [cache_size](#) will be used.

router_db_type

`mnesia | redis | sql`

Database backend to use for routing information. The default value is picked from [default_ram_db](#) option, or if it's not set, `mnesia` will be used.

router_use_cache

`true | false`

Same as [use_cache](#), but applied to routing table cache only. If not set, the value from [use_cache](#) will be used.

rpc_timeout

`timeout()`

A timeout for remote function calls between nodes in an ejabberd cluster. You should probably never change this value since those calls are used for internal needs only. The default value is `5` seconds.

s2s_access

`Access`

This [Access Rule](#) defines to what remote servers can s2s connections be established. The default value is `all`; no restrictions are applied, it is allowed to connect s2s to/from all remote servers.

s2s_cafile

`Path`

A path to a file with CA root certificates that will be used to authenticate s2s connections. If not set, the value of [ca_file](#) will be used.

You can use [host_config](#) to specify this option per-vhost.

s2s_ciphers

`[Cipher, ...]`

A list of OpenSSL ciphers to use for s2s connections. The default value is shown in the example below:

Example:

```
s2s_ciphers:
- HIGH
- "!aNULL"
- "!eNULL"
- "!3DES"
- "@STRENGTH"
```

s2s_dhfile

Path

Full path to a file containing custom DH parameters to use for s2s connections. Such a file could be created with the command `"openssl dhparam -out dh.pem 2048"`. If this option is not specified, 2048-bit MODP Group with 256-bit Prime Order Subgroup will be used as defined in RFC5114 Section 2.3.

s2s_dns_retries

Number

DNS resolving retries. The default value is `2`.

s2s_dns_timeout

timeout()

The timeout for DNS resolving. The default value is `10` seconds.

s2s_max_retry_delay

timeout()

The maximum allowed delay for s2s connection retry to connect after a failed connection attempt. The default value is `300` seconds (5 minutes).

s2s_protocol_options

[Option, ...]

List of general SSL options to use for s2s connections. These map to OpenSSL's `set_options()`. The default value is shown in the example below:

Example:

```
s2s_protocol_options:
- no_sslv3
- cipher_server_preference
- no_compression
```

s2s_queue_type

ram | file

The type of a queue for s2s packets. See description of [queue_type](#) option for the explanation. The default value is the value defined in [queue_type](#) or `ram` if the latter is not set.

s2s_timeout

timeout()

A time to wait before closing an idle s2s connection. The default value is `1` hour.

s2s_tls_compression

true | false

Whether to enable or disable TLS compression for s2s connections. The default value is `false`.

s2s_use_starttls

true | false | optional | required

Whether to use STARTTLS for s2s connections. The value of `false` means STARTTLS is prohibited. The value of `true` or `optional` means STARTTLS is enabled but plain connections are still allowed. And the value of `required` means that only STARTTLS connections are allowed. The default value is `false` (for historical reasons).

s2s_zlib

true | false

Whether to use `zlib` compression (as defined in [XEP-0138](#)) or not. The default value is `false`. WARNING: this type of compression is nowadays considered insecure.

shaper

{ShaperName: Rate}

The option defines a set of [shapers](#). Every shaper is assigned a name `ShaperName` that can be used in other parts of the configuration file, such as [shaper_rules](#) option. The shaper itself is defined by its `Rate`, where `Rate` stands for the maximum allowed incoming rate in **bytes** per second. When a connection exceeds this limit, ejabberd stops reading from the socket until the average rate is again below the allowed maximum. In the example below shaper `normal` limits the traffic speed to 1,000 bytes/sec and shaper `fast` limits the traffic speed to 50,000 bytes/sec:

Example:

```
shaper:
  normal: 1000
  fast: 50000
```

shaper_rules

{ShaperRuleName: {Number|ShaperName: ACLName|ACLDefinition}}

This option defines [shaper rules](#) to use for matching user/hosts. Semantics is similar to [access_rules](#) option, the only difference is that instead using `allow` or `deny`, a name of a shaper (defined in [shaper](#) option) or a positive number should be used.

Example:

```
shaper_rules:
  connections_limit:
    10:
      user: peter@example.com
    100: admin
    5: all
  download_speed:
    fast: admin
    slow: anonymous_users
    normal: all
  log_days: 30
```

sm_cache_life_time

timeout()

Same as [cache_life_time](#), but applied to client sessions table cache only. If not set, the value from [cache_life_time](#) will be used.

sm_cache_missed

true | false

Same as [cache_missed](#), but applied to client sessions table cache only. If not set, the value from [cache_missed](#) will be used.

sm_cache_size

pos_integer() | infinity

Same as [cache_size](#), but applied to client sessions table cache only. If not set, the value from [cache_size](#) will be used.

sm_db_type

mnesia | redis | sql

Database backend to use for client sessions information. The default value is picked from [default_ram_db](#) option, or if it's not set, [mnesia](#) will be used.

sm_use_cache

true | false

Same as [use_cache](#), but applied to client sessions table cache only. If not set, the value from [use_cache](#) will be used.

sql_connect_timeout

timeout()

A time to wait for connection to an SQL server to be established. The default value is `5` seconds.

sql_database

DatabaseName | PathSqliteFile

 improved in [26.04](#)

An SQL database name, being the default value `ejabberd`. If option [sql_type](#) is set to `sqlite`, this must be a full path to a database file, being the default value: `"@DATABASE_PATH@/sqlite/@HOST@.sqlite"`

sql_flags

[mysql_alternative_upsert]

 added in [24.02](#)

This option accepts a list of SQL flags, and is empty by default. `mysql_alternative_upsert` forces the alternative upsert implementation in MySQL.

sql_keepalive_interval

timeout()

An interval to make a dummy SQL request to keep alive the connections to the database. There is no default value, so no keepalive requests are made.

sql_odbc_driver

Path

 added in 20.12

Path to the ODBC driver to use to connect to a Microsoft SQL Server database. This option only applies if the [sql_type](#) option is set to `mssql` and [sql_server](#) is not an ODBC connection string. The default value is: `libtdsodbc.so`

sql_password

Password

The password for SQL authentication. The default is empty string.

sql_pool_size

Size

Number of connections to the SQL server that ejabberd will open for each virtual host. The default value is `10`. WARNING: for SQLite this value is `1` by default and it's not recommended to change it due to potential race conditions.

sql_port

`1..65535`

The port where the SQL server is accepting connections. The default is `3306` for MySQL, `5432` for PostgreSQL and `1433` for MS SQL. The option has no effect for SQLite.

sql_prepared_statements

`true` | `false`

 added in 20.01

This option is `true` by default, and is useful to disable prepared statements. The option is valid for PostgreSQL and MySQL.

sql_query_timeout

`timeout()`

A time to wait for an SQL query response. The default value is `60` seconds.

sql_queue_type

`ram` | `file`

The type of a request queue for the SQL server. See description of [queue_type](#) option for the explanation. The default value is the value defined in [queue_type](#) or `ram` if the latter is not set.

sql_schema_multihost

`true` | `false`

 renamed in 25.10

Whether to use the [multihost SQL schema](#). All schemas are located at <https://github.com/processone/ejabberd/tree/26.09/sql>. There are two schemas available. The legacy `singlehost` schema stores one XMPP domain into one ejabberd database. The

`multihost` schema can handle several XMPP domains in a single ejabberd database. The `multihost` schema is preferable when serving several XMPP domains and/or changing domains from time to time. This avoid need to manage several databases and handle complex configuration changes. The default depends on `.configure` flag `--enable-sql-schema-multihost` which is set at compile time.

sql_server

Host | IP Address | ODBC Connection String | Unix Socket Path

 improved in 24.06

The hostname or IP address of the SQL server. For `sql_type` `mssql` or `odbc` this can also be an ODBC connection string. When `sql_type` is `mysql` or `pgsql`, this can be the path to a unix domain socket expressed like: `"unix:/path/to/socket"`. The default value is `localhost`.

sql_ssl

true | false

 improved in 20.03

Whether to use SSL encrypted connections to the SQL server. The option is only available for MySQL, MS SQL and PostgreSQL. The default value is `false`.

sql_ssl_cafile

Path

A path to a file with CA root certificates that will be used to verify SQL connections. Implies `sql_ssl` and `sql_ssl_verify` options are set to `true`. There is no default which means certificate verification is disabled. This option has no effect for MS SQL.

sql_ssl_certfile

Path

A path to a certificate file that will be used for SSL connections to the SQL server. Implies `sql_ssl` option is set to `true`. There is no default which means ejabberd won't provide a client certificate to the SQL server. This option has no effect for MS SQL.

sql_ssl_verify

true | false

Whether to verify SSL connection to the SQL server against CA root certificates defined in `sql_ssl_cafile` option. Implies `sql_ssl` option is set to `true`. This option has no effect for MS SQL. The default value is `false`.

sql_start_interval

timeout()

A time to wait before retrying to restore failed SQL connection. The default value is `30` seconds.

sql_type

mssql | mysql | odbc | pgsql | sqlite

The type of an SQL connection. The default is `odbc`.

sql_username

Username

A user name for SQL authentication. The default value is `ejabberd`.

trusted_proxies

`all` | `[Network1, Network2, ...]`

Specify what proxies are trusted when an HTTP request contains the header `X-Forwarded-For`. You can specify `all` to allow all proxies, or specify a list of IPs, possibly with masks. The default value is an empty list. Using this option you can know the real IP of the request, for admin purpose, or security configuration (for example using [mod_fail2ban](#)). IMPORTANT: The proxy MUST be configured to set the `X-Forwarded-For` header if you enable this option as, otherwise, the client can set it itself and as a result the IP value cannot be trusted for security rules in ejabberd.

update_sql_schema

`true` | `false`

 updated in [24.06](#)

Allow ejabberd to update SQL schema in MySQL, PostgreSQL and SQLite databases. This option was added in ejabberd [23.10](#), and enabled by default since [24.06](#). The default value is `true`.

update_sql_schema_timeout

`timeout()`

 added in [24.07](#)

Time allocated to SQL schema update queries. The default value is set to 5 minutes.

use_cache

`true` | `false`

Enable or disable cache. The default is `true`. Several modules have a similar option; and some core ejabberd parts support similar options too, see [auth_use_cache](#), [oauth_use_cache](#), [router_use_cache](#), and [sm_use_cache](#).

validate_stream

`true` | `false`

Whether to validate any incoming XML packet according to the schemas of [supported XMPP extensions](#). WARNING: the validation is only intended for the use by client developers - don't enable it in production environment. The default value is `false`.

version

`string()`

The option can be used to set custom ejabberd version, that will be used by different parts of ejabberd, for example by [mod_version](#) module. The default value is obtained at compile time from the underlying version control system.

websocket_origin

`ignore` | `URL`

This option enables validation for `origin` header to protect against connections from other domains than given in the configuration file. In this way, the lower layer load balancer can be chosen for a specific ejabberd implementation while still providing a secure WebSocket connection. The default value is `ignore`. An example value of the `URL` is `"https://test.example.org:8081"`.

websocket_ping_interval

`timeout()`

Defines time between pings sent by the server to a client (WebSocket level protocol pings are used for this) to keep a connection active. If the client doesn't respond to two consecutive pings, the connection will be assumed as closed. The value of `0` can be used to disable the feature. This option makes the server sending pings only for connections using the RFC compliant protocol. For older style connections the server expects that whitespace pings would be used for this purpose. The default value is `60` seconds.

websocket_timeout

`timeout()`

Amount of time without any communication after which the connection would be closed. The default value is `300` seconds.

Modules Options

Please note

This section describes modules options of ejabberd [26.09](#). If you are using an old ejabberd release, please refer to the corresponding archived version of this page in the [Archive](#).

The modules that changed in this version are marked with .

mod_adhoc

ad-hoc command

Command that can be executed by an XMPP client using XEP-0050.

This module implements [XEP-0050: Ad-Hoc Commands](#). It's an auxiliary module and is only needed by some of the other modules.

Available options:

- **report_commands_node:** `true` | `false`

Provide the Commands item in the Service Discovery. Default value: `false`.

mod_adhoc_api

 added in [25.03](#)

Execute [API commands](#) in a XMPP client using [XEP-0050: Ad-Hoc Commands](#). This module requires [mod_adhoc](#) (to execute the commands), and recommends [mod_disco](#) (to discover the commands).

Available options:

- **default_version:** `integer()` | `string()`

What API version to use. If setting an ejabberd version, it will use the latest API version that was available in that [c2s](#) ejabberd version. For example, setting `"24.06"` in this option implies `2`. The default value is the latest version.

Example:

```
acl:
  admin:
    user: jan@localhost

api_permissions:
  "adhoc commands":
    from: mod_adhoc_api
    who: admin
    what:
      - "[tag:roster]"
      - "[tag:session]"
      - stats
      - status

modules:
  mod_adhoc_api:
    default_version: 2
```

mod_admin_extra

This module provides additional administrative commands.

Details for some commands:

[ban_account](#) API: This command kicks all the connected sessions of the account from the server. It also changes their password to a randomly generated one, so they can't login anymore unless a server administrator changes their password again. It is possible to define the reason of the ban. The new password also includes the reason and the date and time of the ban. See an example below.

[push_roster](#) API (and [push_roster_all](#) API): The roster file must be placed, if using Windows, on the directory where you installed ejabberd: `C:/Program Files/ejabberd` or similar. If you use other Operating System, place the file on the same directory where the .beam files are installed. See below an example roster file.

[srg_create](#) API: If you want to put a group Name with blank spaces, use the characters `" "` and `' '` to define when the Name starts and ends. See an example below.

The module has no options.

Examples:

With this configuration, vCards can only be modified with `mod_admin_extra` commands:

```
acl:
  adminextraresource:
    - resource: "modadminextraf8x,31ad"
access_rules:
  vcard_set:
    - allow: adminextraresource
modules:
  mod_admin_extra: {}
  mod_vcard:
    access_set: vcard_set
```

Content of roster file for [push_roster](#) API:

```
[{<<"bob">>, <<"example.org">>, <<"workers">>, <<"Bob">>},
{<<"mart">>, <<"example.org">>, <<"workers">>, <<"Mart">>},
{<<"Rich">>, <<"example.org">>, <<"bosses">>, <<"Rich">>}].
```

With this call, the sessions of the local account which JID is `boby@example.org` will be kicked, and its password will be set to something like `BANNED_ACCOUNT-20080425T21:45:07-2176635-Spammed_rooms`

```
ejabberdctl vhost example.org ban_account boby "Spammed rooms"
```

Call to [srg_create](#) API using double-quotes and single-quotes:

```
ejabberdctl srg_create g1 example.org "'Group number 1'" this_is_g1 g1
```

API Tags: [accounts](#), [erlang](#), [last](#), [private](#), [purge](#), [roster](#), [session](#), [shared_roster_group](#), [stanza](#), [statistics](#), [vcard](#)

mod_admin_update_sql

This module can be used to convert your existing SQL database from the singlehost to the multihost schema. Check the section [Singlehost or Multihost](#) for details. Please note that only MS SQL, MySQL, and PostgreSQL are supported. When the module is loaded use [update_sql](#) API.

The module has no options.

API Tags: [sql](#)

mod_announce

This module enables configured users to broadcast announcements and to set the message of the day (MOTD). Configured users can perform these actions with an XMPP client either using Ad-Hoc Commands or sending messages to specific JIDs. Equivalent API commands are also available.

**Note**

This module can be resource intensive on large deployments as it may broadcast a lot of messages. This module should be disabled for instances of ejabberd with hundreds of thousands users.

To send announcements using [XEP-0050: Ad-Hoc Commands](#), this module requires [mod_adhoc](#) (to execute the commands), and recommends [mod_disco](#) (to discover the commands).

To send announcements by sending messages to specific JIDs, these are the destination JIDs:

- `example.org/announce/all`: Send the message to all registered users in that vhost. If the user is online and connected to several resources, only the resource with the highest priority will receive the message. If the registered user is not connected, the message is stored offline in assumption that offline storage (see [mod_offline](#)) is enabled.
- `example.org/announce/online`: Send the message to all connected users. If the user is online and connected to several resources, all resources will receive the message.
- `example.org/announce/motd`: Set the message of the day (MOTD) that is sent to users when they login. Also sends the message to all connected users (similar to `announce/online`).
- `example.org/announce/motd/update`: Set the message of the day (MOTD) that is sent to users when they login. This does not send the message to any currently connected user.
- `example.org/announce/motd/delete`: Remove the existing message of the day (MOTD) by sending a message to this JID.

There are similar destination JIDs to apply to all virtual hosts in ejabberd:

- `example.org/announce/all-hosts/all`: send to all registered accounts
- `example.org/announce/all-hosts/online`: send to online sessions
- `example.org/announce/all-hosts/motd`: set MOTD and send to online
- `example.org/announce/all-hosts/motd/update`: update MOTD
- `example.org/announce/all-hosts/motd/delete`: delete MOTD

Available options:

- **access:** `AccessName`
This option specifies who is allowed to send announcements and to set the message of the day. The default value is `none` (i.e. nobody is able to send such messages).
- **cache_life_time:** `timeout()`
Same as top-level [cache_life_time](#) option, but applied to this module only.
- **cache_missed:** `true | false`
Same as top-level [cache_missed](#) option, but applied to this module only.
- **cache_size:** `pos_integer() | infinity`
Same as top-level [cache_size](#) option, but applied to this module only.
- **db_type:** `mnesia | sql`
Same as top-level [default_db](#) option, but applied to this module only.
- **use_cache:** `true | false`
Same as top-level [use_cache](#) option, but applied to this module only.

API Tags: [announce](#)

mod_antispam



added in 25.07

Filter spam messages and subscription requests received from remote servers based on [Real-Time Block Lists \(RTBL\)](#), lists of known spammer JIDs and/or URLs mentioned in spam messages. Traffic classified as spam is rejected with an error (and an *[info]* message is logged) unless the sender is subscribed to the recipient's presence.

Available options:

- **access_spam:** `Access`

Access rule that controls what accounts may receive spam messages. If the rule returns `allow` for a given recipient, spam messages aren't rejected for that recipient. The default value is `none`, which means that all recipients are subject to spam filtering verification.

- **cache_size:** `pos_integer()`

Maximum number of JIDs that will be cached due to sending spam URLs. If that limit is exceeded, the least recently used entries are removed from the cache. Setting this option to `0` disables the caching feature. Note that separate caches are used for each virtual host, and that the caches aren't distributed across cluster nodes. The default value is `10000`.

- **rtbl_services:** `[Service]`

Query a RTBL service to get domains to block, as provided by [xmppbl.org](#). Please note right now this option only supports one service in that list. For blocking spam and abuse on MUC channels, please use `mod_muc_rtbl` for now. If only the host is provided, the default node names will be assumed. If the node name is different than `spam_source_domains`, you can setup the custom node name with the option `spam_source_domains_node`. The default value is an empty list of services.

Example:

```
rtbl_services:
- pubsub.server1.localhost:
  spam_source_domains_node: actual_custom_pubsub_node
```

- **spam_domains_file:** `none | Path`

Path to a plain text file containing a list of known spam domains, one domain per line. Messages and subscription requests sent from one of the listed domains are classified as spam if sender is not in recipient's roster. This list of domains gets merged with the one retrieved by an RTBL host if any given. Use an absolute path, or the `@CONFIG_PATH@` [predefined keyword](#) if the file is available in the configuration directory. The default value is `none`.

- **spam_dump_file:** `false | true | Path`

Path to the file to store blocked messages. Use an absolute path, or the `@LOG_PATH@` [predefined keyword](#) to store logs in the same place that the other ejabberd log files. If set to `false`, it doesn't dump stanzas, which is the default. If set to `true`, it stores in `"@LOG_PATH@/spam_dump_@HOST@.log"`.

- **spam_jids_file:** `none | Path`

Path to a plain text file containing a list of known spammer JIDs, one JID per line. Messages and subscription requests sent from one of the listed JIDs are classified as spam. Messages containing at least one of the listed JIDs are classified as spam as well. Furthermore, the sender's JID will be cached, so that future traffic originating from that JID will also be classified as spam. Use an absolute path, or the `@CONFIG_PATH@` [predefined keyword](#) if the file is available in the configuration directory. The default value is `none`.

- **spam_urls_file:** `none | Path`

Path to a plain text file containing a list of URLs known to be mentioned in spam message bodies. Messages containing at least one of the listed URLs are classified as spam. Furthermore, the sender's JID will be cached, so that future traffic originating from that JID will be classified as spam as well. Use an absolute path, or the `@CONFIG_PATH@` [predefined keyword](#) if the file is available in the configuration directory. The default value is `none`.

- **whitelist_domains_file:** `none | Path`

Path to a file containing a list of domains to whitelist from being blocked, one per line. If either it is in `spam_domains_file` or more realistically in a domain sent by a RTBL host (see option `rtbl_services`) then this domain will be ignored and stanzas from there won't be blocked. Use an absolute path, or the `@CONFIG_PATH@` [predefined keyword](#) if the file is available in the configuration directory. The default value is `none`.

Example:

```
modules:
  mod_antispam:
    rtbl_services:
      - xmppbl.org
```

```
spam_jids_file: "@CONFIG_PATH@/spam_jids.txt"
spam_dump_file: "@LOG_PATH@/spam/host-@HOST@.log"
```

API Tags: spam

mod_auth_fast

 improved in 26.07

The module adds support for [XEP-0484: Fast Authentication Streamlining Tokens](#) that allows users to authenticate using self-managed tokens.

Available options:

- **db_type:** mnesia | sql
 improved in 26.07 Same as top-level [default_db](#) option, but applied to this module only.
- **token_lifetime:** timeout()
 Time that tokens will be kept, measured from it's creation time. Default value set to 30 days
- **token_refresh_age:** timeout()
 This time determines age of token, that qualifies for automatic refresh. Default value set to 1 day

Example:

```
modules:
  mod_auth_fast:
    token_lifetime: 14days
```

mod_avatar

The purpose of the module is to cope with legacy and modern XMPP clients posting avatars. The process is described in [XEP-0398: User Avatar to vCard-Based Avatars Conversion](#).

Also, the module supports conversion between avatar image formats on the fly.

The module depends on [mod_vcard](#), [mod_vcard_xupdate](#) and [mod_pubsub](#).

Available options:

- **convert:** {From: To}
 Defines image conversion rules: the format in `From` will be converted to format in `To`. The value of `From` can also be `default`, which is match-all rule. NOTE: the list of supported formats is detected at compile time depending on the image libraries installed in the system.

Example:

```
convert:
  webp: jpg
  default: png
```

- **rate_limit:** Number
 Limit any given JID by the number of avatars it is able to convert per minute. This is to protect the server from image conversion DoS. The default value is 10.

mod_block_strangers

This module blocks and logs any messages coming from an unknown entity. If a writing entity is not in your roster, you can let this module drop and/or log the message. By default you'll just not receive message from that entity. Enable this module if you want to drop SPAM messages.

Available options:

- **access:** `AccessName`

The option is supposed to be used when `allow_local_users` and `allow_transports*` are not enough. It's an Access Rule where `*deny` means the stanza will be rejected; there's an exception if option `captcha` is configured. And `allow` means the sender is whitelisted and the stanza will pass through. The default value is `none`, which means nothing is whitelisted.

- **allow_local_users:** `true | false`

This option specifies if strangers from the same local host should be accepted or not. The default value is `true`.

- **allow_transports:** `true | false`

If set to `true` and some server's JID is in user's roster, then messages from any user of this server are accepted even if no subscription present. The default value is `true`.

- **captcha:** `true | false`

Whether to generate CAPTCHA challenges in response to incoming presence subscription requests from strangers. See also section [CAPTCHA](#) of the Configuration Guide. The default value is `false`.

- **drop:** `true | false`

This option specifies if strangers messages should be dropped or not. The default value is `true`.

- **log:** `true | false`

This option specifies if strangers' messages should be logged (as info message) in ejabberd.log. The default value is `false`.

mod_blocking

The module implements [XEP-0191: Blocking Command](#).

This module depends on [mod_privacy](#) where all the configuration is performed.

The module has no options.

mod_bosh

This module implements XMPP over BOSH as defined in [XEP-0124](#) and [XEP-0206](#). BOSH stands for Bidirectional-streams Over Synchronous HTTP. It makes it possible to simulate long lived connections required by XMPP over the HTTP protocol. In practice, this module makes it possible to use XMPP in a browser without WebSocket support and more generally to have a way to use XMPP while having to get through an HTTP proxy.

Available options:

- **cache_life_time:** `timeout()`
Same as top-level [cache_life_time](#) option, but applied to this module only.
- **cache_missed:** `true` | `false`
Same as top-level [cache_missed](#) option, but applied to this module only.
- **cache_size:** `pos_integer()` | `infinity`
Same as top-level [cache_size](#) option, but applied to this module only.
- **json:** `true` | `false`
This option has no effect.
- **max_concat:** `pos_integer()` | `infinity`
This option limits the number of stanzas that the server will send in a single bosh request. The default value is `unlimited`.
- **max_inactivity:** `timeout()`
The option defines the maximum inactivity period. The default value is `30` seconds.
- **max_pause:** `pos_integer()`
Indicate the maximum length of a temporary session pause (in seconds) that a client can request. The default value is `120`.
- **prebind:** `true` | `false`
If enabled, the client can create the session without going through authentication. Basically, it creates a new session with anonymous authentication. The default value is `false`.
- **queue_type:** `ram` | `file`
Same as top-level [queue_type](#) option, but applied to this module only.
- **ram_db_type:** `mnesia` | `sql` | `redis`
Same as top-level [default_ram_db](#) option, but applied to this module only.
- **use_cache:** `true` | `false`
Same as top-level [use_cache](#) option, but applied to this module only.

Example:

```
listen:
-
  port: 5222
  module: ejabberd_c2s
-
  port: 5443
  module: ejabberd_http
  request_handlers:
    /bosh: mod_bosh

modules:
  mod_bosh: {}
```

mod_caps

This module implements [XEP-0115: Entity Capabilities](#). The main purpose of the module is to provide PEP functionality (see [mod_pubsub](#)).

Available options:

- **cache_life_time:** `timeout()`
Same as top-level [cache_life_time](#) option, but applied to this module only.
- **cache_missed:** `true` | `false`
Same as top-level [cache_missed](#) option, but applied to this module only.
- **cache_size:** `pos_integer()` | `infinity`
Same as top-level [cache_size](#) option, but applied to this module only.
- **db_type:** `mnesia` | `sql`
Same as top-level [default_db](#) option, but applied to this module only.
- **use_cache:** `true` | `false`
Same as top-level [use_cache](#) option, but applied to this module only.

mod_carboncopy

The module implements [XEP-0280: Message Carbons](#). The module broadcasts messages on all connected user resources (devices).

The module has no options.

mod_client_state

This module allows for queueing certain types of stanzas when a client indicates that the user is not actively using the client right now (see [XEP-0352: Client State Indication](#)). This can save bandwidth and resources.

A stanza is dropped from the queue if it's effectively obsoleted by a new one (e.g., a new presence stanza would replace an old one from the same client). The queue is flushed if a stanza arrives that won't be queued, or if the queue size reaches a certain limit (currently 100 stanzas), or if the client becomes active again.

Available options:

- **queue_chat_states:** `true` | `false`
Queue "standalone" chat state notifications (as defined in [XEP-0085: Chat State Notifications](#)) while a client indicates inactivity. The default value is `true`.
- **queue_pep:** `true` | `false`
Queue PEP notifications while a client is inactive. When the queue is flushed, only the most recent notification of a given PEP node is delivered. The default value is `true`.
- **queue_presence:** `true` | `false`
While a client is inactive, queue presence stanzas that indicate (un)availability. The default value is `true`.

mod_configure



improved in [25.10](#)

The module provides server configuration functionalities using [XEP-0030: Service Discovery](#) and [XEP-0050: Ad-Hoc Commands](#):

- List and discover outgoing s2s, online client sessions and all registered accounts
- Most of the ad-hoc commands defined in [XEP-0133: Service Administration](#)
- Additional custom ad-hoc commands specific to ejabberd

Ad-hoc commands from XEP-0133 that behave differently to the XEP:

- `get-user-roster` : returns standard fields instead of roster items that client cannot display

Those ad-hoc commands from XEP-0133 do not include in the response the client that executed the command:

- `get-active-users-num`
- `get-idle-users-num`
- `get-active-users`
- `get-idle-users`

Those ad-hoc commands from XEP-0133 are not implemented:

- `edit-blacklist`
- `edit-whitelist`
- `edit-admin`

This module requires [mod_adhoc](#) (to execute the commands), and recommends [mod_disco](#) (to discover the commands).

Please notice that all the ad-hoc commands implemented by this module have an equivalent [API Command](#) that you can execute using [mod_adhoc_api](#) or any other API frontend.

Available options:

- **access:** `AccessName`
 added in [25.03](#) This option defines which access rule will be used to control who is allowed to access the features provided by this module. The default value is `configure`.

Example:

```
acl:
  admin:
    user: sun@localhost

access_rules:
  configure:
    allow: admin

modules:
  mod_configure:
    access: configure
```

mod_conversejs

-  improved in [25.07](#)

This module serves a simple page for the [Converse](#) XMPP web browser client.

To use this module, in addition to adding it to the `modules` section, you must also enable it in `listen` → `ejabberd_http` → [request_handlers](#).

Make sure either [mod_bosh](#) or [ejabberd_http_ws](#) are enabled in at least one `request_handlers`.

When `conversejs_css` and `conversejs_script` are `auto`, by default they point to the public Converse client.

When this module is enabled in `modules`, it adds automatically a request handler and link in WebAdmin. .

This module is available since ejabberd [21.12](#).

Available options:

- **bosh_service_url:** `auto` | `BoshURL`
BOSH service URL to which Converse can connect to. The keyword `@HOST@` is replaced with the real virtual host name. If set to `auto`, it will build the URL of the first configured BOSH request handler. The default value is `auto`.
- **conversejs_css:** `auto` | `URL`
Converse CSS URL. The keyword `@HOST@` is replaced with the hostname. The default value is `auto`.
- **conversejs_options:** `{Name: Value}`
 added in 22.05 Specify additional options to be passed to Converse. See [Converse configuration](#). Only boolean, integer and string values are supported; lists are not supported.
- **conversejs_plugins:** `[Filename]`
List of additional local files to include as scripts in the homepage. Please make sure those files are available in the path specified in `conversejs_resources` option, in subdirectory `plugins/`. Notice that `libsignal` is only required for OMEMO encryption in Converse 13 or lower; Converse 14 uses and includes `libomemo`. If using the public Converse client, then `"libsignal"` gets replaced with the URL of the public library. The default value is `[]`.
- **conversejs_resources:** `Path`
 added in 22.05 Local path to the Converse files. If not set, the public Converse client will be used instead.
- **conversejs_script:** `auto` | `URL`
Converse main script URL. The keyword `@HOST@` is replaced with the hostname. The default value is `auto`.
- **default_domain:** `Domain`
Specify a domain to act as the default for user JIDs. The keyword `@HOST@` is replaced with the hostname. The default value is `@HOST@`.
- **websocket_url:** `auto` | `WebSocketURL`
A WebSocket URL to which Converse can connect to. The `@HOST@` keyword is replaced with the real virtual host name. If set to `auto`, it will build the URL of the first configured WebSocket request handler. The default value is `auto`.

Examples:

Manually setup WebSocket url, and use the public Converse client:

```
listen:
-
  port: 5280
  module: ejabberd_http
  request_handlers:
    /bosh: mod_bosh
    /websocket: ejabberd_http_ws
    /conversejs: mod_conversejs

modules:
  mod_bosh: {}
  mod_conversejs:
    websocket_url: "ws://@HOST@:5280/websocket"
```

Host Converse locally and let auto detection of WebSocket and Converse URLs:

```
listen:
-
  port: 443
  module: ejabberd_http
  tls: true
  request_handlers:
    /websocket: ejabberd_http_ws
    /conversejs: mod_conversejs

modules:
  mod_conversejs:
    conversejs_resources: "/home/ejabberd/conversejs-x.y.z/package/dist"
    conversejs_plugins: ["libsignal-protocol.min.js"]
    # File path is: /home/ejabberd/conversejs-x.y.z/package/dist/plugins/libsignal-protocol.min.js
```

Configure some additional options for Converse

```
modules:
  mod_conversejs:
    websocket_url: auto
    conversejs_options:
```

```

auto_away: 30
clear_cache_on_logout: true
i18n: "pt"
locked_domain: "@HOST@"
message_archiving: always
theme: dracula

```

mod_delegation

This module is an implementation of [XEP-0355: Namespace Delegation](#). Only admin mode has been implemented by now. Namespace delegation allows external services to handle IQ using specific namespace. This may be applied for external PEP service.

Warning

Security issue: Namespace delegation gives components access to sensitive data, so permission should be granted carefully, only if you trust the component.

Note

This module is complementary to [mod_privilege](#) but can also be used separately.

Available options:

- **namespaces:** {Namespace: Options}
If you want to delegate namespaces to a component, specify them in this option, and associate them to an access rule. The options are:
- **access:** AccessName The option defines which components are allowed for namespace delegation. The default value is `none`.
- **filtering:** Attributes The list of attributes. Currently not used.

Examples:

Make sure you do not delegate the same namespace to several services at the same time. As in the example provided later, to have the `sat-pubsub.example.org` component perform correctly disable the [mod_pubsub](#) module.

```

access_rules:
  external_pubsub:
    allow: external_component
  external_mam:
    allow: external_component

acl:
  external_component:
    server: sat-pubsub.example.org

modules:
  mod_delegation:
    namespaces:
      urn:xmpp:mam:1:
        access: external_mam
      http://jabber.org/protocol/pubsub:
        access: external_pubsub

```

mod_disco

This module adds support for [XEP-0030: Service Discovery](#). With this module enabled, services on your server can be discovered by XMPP clients.

Available options:

- **extra_domains:** [Domain, ...]
With this option, you can specify a list of extra domains that are added to the Service Discovery item list. The default value is an empty list.
- **name:** Name
A name of the server in the Service Discovery. This will only be displayed by special XMPP clients. The default value is `ejabberd`.
- **server_info:** [Info, ...]
Specify additional information about the server, as described in [XEP-0157: Contact Addresses for XMPP Services](#). Every `Info` element in the list is constructed from the following options:
- **modules:** all | [Module, ...] The value can be the keyword `all`, in which case the information is reported in all the services, or a list of ejabberd modules, in which case the information is only specified for the services provided by those modules.
- **name:** Name The field `var` name that will be defined. See XEP-0157 for some standardized names.
- **urls:** [URI, ...] A list of contact URIs, such as HTTP URLs, XMPP URIs and so on.

Example:

```
server_info:
-
  modules: all
  name: abuse-addresses
  urls: ["mailto:abuse@shakespeare.lit"]
-
  modules: [mod_muc]
  name: "Web chatroom logs"
  urls: ["http://www.example.org/muc-logs"]
-
  modules: [mod_disco]
  name: feedback-addresses
  urls:
    - http://shakespeare.lit/feedback.php
    - mailto:feedback@shakespeare.lit
    - xmpp:feedback@shakespeare.lit
-
  modules:
    - mod_disco
    - mod_vcard
  name: admin-addresses
  urls:
    - mailto:xmpp@shakespeare.lit
    - xmpp:admins@shakespeare.lit
```

mod_fail2ban

The module bans IPs that show the malicious signs. Currently only C2S authentication failures are detected.

Unlike the standalone program, `mod_fail2ban` clears the record of authentication failures after some time since the first failure or on a successful authentication. It also does not simply block network traffic, but provides the client with a descriptive error message.

**Warning**

You should not use this module behind a proxy or load balancer. ejabberd will see the failures as coming from the load balancer and, when the threshold of auth failures is reached, will reject all connections coming from the load balancer. You can lock all your user base out of ejabberd when using this module behind a proxy.

Available options:

- **access:** `AccessName`
Specify an access rule for whitelisting IP addresses or networks. If the rule returns `allow` for a given IP address, that address will never be banned. The `AccessName` should be of type `ip`. The default value is `none`.
- **c2s_auth_ban_lifetime:** `timeout()`
The lifetime of the IP ban caused by too many C2S authentication failures. The default value is `1` hour.
- **c2s_max_auth_failures:** `Number`
The number of C2S authentication failures to trigger the IP ban. The default value is `20`.

API Tags: [accounts](#)**mod_host_meta** added in [22.05](#)

This module serves small `host-meta` files as described in [XEP-0156: Discovering Alternative XMPP Connection Methods](#).

To use this module, in addition to adding it to the `modules` section, you must also enable it in `listen` → `ejabberd_http` → `request_handlers`.

Notice it only works if `ejabberd_http` has `tls` enabled.

Available options:

- **bosh_service_url:** `undefined | auto | BoshURL`
BOSH service URL to announce. The keyword `@HOST@` is replaced with the real virtual host name. If set to `auto`, it will build the URL of the first configured BOSH request handler. The default value is `auto`.
- **websocket_url:** `undefined | auto | WebSocketURL`
WebSocket URL to announce. The keyword `@HOST@` is replaced with the real virtual host name. If set to `auto`, it will build the URL of the first configured WebSocket request handler. The default value is `auto`.

Example:

```
listen:
-
  port: 443
  module: ejabberd_http
  tls: true
  request_handlers:
    /bosh: mod_bosh
    /websocket: ejabberd_http_ws
    /.well-known/host-meta: mod_host_meta
    /.well-known/host-meta.json: mod_host_meta

modules:
  mod_bosh: {}
  mod_host_meta:
    bosh_service_url: "https://@HOST@:5443/bosh"
    websocket_url: "wss://@HOST@:5443/websocket"
```

mod_http_api

This module provides a ReST interface to call [ejabberd API](#) commands using JSON data.

To use this module, in addition to adding it to the `modules` section, you must also enable it in `listen` → `ejabberd_http` → `request_handlers`.

To use a specific API version N, when defining the URL path in the `request_handlers`, add a `vN`. For example: `/api/v2: mod_http_api`.

To run a command, send a POST request to the corresponding URL: `http://localhost:5280/api/COMMAND-NAME`

Available options:

- **default_version:** `integer() | string()`



added in [24.12](#) What API version to use when none is specified in the URL path. If setting an ejabberd version, it will use the latest API version that was available in that ejabberd version. For example, setting `"24.06"` in this option implies `2`. The default value is the latest version.

Example:

```
listen:
-
  port: 5280
  module: ejabberd_http
  request_handlers:
    /api: mod_http_api

modules:
  mod_http_api:
    default_version: 2
```

mod_http_fileserver



improved `docroot` in [26.01](#)

This simple module serves files from the local disk over HTTP.

Available options:

- **accesslog:** Path

File to log accesses using an Apache-like format. No log will be recorded if this option is not specified.

- **content_types:** {Extension: Type}

Specify mappings of extension to content type. There are several content types already defined. With this option you can add new definitions or modify existing ones. The default values are:

Example:

```
content_types:
  .avi: video/avi
  .bmp: image/bmp
  .bz2: application/x-bzip2
  .css: text/css
  .gif: image/gif
  .gz: application/x-gzip
  .html: text/html
  .ico: image/vnd.microsoft.icon
  .jar: application/java-archive
  .jpeg: image/jpeg
  .jpg: image/jpeg
  .js: text/javascript
  .json: application/json
  .m4a: audio/mp4
  .map: application/json
  .mp3: audio/mpeg
  .mp4: video/mp4
  .mpeg: video/mpeg
  .mpg: video/mpeg
  .ogg: application/ogg
  .pdf: application/pdf
  .png: image/png
  .rtf: application/rtf
  .svg: image/svg+xml
  .tiff: image/tiff
  .ttf: font/ttf
  .txt: text/plain
  .wav: audio/wav
  .wasm: application/wasm
  .webp: image/webp
  .woff: font/woff
  .woff2: font/woff2
  .xml: application/xml
  .xpi: application/x-xpinstall
  .xul: application/vnd.mozilla.xul+xml
  .xz: application/x-xz
  .zip: application/zip
```

- **custom_headers:** {Name: Value}

Indicate custom HTTP headers to be included in all responses. There are no custom headers by default.

- **default_content_type:** Type

Specify the content type to use for unknown extensions. The default value is `application/octet-stream`.

- **directory_indices:** [Index, ...]

Indicate one or more directory index files, similarly to Apache's `DirectoryIndex` variable. When an HTTP request hits a directory instead of a regular file, those directory indices are looked in order, and the first one found is returned. The default value is an empty list.

- **docroot:** PathDir | {PathURL, PathDir}



improved in 26.01 Directory to serve the files from, or a map with several URL path (as specified in [request_handlers](#)) and their corresponding directory. This is a mandatory option.

Example:

```
listen:
-
  port: 5280
  module: ejabberd_http
  request_handlers:
    /pub/content: mod_http_fileserver
    /share: mod_http_fileserver
    /: mod_http_fileserver
  modules:
    mod_http_fileserver:
      docroot:
        /pub/content: /var/service/www
        /share: /usr/share/javascript
        /: /var/www
```

- **must_authenticate_with:** [{Username, Hostname}, ...]

List of accounts that are allowed to use this service. Default value: `[]`.